

Christopher Young

Sr ML Engineer | LLM Systems Specialist | AI & Backend Architect

✉ chris_yng1@hotmail.com ☎ 9566705277 📍 Princeton, New Jersey

🌐 [linkedin.com/in/chris-young-402541375](https://www.linkedin.com/in/chris-young-402541375)

Profile

Senior Machine Learning Engineer with **10+ years of experience** designing, training, and deploying ML systems end-to-end—from data ingestion to production MLOps. Deep hands-on background in **Python, PyTorch/TensorFlow, cloud platforms (AWS/Azure/GCP), and modern MLOps stacks (Docker, Kubernetes, CI/CD, MLflow)**.

Delivered solutions across **cloud hosting, healthcare, and public sector**: demand forecasting, anomaly detection, NLP for support and clinical data, recommendation systems, and more. Comfortable owning projects from idea → prototype → production, collaborating with product, data, and backend teams in **fully remote, distributed environments**—a pattern now common in remote ML/AI job descriptions.

Professional Experience

Senior Machine Learning Engineer, A2 Hosting

03/2018 – 10/2025

- Designed and deployed **end-to-end ML systems** for the hosting platform, including:
 - **Infrastructure anomaly detection** on logs/metrics (CPU, I/O, latency) to flag incidents early.
 - **Capacity and demand forecasting** for servers and storage to reduce over-provisioning.
 - **Customer churn and upsell models** to recommend plan upgrades and retention offers.
- Built **MLOps pipelines on AWS**, using S3 for data lake storage, ECS/EKS + Docker for model serving, and CI/CD pipelines (GitHub Actions) to automate testing, packaging, and deployment.
- Implemented **real-time scoring microservices** (FastAPI/Flask) that exposed REST endpoints consumed by the core hosting platform and customer dashboard.
- Introduced **MLflow/Kubeflow** for experiment tracking and model registry, improving reproducibility and collaboration with other engineers/data scientists.
- Developed **LLM-assisted support tools**, such as:
 - A **ticket routing model** to auto-classify support tickets (billing, performance, DNS, email) and reduce manual triage.
 - A **knowledge-base QA bot** (RAG style) over internal docs and runbooks to help support agents resolve issues faster.
- Collaborated closely with SRE and platform teams to integrate model outputs into **monitoring and alerting** Prometheus/Grafana), ensuring models improved reliability instead of adding noise.
- Drove **data quality initiatives**, defining SLAs for data pipelines, validation checks, and schema versioning to keep production models stable.

SW / Machine Learning Engineer, Medforce

01/2016 – 01/2018

- Worked with healthcare stakeholders (clinicians, operations, billing) to define and deliver **AI/ML solutions** for:
 - **Clinical text NLP**: extracting diagnoses, medications, and procedures from unstructured notes and reports.
 - **Risk and readmission prediction**: supervised models to estimate patient risk, feeding into care-management workflows.
- Built **data pipelines** ingesting EHR, claims, and lab data into a centralized warehouse (SQL + Spark), ensuring PHI handling aligned with security and privacy policies.

- Prototyped and productionized models using **Python, scikit-learn, and TensorFlow**, wrapping them in REST APIs for integration with internal portals and clinician dashboards.
- Implemented **model monitoring** (performance drift, data skew) and scheduled retraining jobs via Airflow, keeping models up-to-date with new patient cohorts.
- Collaborated with DevOps to deploy services on **Azure/AWS** using Docker and infrastructure-as-code templates, in line with how many current remote AI roles operate.

SW / Machine Learning Engineer, City of Lompic

01/2013 – 01/2016

- Built analytics and ML models to support **city operations and planning**, including:
 - **Traffic and congestion forecasting** from sensor and event data.
 - **Utility consumption prediction** (water, electricity) to support capacity planning.
 - **Service request prioritization** using text classification on 311-style requests.
- Created **ETL pipelines** in Python and SQL to bring data from disparate city systems (traffic sensors, billing systems, spreadsheets) into a unified analytics database.
- Developed dashboards and simple APIs to expose model outputs to planners and operations teams, working closely with non-technical stakeholders.
- Helped introduce **basic MLOps practices** (version control, environments, automated retraining scripts) that laid the groundwork for later cloud migrations.

Education

Bachelor's degree in Computer Science, Princeton University

05/2008 – 05/2013

Skills

Programming Languages

Python, Java, Javascript, Typescript, SQL, Bash

Data Science & ML

AWS Services, MLFlow, Pandas, Scipy, Numpy, Tensorflow, TFX, Scikit-Learn, PyTorch, PySpark, Snowflake, Redshift, OpenCV, PowerBI, Tableau, Sagemaker, Pandas, Numpy

DevOps & MLOps

AWS SageMaker, Azure AI, GCP Vertex, Bedrock, MLflow, Docker, Kubernetes, Terraform, GitHub Actions, Airflow, Jenkins

Backend Engineering

Python, FastAPI, Flask, REST, Kafka, Redis, GraphQL, Postgres, MongoDB, AWS Lambda, ECS, EKS, Cloud Run

LLM & GenAI Stack

LangChain, LangGraph, Autogen, RAG, OpenAI API, Claude, Cohere, Bedrock, G-Eval, Sentence Transformers, FAISS, Pinecone, BERT, GPT-3.5/, LlamaIndex, OpenSearch