

DEVI SUHITHA CHUNDRU

Chicago, IL • +1-309-240-2979 • devisuhithach@gmail.com • [linkedin.com/in/devisuhithachundru](https://www.linkedin.com/in/devisuhithachundru)

SUMMARY

Seasoned Data Engineer with 5+ years of experience delivering scalable, cloud-native solutions on Azure and Google Cloud Platform. Designed and implemented robust ETL/ELT pipelines using SQL, Python, Airflow, and Azure Data Factory to process data from APIs, databases, files, and streaming sources. Built and optimized data models using Pandas, PySpark, and Polars to create analytics-ready datasets that support key business decisions. Developed and maintained data warehouses in BigQuery, Snowflake, and Redshift. Managed version control workflows with Git and GitHub. Applied containerization using Docker and Kubernetes, and implemented Infrastructure as Code using Terraform. Created data catalogs, lineage documentation, and user guides to support platform users. Developed dashboards and visualizations in Tableau, Power BI, and Looker. Known for strong collaboration, clear communication, and delivering secure, compliant data solutions in fast-paced environments.

SKILLS

- **Cloud Platforms:** Google Cloud Platform (GCP), Microsoft Azure, AWS (for transferable context)
- **Programming Languages:** SQL, Python
- **Data Engineering & Pipelines:** ETL/ELT, batch & streaming data processing, data integration (APIs, files, databases, streaming), Airflow, Azure Data Factory
- **Data Warehousing & Modeling:** Snowflake, BigQuery, Redshift, dimensional modeling, schema design
- **Data Processing Frameworks:** Pandas, PySpark, Polars
- **Containerization & IaC:** Docker, Kubernetes, Terraform, Pulumi
- **Version Control & DevOps:** Git, GitHub, GitLab
- **Visualization & Documentation:** Tableau, Power BI, Looker, data catalogs, data lineage, user guides
- **Security & Compliance:** IAM, data security best practices, compliance-aware development
- **Collaboration & SDLC:** Agile development, requirements gathering, testing, deployment

EXPERIENCE

Ryan Specialty, Chicago, IL

Jul 2024 – Present

Azure Data Engineer

- Designed and implemented scalable ETL/ELT pipelines on Azure using Databricks, Spark, and Airflow to process high-volume structured and semi-structured data from APIs, files, and streaming sources.
- Built analytics-ready datasets by converting logical models into physical structures in Azure Data Lake Storage and Snowflake, enabling accurate and timely business reporting.
- Integrated disparate data systems including Kafka streams, REST APIs, and relational databases to support unified, real-time analytics and operational dashboards.
- Developed infrastructure as code using Terraform and automated deployments via Azure DevOps CI/CD, improving release speed and reducing deployment errors by 60%.
- Created and maintained secure, compliant data pipelines by implementing IAM controls, logging, and auditing policies aligned with company data governance standards.
- Participated in end-to-end software development lifecycle: from gathering data requirements and modeling through testing, deployment, and maintenance of data products.
- Authored clear documentation, data catalog entries, and lineage tracking to support data platform users and improve data discoverability across teams.
- Collaborated cross-functionally with business stakeholders and engineering teams to optimize data models for reporting, forecasting, and decision support.
- Developed visualizations and operational insights using Power BI and Looker for leadership to monitor KPIs and data quality trends.
- Stayed up-to-date with cloud and data engineering trends, continuously enhancing ingestion frameworks, automation strategies, and orchestration workflows.

Capital One, Chicago, IL

Jul 2023 – May 2024

AWS Data Engineer

- Collaborated with cross-functional teams to implement cloud-native data solutions based on high-level architectural designs and evolving business requirements.
- Built and maintained ELT pipelines using AWS Glue, Lambda, and Kinesis to ingest and transform structured and semi-structured data from APIs, files, and databases into Redshift and S3-based data lakes.
- Designed and implemented logical and physical data models in Redshift, optimizing performance for analytics and enabling faster decision-making across fraud and risk domains.
- Integrated data from disparate sources using DMS and Kinesis Data Streams, creating unified, analytics-ready datasets to support critical business operations.
- Automated infrastructure provisioning with Terraform and AWS CloudFormation, reducing environment setup time and enabling consistent, reproducible deployments.
- Orchestrated workflows using Apache Airflow and AWS Step Functions, improving data reliability and reducing end-to-end pipeline failures by 30%.
- Enforced IAM roles, RBAC policies, and KMS encryption to ensure secure and compliant data flows in line with company and regulatory standards (HIPAA, GDPR).
- Authored data catalogs, lineage documentation, and user-facing guides to support data discovery, transparency, and effective onboarding for analytics users.
- Delivered streaming and batch data solutions enabling real-time insights, enhancing responsiveness to operational events.
- Collaborated with cross-functional teams to translate high-level business needs into scalable data solutions, contributing to the full software development lifecycle.
- Explored modern data architecture trends including Docker-based containerization and Iceberg table formats to inform future platform evolution.

DXC Technology, Chennai, IN

Jun 2021 – Dec 2022

GCP Data Engineer

- Engineered data integration pipelines on Google Cloud Platform using Cloud Storage, BigQuery, and Data Fusion to consolidate data from APIs, files, relational databases, and streaming platforms like Kafka and Cloud Pub/Sub.
- Translated business requirements into logical data models and developed corresponding physical structures in BigQuery to enable fast, reliable access to analytics-ready datasets.
- Delivered production-grade data orchestration workflows using Cloud Composer (Airflow), automating ingestion, transformation, and delivery processes across multiple business units.
- Curated and optimized datasets for operational dashboards by building visualizations with Tableau, Power BI, and Looker, supporting timely, data-informed decisions.
- Facilitated high-quality data delivery by implementing validation frameworks and profiling logic, resulting in a 25% reduction in data quality issues across reporting layers.
- Authored data catalog entries, lineage documentation, and detailed user guides to support platform users and ensure clarity in data sourcing and transformation logic.
- Maintained secure and compliant data environments by applying IAM policies, encryption standards, and Cloud DLP to protect sensitive PII and ensure adherence to company policy.
- Applied version control best practices using Git and GitHub to track data pipeline changes, collaborate across teams, and ensure reproducibility in deployment workflows.
- Provisioned and managed infrastructure using Terraform, automating GCP resource creation while enabling scalable, cost-effective environments for analytics teams.
- Incorporated Polars and Pandas into transformation logic to accelerate data wrangling tasks, improving runtime efficiency of batch workflows by 20%.
- Participated in Agile development cycles, contributing to sprint planning, backlog grooming, and cross-functional discussions to align platform enhancements with evolving data needs.

- Actively monitored industry trends and evaluated tools such as Dataform, Vertex AI, and containerized deployments with Docker to modernize data architecture and pipeline design.

Bank Of America, Chennai, IN

Jul 2019 – May 2021

Data Engineer

- Constructed Java APIs (Commerce API) to process JSON payloads and interface with Cassandra, increasing customer transaction processing efficiency by 40% and streamlining financial data workflows while ensuring secure private data handling.
- Engineered scalable data pipelines using Azure Databricks, Python, Scala, and SQL. Leveraged Databricks Unity Catalog to manage data governance, ensuring 100% compliance with regulatory requirements and optimizing financial reporting workflows.
- Improved ETL workflows by utilizing Python, Spark SQL, S3, and Redshift, enhancing reporting accuracy by 30% and incorporating encryption at rest and in transit to meet PII data protection standards (HIPAA, GDPR).
- Analyzed large datasets using Hive and HBase to extract key customer transaction metrics, facilitating real-time business decisions and reducing response times by 25%.
- Optimized IaaS and PaaS services within Azure, integrating multi-cloud platforms (AWS, Azure) to enable seamless access to financial data and boosting operational efficiency by 20%.
- Implemented real-time data ingestion pipelines with Kafka and Spark Streaming, supporting enterprise-scale ingestion and real-time fraud detection, enhancing data monitoring accuracy by 30%.
- Designed framework-based Spark applications with PySpark and Spark SQL to transform large-scale customer datasets into actionable insights, driving a 25% increase in banking product engagement.
- Managed the Software Development Life Cycle (SDLC) for financial data pipelines, utilizing Apache Airflow for workflow orchestration and achieving a 15% improvement in delivery timelines.
- Refined financial data processing workflows using Pandas and SQL, enabling seamless integration and robust analysis of customer and financial data to support strategic decision-making.
- Streamlined data ingestion automation with Azure Data Factory and Databricks Workbooks, reducing manual intervention by 40% and accelerating operational efficiency in data processing.
- Developed real-time streaming solutions with Apache Flink, Kafka, and Hive, delivering deeper insights into customer behavior and supporting proactive fraud detection with faster response times.
- Migrated historical financial data to Azure Data Lake and Azure Databricks, employing framework-based pipelines to improve scalability and reduce data handling costs by 25%.
- Integrated SAP APIs into financial data workflows, enhancing enterprise-level data exchange and increasing cross-functional reporting accuracy by 20%.

EDUCATION

Bradley University

Master of Science in Data Science and Analytics, Computer Science

May 2024

Peoria, IL