

Harshavardhan A

+1 (940) 629-6364

harshaavardhan0@gmail.com

<https://www.linkedin.com/in/harshavardhan-a-927b01303/>

Professional Summary:

- Experienced Data Engineer with 8+ years in the field, specializing in developing scalable data processing and analytics solutions using Spark, PySpark, Hive, and Pig.
- Expertise in Spark SQL, PySpark, and Spark Streaming for real-time data processing, significantly improving data ingestion and analysis capabilities.
- Proficient in managing big data with HBase and Sqoop, ensuring efficient data storage and transfer across distributed systems.
- Demonstrated ability in Kafka for building robust data pipelines, enhancing data flow between applications and systems.
- Skilled in cloud computing environments, particularly AWS (EC2, S3, Redshift) and GCP (BigQuery, Dataflow, Dataproc), to deploy scalable and resilient data infrastructures.
- Proficient in Scala development within Scala IDE (Eclipse) and Linux Shell Scripting, optimizing data processing tasks and automation.
- Strong background in Python and PySpark programming, automating data processes and integrating data systems seamlessly.
- Expertise in Snowflake for cloud data warehousing and QlikView for data visualization, delivering comprehensive business intelligence solutions.
- Proficient in implementing CI/CD pipelines using GIT, Bitbucket, Jenkins, and Ansible, improving deployment processes and system reliability.
- Collaborated with the database team to implement SQL best practices, optimize queries, and enhance Power BI performance, resulting in a 50% reduction in report load times, a 20% improvement in response time, and a 15% decrease in data extraction load time.
- Deep understanding of Cloud architecture and management tools, including CloudFormation, CloudWatch, and ELK Stack, ensuring high availability and performance monitoring.
- Expertise in GCP's cloud services, such as G-Cloud Function, Apache Beam, and Cloud Pub/Sub, facilitating efficient data integration and real-time analytics.
- Skilled in Azure data services, including Azure Data Factory, Databricks, and Azure Data Lake, implementing advanced analytics solutions in the Azure cloud.
- Extensive experience with Unix/Linux systems administration, enhancing system security and performance for data applications.
- Expertise in Microservices architecture and Java development, building scalable, service-oriented applications and data systems.
- Demonstrated success in agile project environments, adept at using Jira for project management and maintaining high standards of software delivery.
- Solid foundation in data security and compliance, employing best practices in SSH, YAML, and WebLogic for secure application deployment.

Technical Skills:

Big Data Technologies:	PySpark, Spark, Hive, Pig, HBase, Sqoop, Kafka, Cloudera, Apache Hadoop, MapReduce, Cassandra, Zookeeper
Cloud Platform Services:	AWS (EC2, S3, EBS, ELB, RDS, SNS, SQS, VPC, Glacier, CloudFormation, Step Functions, API Gateway, CloudFront, CloudWatch, ELK Stack), GCP (Big Query, GCS Bucket, G-Cloud Function, Apache Beam, Cloud Dataflow, Data proc, Cloud Pub/Sub, Cloud Spanner, Cloud SQL, Data Catalog, Databricks), Azure (Data Factors)

Programming Languages &Scripting:	Pysark, Python, Scala, Java 1.7, Linux Shell Scripting, PowerShell, Shell Scripting, JavaScript.
Data Management & Analytics:	Snowflake, QlikView, Power BI, SQL Server, DynamoDB, MySQL, PostgreSQL, Cosmos DB, Redshift, GCP Dataflow, GCP Dataproc, Azure SQL, JDBC , PostgreSQL.
Development Tools & CI/CD:	Scala IDE (Eclipse), GIT, Maven, Jira, Bitbucket, Ansible, Terraform, Azure DevOps, JBOSS, WebSphere, WebLogic, SSH, YAML, Git, Jenkins, Terraform.
Operating Systems:	Unix, Linux, CentOS, Windows
Miscellaneous:	JSON, Microservices, Agile, Sencha ExtJS, Hibernate, JMS, Project Server, OpenShift

Professional Experience:

Client: Vanguard (Remote)

Nov'22-Present

Role: AWS Data Engineer

Responsibilities:

- Designed and implemented scalable data pipelines using Snowflake and PySpark for large-scale data processing and analytics, replacing traditional big data tools like Spark and Hive.
- Developed and managed Snowflake data sharing and secure data exchange capabilities for cross-organization collaboration.
- Utilized Snowflake Streams, Tasks, and PySpark for real-time data processing and change data capture (CDC) to support real-time analytics.
- Architected and deployed Snowflake data warehousing solutions, enabling advanced data analysis and reporting for business stakeholders.
- Configured and maintained Snowflake roles, permissions, and security policies to ensure data governance and compliance.
- Leveraged Snowflake's Time Travel feature for data recovery and historical data analysis, reducing downtime and improving data reliability.
- Built and optimized ETL pipelines using PySpark and AWS Glue to process large datasets efficiently.
- Designed and implemented Snowflake data lakes on AWS S3, enabling cost-effective storage and analytics on semi-structured and structured data.
- Utilized Snowflake's JSON and Variant data types for processing semi-structured data, enhancing flexibility in data pipelines.
- Automated data pipeline orchestration using AWS Step Functions, PySpark, and Snowflake Tasks, improving workflow efficiency.
- Developed dynamic, data-driven reports and dashboards using Snowflake's native integration with Tableau and QlikView for actionable business insights.
- Managed and optimized Snowflake queries and workloads using query profiling and performance tuning techniques.
- Integrated Snowflake with AWS Redshift and PySpark for hybrid data warehousing solutions, enabling cross-platform analytics.
- Implemented Snowflake's data cloning feature for rapid environment setup and testing, reducing development cycle times.
- Automated infrastructure provisioning for Snowflake and AWS using Terraform and CI/CD tools, ensuring consistent and repeatable deployments.

- Collaborated with cross-functional teams using GIT, Jira, and Confluence for project management and timely delivery of data solutions.
- Designed and developed robust, scalable solutions considering security, reliability, performance, and cost.
- Conducted performance tuning of Snowflake data stores and PySpark applications for optimal resource utilization and query performance.
- Engaged in continuous learning to stay updated with emerging Snowflake, PySpark, and AWS technologies and best practices.
- Integrated Snowflake with AWS S3 and PySpark for seamless data ingestion and storage, enabling efficient ELT (Extract, Load, Transform) processes.

Environment:

Snowflake, PySpark, AWS S3, AWS Lambda, AWS Glue, AWS Redshift, AWS Step Functions, AWS CloudWatch, Terraform, Power BI, JSON, Python, SQL, GIT, Jira, Confluence, AWS IAM, AWS VPC, AWS SNS, AWS SQS, ELK Stack, Unix/Linux.

Client: Health Partners, MN

May'20-Oct 22

Role: GCP Data Engineer

Responsibilities:

- Deployed and managed big data solutions using GCP BigQuery and PySpark for Health Partners' analytics needs, enabling efficient data analysis and reporting.
- Implemented data storage solutions using GCS Bucket, ensuring secure and scalable storage for Health Partners' vast data repositories.
- Developed serverless data processing functions with G-Cloud Function, enhancing Health Partners' ability to process data efficiently without managing servers.
- Utilized Apache Beam, PySpark, and Cloud Dataflow for building and running data processing pipelines, significantly improving Health Partners' data integration and ETL processes.
- Automated cloud resource management and script execution using Cloud Shell and GSUtil, streamlining Health Partners' cloud operations.
- Configured and maintained Dataproc clusters for running PySpark, Spark, and Hive workloads, optimizing Health Partners' data processing capabilities.
- Managed VM Instances for specific data processing tasks, providing scalable compute resources for Health Partners.
- Implemented and maintained databases in Cloud SQL, MySQL, Postgres, and SQL Server, supporting Health Partners' application data needs.
- Developed data transformation scripts in Python, PySpark, and Scala, enhancing Health Partners' data pipeline flexibility and efficiency.
- Optimized Health Partners' data warehouse operations using Spark SQL and PySpark for complex data queries and analyses.
- Prepared and cleaned data using GCP Data Prep, ensuring high-quality data for analysis and reporting at Health Partners.
- Orchestrated data pipelines using Cloud Composer and PySpark, automating Health Partners' data workflows and scheduling tasks efficiently.
- Ensured high availability and performance monitoring using Cloud Composer and CloudWatch.
- Created automated reporting dashboards integrated with business intelligence tools.
- Integrated Cloud Pub/Sub for event-driven data processing, enhancing Health Partners' real-time data capabilities.
- Managed large-scale data migrations using Cloud Storage Transfer Service, ensuring seamless data integration for Health Partners' global operations.
- Implemented and maintained globally distributed databases using Cloud Spanner, providing high availability and consistency for Health Partners' critical applications.

- Enhanced Health Partners' data governance and cataloging with Data Catalog, improving data discovery and metadata management.
- Leveraged GCP Databricks and PySpark for advanced data analytics and machine learning projects, driving Health Partners' innovation in data solutions.
- Utilized Cloud SQL for relational database management, supporting Health Partners' application development with robust data storage solutions.
- Architected scalable and reliable data processing workflows using GCP Dataflow, GCP Dataproc, and PySpark, meeting Health Partners' data analytics needs.
- Monitored and optimized cloud resource usage and performance, ensuring cost-effective operations for Health Partners.
- Developed custom data models and algorithms in Python, PySpark, and Scala, enhancing Health Partners' predictive analytics and machine learning capabilities.
- Collaborated with Health Partners' IT and business teams to define data requirements and build solutions that support strategic initiatives.
- Ensured data security and compliance with industry standards and regulations, protecting Health Partners' sensitive information.
- Stayed abreast of new GCP technologies and services, continuously improving Health Partners' cloud data engineering practices.
- Provided technical leadership and mentoring to junior data engineers, fostering a culture of excellence.

Environment: GCP, GCP BigQuery, GCS Bucket, G-Cloud Function, Apache Beam, PySpark, Cloud Dataflow, Cloud Shell, GSUtil, VM Instances, Cloud SQL, MySQL, Postgres, SQL Server, Python, Scala, Spark, Hive, Spark SQL, GCP Data Prep, GCP Dataflow, GCP Dataproc, Cloud Composer, Cloud Pub/Sub, Cloud Storage Transfer Service, Cloud Spanner, Cloud SQL, Data Catalog, GCP Databricks.

Client: Village, MD

Jan'18 to Apr'20

Role: AWS Data Engineer

Responsibilities:

- Led the migration of Village's data warehousing solutions to Amazon S3 and AWS Redshift, enhancing data storage and accessibility across global markets.
- Developed and deployed data pipelines in AWS Glue, automating data movement and transformation processes for Village's sales and marketing data.
- Utilized AWS EMR (Elastic MapReduce) to build and optimize data models, supporting Village's analytics on health and wellness product performance.
- Implemented real-time data processing using Amazon Kinesis and AWS Lambda, enabling dynamic pricing and inventory management.
- Configured and managed Hive and HBase on Amazon EMR for complex data queries, supporting Village's customer behavior analysis.
- Automated data ingestion from various sources using AWS Glue and Amazon Kinesis Data Firehose, streamlining Village's data consolidation efforts.
- Managed Village's Amazon S3 and DynamoDB for scalable, secure data storage, supporting global e-commerce operations.
- Employed EMR and Spark for processing large datasets in Amazon S3, optimizing Village's product recommendation engine.
- Developed SQL databases in Amazon RDS and Amazon Aurora, ensuring robust data

management for Village's sales and distribution networks.

- Architected secure data exchange mechanisms using SSH and Amazon ECR (Elastic Container Registry), enhancing Village's data integrity and security.
- Orchestrated continuous deployment and integration pipelines using AWS Code Pipeline and Git, improving Village's software development lifecycle.
- Managed application settings and configurations using YAML in AWS Elastic Beanstalk deployments, enhancing operational efficiency.
- Automated backend processes through AWS Lambda, enabling efficient data transactions for Village's loyalty programs.
- Integrated Amazon Quick Sight with Amazon S3 for dynamic reporting and dashboards, providing actionable insights into Village's product sales.
- Scripted complex Python applications for data analysis and machine learning models, driving Village's customer segmentation strategies.
- Configured Apache Tomcat application server environments for enterprise applications, supporting Village's online retail platform.
- Utilized Amazon RDS for developing, testing, and deploying scalable applications, enhancing Village's business intelligence solutions.
- Migrated data warehousing solutions to AWS Redshift and S3 for scalability, developed ETL processes using AWS Glue and PySpark, implemented RESTful API integrations for seamless data exchange, and optimized SQL-based data models in PostgreSQL and DynamoDB.
- Monitored and optimized data storage costs and performance in AWS, ensuring efficient resource use across Village's cloud environments.
- Collaborated with global teams using Jira and Maven, coordinating development efforts for Village's digital transformation projects.
- Implemented data security and compliance measures, aligning Village's data operations with GDPR and other regulatory standards.
- Developed custom ETL frameworks using Spark and AWS Glue in Amazon EMR, facilitating efficient data processing and analysis.
- Optimized Village's customer experience by leveraging data insights from DynamoDB, improving personalization and customer service.
- Designed and executed data backup and recovery strategies in AWS, safeguarding Village's critical business data against loss.
- Enabled data democratization for Village's decision-makers by integrating AWS Glue with business analytics tools, fostering a data-driven culture.
- Contributed to Village's IT infrastructure scalability by deploying AWS Elastic Beanstalk and AWS Lambda, supporting global expansion initiatives.

Environment: AWS Glue, Amazon EMR, Amazon S3, AWS Redshift, Spark, Hive, HBase, AWS Kinesis, AWS Lambda, DynamoDB, Amazon RDS, Amazon Aurora, Amazon Quick Sight, AWS Code Pipeline, AWS Elastic Beanstalk, AWS ECR, SSH, YAML, Apache Tomcat, Python, AWS Code Commit, Git, Maven, Jira.

Client: Byte ridge Software Private Limited

Aug'15 to Oct'17

Role: Hadoop Developer

Responsibilities:

- Designed and implemented Big Data solutions using Apache Hadoop, focusing on scalability and efficiency. Utilized Hive and Pig for data manipulation and querying, enhancing data analysis processes.
- Managed large-scale data storage with Cassandra and HBase, ensuring high availability and performance of NoSQL databases in support of Big Data initiatives.
- Developed and optimized Sqoop scripts for efficient data transfer between Hadoop and relational databases, streamlining data ingestion and integration workflows.
- Configured and maintained Hadoop clusters using Cloudera on CentOS, ensuring robust and scalable data processing environments. Monitored cluster health and

performance using Zookeeper.

- Leveraged Sencha ExtJS and Ajax for creating rich interactive web applications, improving user experience and functionality in Big Data projects.
- Implemented data processing pipelines in Apache Hadoop environment, utilizing Hive, Pig, and custom Java applications to support diverse data analysis requirements.
- Enhanced data security and integrity by integrating HBase and Cassandra with Hadoop ecosystems, applying best practices in NoSQL database management.
- Automated data workflow processes using Sqoop and Oozie, significantly reducing manual effort and improving reliability of data pipelines.
- Developed dynamic web applications using JavaScript, Hibernate, and Ajax, facilitating efficient data visualization and manipulation for Big Data analytics.
- Integrated JMS with WebLogic Application Server for messaging services, enabling reliable communication between distributed systems in Big Data environments.
- Utilized Eclipse and Web Services for developing and debugging complex Hadoop applications, enhancing developer productivity and project collaboration.
- Administered Unix and Windows servers hosting Big Data applications, ensuring optimal performance and security of the Hadoop ecosystem.

Environment: Hive, pig, Apache Hadoop, Cassandra, Sqoop, Big Data, HBase, Zookeeper, Cloudera, CentOS, No SQL, sencha extjs, java script, ajax, Hibernate, JMS, web logic Application server, Eclipse, Web services, Project Server, Unix, Windows.