

JAI PONNA

PHONE: 603-404-7715

EMAIL: JK.PONNA@GMAIL.COM

READY FOR IN-PERSON INTERVIEW

SUMMARY:

- 10+ years Senior Data Engineer with experience in AWS & Azure cloud

SKILLS:

- Data Warehouse: Snowflake | Databricks | AWS Redshift
- Apache: Airflow | Flink | Kafka | DBT, Spark, Hive, Tableau
- AWS cloud: ECS | EC2 | S3 | EKS | Cloudformation | Glue | Kinesis
- Azure Cloud: Data Factory | Synapse Analytics | DevOps | Power BI
- SQL: PostgreSQL | Microsoft SQL | MySQL
- NoSQL: Cassandra | AWS DynamoDB
- Programming languages: Python | NumPy | Pandas | Scala
- Linux: Ubuntu | CentOS | Bash
- CI/CD: Jenkins | Docker | Kubernetes | Git | Bitbucket | CodePipelines

EXPERIENCE:

Senior Data Engineer @ AKRA Investment Services LLC, Boston, MA 10/2024 – Till date

- Architected and implemented a scalable data pipeline using **AWS** cloud services such as **EC2** and **ECS** with **Fargate** to reduce the time of managing infrastructure by 50%.
- Orchestrated **Airflow** for building pipelines, **Kafka** for data batch processing and **Flink** for data stream processing, resulting in a 40% reduction in data processing time for streaming data lag times that originally lagged in previous pipelines.
- Developed high-performance **ETL** jobs to process daily financial data from various sources into **Snowflake** data warehouse and **PostgreSQL** database improving query performance by 60%.
- Implemented a robust **CI/CD** pipeline using **Git**, **Codecommit**, **Codepipelines** and **Jenkins** for automated deployment of data engineering code to the **ECS**, reducing deployment time by 70% and minimizing errors.

Software Data Engineer @ Contactspace Ltd, Sydney, Australia 07/2021 – 10/2024

- Architected a real-time data pipeline using AWS **Kinesis** and **Kafka** to ingest and process call center data.
- Leveraged AWS **Glue** for **ETL** jobs for distributed processing with **Spark** hosted on **ECS**, resulting in a 60% reduction in infrastructure management time.
- Implemented a scalable data warehouse solution using **Redshift**, optimizing **SQL** queries for efficient schemas to handle 10 million daily call records, while utilizing **DynamoDB** and **S3** for cost-effective storage of historical data.
- Developed serverless workflow using **Lambda** and **Kubernetes EKS** containers, extracting and transforming of data from **NoSQL DynamoDB**, reducing processing time by 75%.
- Created interactive dashboards using **Tableau**.

- Used **Python** and **Scala** based programs of real-time data streams, providing insights into performance and customer satisfaction metrics for over 500 agents.

Senior Data Engineer @ NSW Dept of Education, Sydney, Australia 09/2016 – 07/2021

- Architected and implemented a scalable data pipeline using Azure Cloud **Data Factory** to process daily student performance data from 500+ schools.
- Leveraged **Databricks** for data warehousing and transformation. Used **Python** for programming.
- Developed an automated **ETL** workflow using **Synapse Analytics** and **Azure Functions** to integrate diverse educational datasets, implementing advanced **SQL** queries for data quality checks and utilizing **Docker** containers for consistent deployment across environments, resulting in a 30% improvement in data processing efficiency.
- Implemented a robust **CI/CD** pipeline using **Azure DevOps** and **Jenkins** for seamless integration and deployment of data engineering solutions reducing deployment time by 50%.
- Created interactive dashboards and reports using **Power BI**, integrating data from **SQL** databases and **Azure Virtual Machines**, while implementing row-level security with **Azure Active Directory** to ensure data privacy for over 100,000 student records.

Software & Cloud Engineer @ Coreware Ltd, London, UK 09/2012 – 07/2016

- Developed a real-time news aggregation pipeline for Reuters using **Python**, **Spark** and **Kafka** processing over 100,000 articles daily from diverse sources and implemented **SQL**-based data quality checks, resulting in a 40% increase in content relevance for subscribers.
- Architected a scalable data lake solution using **Hadoop** ecosystem with **HDFS**, **MapReduce**, **Hive** leveraging **HiveQL** to store and process historical news data, while enabling journalists to identify trending topics 30% faster.
- Implemented a robust **ETL** workflow using **MySQL** and **PostgreSQL** for structured data storage, with **Docker** and **Kubernetes** for containerization and orchestration, improving system reliability by 50% and reducing deployment time by 60%.
- Created an automated data pipeline using **Git** for version control and **Jenkins** for **CI/CD**, deploying to **CentOS** servers, while utilizing **Jupyter** Notebooks for exploratory data analysis and model prototyping, enhancing the news recommendation engine's accuracy by 25%.

EDUCATION:

- Masters in Information Systems 2012 – University of East Anglia, UK
- Bachelors in Computer Science 2009 – Jawaharlal Nehru Tech University, India