

# JAVEED SHAIK

## Data Engineer

California, USA | +1 (917) 379-8162 | [sjaveed.de@gmail.com](mailto:sjaveed.de@gmail.com)

---

### SUMMARY

**Data Engineer** with **6 years** of experience across **finance, consulting, healthcare, and e-commerce** domains. Delivery scope spans **60+ ETL pipelines** constructed using **Apache Spark, NiFi, and Python** to process over **2.4 TB** of data daily from **15+ enterprise sources**. Responsibilities include administration of **200+ orchestration schedules** with **Apache Airflow, AWS Glue, and Azure Data Factory** supporting **fraud analytics and regulatory workflows**. Work includes performance tuning for **Oracle, Redshift, and BigQuery** queries accessed by **120+ reporting users**. Experience spans **cloud-native ingestion platforms** built on **Snowflake, Vertex AI, and GCP**, supporting **ML readiness and data lineage compliance**. Visualization efforts support **30+ executive dashboards** created using **Power BI, Tableau, and Looker** for business units managing **\$1.2B+** in assets.

---

### PROFESSIONAL EXPERIENCE

#### JPMorgan Chase & Co. | USA

##### Data Engineer

*November 2022 – Present*

- Constructed a modular ingestion framework using **Apache Airflow**, orchestrating **180+ daily batch workflows** across trading and compliance, ensuring delivery of **2.4 TB/day** of financial records for downstream risk engines.
- Programmed scalable pipelines using **AWS Glue and Python** to move **12M+ transaction records** into **Amazon Redshift**, reducing ingestion-to-report latency from **6 hours to 45 minutes** across treasury operations.
- Integrated 20+ systems including CRM and loan origination into a unified **Snowflake** data warehouse, improving reconciliation accuracy for audit teams reviewing **1.3B+ cross-system records** quarterly.
- Developed distributed fraud detection jobs using **Apache Spark (PySpark)**, processing **75K+ events/second** and detecting anomalies across 3 internal risk platforms with **5x faster throughput** than legacy methods.
- Applied **advanced SQL** techniques (CTEs, partitioning, window functions) in **Oracle** to optimize query sets powering **120+ analyst dashboards**, reducing cumulative runtime by **over 17 hours/day**.
- Automated anomaly detection and rule-based validation using **Python**, flagging inconsistencies in **500M+ records/month**, eliminating over **85 hours/week** of manual data review for compliance teams.
- Established metadata tracing in **Azure Data Factory**, linking **15+ pipelines** and ensuring end-to-end lineage across datasets supporting **\$5B+ in asset reporting** under Basel regulatory standards.

#### KPMG | India

##### Data Engineer

*January 2020 – October 2021*

- Devised an **Apache NiFi** pipeline to consolidate **12M+ daily records** from 12 disparate client systems, enabling faster cross-audit access for the risk advisory division.
- Produced **ETL workflows in AWS Glue** processing **2 TB/day** for an e-commerce transformation program, reducing batch window by over **5 hours** and supporting SLA-bound analytics updates.
- Assembled a **Snowflake data warehouse** for a ₹1,000 crore infrastructure engagement, enabling concurrent access by 20+ departments and reducing query response lag across **85+ reporting jobs**.
- Shaped dimensional models using **Star Schema** for KPI monitoring across 6 verticals, facilitating dashboard accuracy over **25 tracked metrics** for consulting leadership.
- Executed **CI/CD automation in Jenkins**, cutting deployment time from 90 to under 10 minutes and reducing configuration error frequency from 15 to fewer than 2 incidents per month.

#### Trigent Software | India

##### Data Analyst / Data Engineer

*May 2018 – December 2019*

- Delivered a scalable ETL solution with **Apache Airflow** to ingest structured data from multiple clients, replacing over **140 manual hours/month** with automated pipelines in a finance platform.
- Processed **2 TB/day** of insurance data using **Apache Spark (PySpark)** on **HDFS**, improving legacy batch cycle performance for a U.S.-based modernization project by over **6 hours**.
- Managed **data lake integrations on AWS S3 and EMR**, optimizing query execution over **500M+ records**, and supporting real-time insights for a supply chain analytics use case.
- Deployed **data replication pipelines using Informatica**, enabling consistent synchronization between hybrid databases and maintaining SLA compliance across **4 live production systems**.

- Modeled retail analytics dashboards with **Snowflake schema** in **Amazon Redshift**, supporting **300+ business users** and reducing report generation time across 40+ KPIs.

---

## TECHNICAL SKILLS

**Programming & Scripting:** Python, SQL, Scala, Shell Scripting

**Big Data Technologies:** Apache Spark (PySpark/Scala), Hadoop, Hive, HDFS, Apache NiFi

**ETL & Data Orchestration:** Apache Airflow, AWS Glue, Talend, Informatica, Azure Data Factory, Google Cloud Dataflow

**Cloud Platforms:** AWS (S3, EMR, Redshift, Glue, Lambda), Azure (Data Factory, Synapse), GCP (BigQuery, Vertex AI, Pub/Sub, Dataproc)

**Data Warehousing:** Snowflake, Amazon Redshift, Google BigQuery, Teradata

**Databases:** PostgreSQL, MySQL, Oracle, MongoDB, Cosmos DB

**CI/CD & DevOps:** Docker, Jenkins, Terraform, Git, GitHub, GitLab

**Data Modeling & Governance:** Star Schema, Snowflake Schema, Data Lineage, Data Quality, RBAC

**Analytics & Visualization:** Power BI, Tableau, Looker, Seaborn, Superset

**Machine Learning Integration:** TensorFlow, Scikit-learn, XGBoost, MLflow, Vertex Pipelines

**Security & Compliance:** IAM, VPC Service Controls, Data Encryption, GDPR, PCI-DSS, SOX

---

## PROJECTS

### Scalable Fraud Detection Pipeline – GCP + ML Deployment

- Deployed a fraud analytics pipeline using **Google Pub/Sub**, **Cloud Dataflow**, and **BigQuery**, processing **1B+ transactions/day** with sub-second latency, and integrated **Vertex AI Pipelines** for real-time scoring across 8 endpoints with model versioning and drift monitoring.

### Cloud-Native ETL Platform Modernization

- Built lineage-aware ETL workflows across 5 enterprise systems using **Apache Spark**, **Airflow**, and **AWS Glue**, while introducing **CI/CD automation** with **Docker**, **Terraform**, and **Jenkins** to reduce batch runtime by 6 hours and enable controlled multi-environment deployment.

### Predictive Churn Modeling Engine – Azure

- Formulated feature engineering pipelines with **Azure Data Factory**, **Delta Lake**, and **PySpark**, managing training data from **100K+ user records**, and operationalized churn prediction for 12 segments using **Scikit-learn**, **XGBoost**, and **MLflow** with automated tracking and model registry.

---

## EDUCATION

|                                                                 |             |
|-----------------------------------------------------------------|-------------|
| <b>International Technological University   California, USA</b> | <b>2024</b> |
| Master of Science in Computer Science                           |             |
| <b>K L University   Andhra Pradesh, India</b>                   | <b>2017</b> |
| Bachelor of Technology in Biotechnology                         |             |