

Lokesh Reddy

Data Engineer

214-899-0998 | Frisco, TX | lokeshreddy.dataengr@gmail.com

SUMMARY

- Data Engineer with **4+ years** of hands-on experience designing, building, and optimizing **scalable data pipelines** and **ETL processes** across **cloud** and **on-premises** environments.
- Proven track record of leveraging tools like **Apache Airflow**, **Azure Data Factory (ADF)**, **Databricks**, and **Informatica PowerCenter** to manage complex data workflows in **healthcare** and **financial** sectors.
- Strong expertise in **SQL (T-SQL, PL/SQL)**, **Python**, and **PySpark** for **data transformation**, **automation**, and **real-time ingestion** using **APIs** and streaming platforms like **Kafka** and **NiFi**.
- Skilled in architecting **data models (star/snowflake schemas)**, integrating **structured/semi-structured data** using **Azure Synapse**, **AWS EMR**, **Hive**, and **Snowflake** for analytics and **BI reporting**.
- Adept at implementing **secure, compliant (HIPAA/PHI)** data practices and **DevOps workflows** using **Git**, **Azure DevOps**, **Jenkins**, **Docker**, and **Kubernetes**.
- Collaborative **Agile** team player with strong communication and **documentation skills**, contributing to **end-to-end data solutions** from design to monitoring.

TECHNICAL SKILLS

- **Languages & Scripting:** SQL (T-SQL, PL/SQL), Python (Pandas, PySpark), Shell Scripting, Scala
- **ETL & Workflow Orchestration:** Apache Airflow, Azure Data Factory (ADF), Informatica PowerCenter, Apache NiFi, AWS Glue, Control-M
- **Data Processing & Analytics:** Apache Spark, Databricks, PySpark, Hive, Impala, HBase, AWS EMR, Azure Synapse Analytics
- **Databases:** SQL Server, Oracle, PostgreSQL, MySQL, Snowflake
- **Cloud Platforms & Services:** Microsoft Azure (ADF, Data Lake Gen2, Synapse, DevOps, Repos), Amazon Web Services (S3, EC2, Glue, Redshift, Lambda)
- **Data Modeling & Warehousing:** Star Schema, Snowflake Schema, Dimensional Modeling
- **Visualization & BI Tools:** Power BI, Tableau, Excel
- **Version Control & CI/CD:** Git, Azure DevOps, Bitbucket, Jenkins
- **Security & Compliance:** HIPAA, PHI, RBAC, Data Masking, Encryption, Kerberos, LDAP, Apache Ranger
- **Containerization & Monitoring:** Docker, Kubernetes, Prometheus, Grafana
- **Agile & Collaboration Tools:** JIRA, Confluence, Agile/Scrum, TDD
- **Testing & Quality Assurance:** PyTest, Great Expectations, Data Validation, Data Lineage Tracking

WORK EXPERIENCE

Data Engineer

Banner Health, Phoenix, AZ

02/2024 – Present

- Built and maintained data pipelines using **Apache Airflow**, **Azure Data Factory (ADF)**, and **Informatica PowerCenter** to support enterprise data workflows and reporting needs.
- Developed SQL scripts and stored procedures in **T-SQL** and **PL/SQL** for data extraction, transformation, and loading across **SQL Server**, **Oracle**, and **PostgreSQL** systems.
- Utilized **Python** and **PySpark** to automate ETL processes, transform large datasets, and perform API-based data ingestion from third-party systems.
- Processed and integrated structured and semi-structured data using **Azure Data Lake (Gen2)** and **Azure Synapse Analytics** for centralized data storage and querying.
- Created data pipelines and transformations in **Databricks** using Spark to enable high-volume data processing and analytics.
- Configured data ingestion flows using **Apache NiFi** for real-time streaming and batch processing across internal systems.
- Ensured secure data handling practices in alignment with **HIPAA** and **PHI** compliance standards, including **data masking**, **encryption**, and **RBAC**.
- Designed and maintained **star schema** and **snowflake schema** data models to optimize reporting and analytics within enterprise data warehouses.
- Built interactive dashboards and self-service reports in **Power BI**, enabling real-time insights for clinical and operational users.
- Conducted unit testing and data quality checks using **PyTest** and **Great Expectations** to ensure pipeline integrity and reliability.
- Utilized **Azure DevOps**, **Git**, and **Azure Repos** for version control, code collaboration, and deployment of data engineering solutions.
- Participated in Agile sprints using **JIRA** for task tracking and collaborated via **Confluence** for technical documentation and process standardization.

- Used **SSIS** and **SSRS** within the Microsoft stack to support legacy ETL workflows and generate scheduled operational reports.

Data Engineer

Goldman Sachs, India

07/2021 – 07/2023

- Designed and maintained scalable ETL/ELT pipelines to process structured and semi-structured data from diverse internal and external sources using **Apache Spark**, **Airflow**, and **Python**.
- Engineered batch and real-time data ingestion workflows utilizing **Kafka**, **NiFi**, and **Hadoop HDFS**, optimizing for performance and reliability.
- Developed data models (star/snowflake schemas) to support analytical and operational use cases across business domains using **PostgreSQL**, **Oracle**, and **SQL Server**.
- Implemented distributed data processing solutions with **PySpark** and **Scala Spark** on **AWS EMR** for large-scale data transformation workloads.
- Performed data warehousing tasks using **Apache Hive**, **Impala**, and **HBase**, supporting both historical data analysis and near-real-time insights.
- Utilized **AWS services** including **S3**, **EC2**, **Lambda**, and **Glue** for cloud-based data storage, compute, and ETL orchestration.
- Employed **Docker** and **Kubernetes** to containerize data pipelines and support environment-agnostic deployments in cloud and hybrid setups.
- Authored complex **SQL queries** for data analysis, validation, and transformation; automated recurring data validation tasks using **Shell scripting** and **cron jobs**.
- Managed source control and CI/CD processes using **Git**, **Bitbucket**, and **Jenkins** to support agile development workflows.
- Ensured secure access to data pipelines and storage systems by integrating **Kerberos**, **LDAP**, and **Apache Ranger** for authentication and authorization.
- Developed data lineage tracking and metadata management solutions using internal frameworks and tools comparable to **Collibra**.
- Built and maintained dashboards and monitoring systems using **Prometheus**, **Grafana**, and **Tableau** to ensure data pipeline reliability and performance.
- Followed Agile methodologies and **Test-Driven Development (TDD)** practices for iterative, quality-focused development.
- Adhered to **secure coding standards** and compliance requirements across all development and deployment activities.

Junior Data Engineer

Wipro, India

01/2021 – 06/2021

- Developed and maintained SQL queries to extract, transform, and load data from multiple relational databases including Oracle and SQL Server for internal reporting systems.
- Automated routine data ingestion tasks using Python scripts, incorporating pandas for data manipulation and data validation.
- Participated in the design and execution of ETL workflows using **Informatica PowerCenter** and **Apache Nifi** for client data integration projects.
- Executed batch jobs and workflow scheduling using **Apache Airflow** and **Control-M**, ensuring smooth daily pipeline runs.
- Assisted in managing data pipelines built on the **Hadoop ecosystem**, including using **Hive** for querying large datasets and **HDFS** for data storage.
- Gained exposure to distributed data processing with **Apache Spark** for handling high-volume client data efficiently.
- Conducted data cleaning and transformation tasks across **MySQL** and **SQL Server** databases, ensuring consistency for downstream analytics.
- Supported cloud-based data processing tasks on **AWS**, utilizing services such as **S3**, **Glue**, and **Redshift** for storage and ETL operations.
- Created basic visualizations and validation dashboards using **Power BI** and Excel to support QA and reporting teams.
- Used **Git** and **Bitbucket** for version control during collaborative data pipeline development and troubleshooting.
- Participated in Agile sprint ceremonies and collaborated with cross-functional teams using **JIRA** and **Confluence** for documentation and task tracking.
- Documented data pipeline structures, job logic, and troubleshooting steps to support knowledge transfer and project continuity.

EDUCATION

Master of Science (M.S.) in Computer Science – University of North Texas, Denton, TX