



Naren
Phone no: +1 (469)-659-0495
Email: narendrakumar000@gmail.com
Linked in URL: [LinkedIn](#)



Professional Summary

- **Over 5 years of experience as a Data Engineer** with expertise in designing, building, and maintaining data pipelines, data warehouses, and analytics solutions across **Finance, Healthcare, and Logistics** domains.
- **Proficient in Snowflake Cloud Data Platform**, including data warehouse design, performance tuning, and advanced SQL capabilities for large-scale data processing.
- Skilled in developing and optimizing **ETL pipelines** using **Python, SQL, Azure Data Factory**, and **GCP Dataflow** to process high-volume structured and semi-structured data.
- Expertise in **data modeling** with **Erwin** and **dbt**, implementing dimensional models (Star and Snowflake schemas) to support analytics and reporting needs.
- Strong knowledge of **cloud platforms**, including **Azure** and **Google Cloud Platform (GCP)**, leveraging services like **BigQuery, Azure Synapse Analytics, Azure Data Lake**, and **Cloud Storage** for scalable data processing.
- Experience in integrating real-time data streams using **Kafka, Pub/Sub**, and batch workflows for high-performance data pipelines.
- Proficient in creating interactive dashboards and reports using **Power BI, Tableau**, and **Google Data Studio**, delivering insights to business stakeholders.
- Hands-on experience with **NoSQL databases** like **Azure Cosmos DB** and **HBase** for managing semi-structured and unstructured data.
- Adept in **CI/CD practices**, automating deployment workflows with tools like **Jenkins** and **GitHub Actions** for reliable and efficient data pipeline management.
- Extensive experience in working with healthcare standards like **HL7** and ensuring compliance with regulatory requirements such as **HIPAA** and **SOX**.
- Skilled in **data transformation, data quality validation**, and **metadata management**, ensuring data integrity and governance across projects.
- Strong knowledge of Agile methodologies, including sprint planning, daily stand-ups, and retrospectives, ensuring timely delivery of data solutions.
- Excellent problem-solving and communication skills, with a proven track record of working collaboratively with cross-functional teams to deliver data-driven solutions that drive business value.

Technical Skills

Operating Systems	Linux, Windows XP/7/8/10, Mac.
Programming languages	Python, Core Java, C, JavaScript, R, Scala.
BigData Technologies	HDFS, Hive, MapReduce, Pig, Sqoop, Flume, Oozie, Hadoop distribution, and HBase, Spark, Spark Streaming, Yarn, Zookeeper, Kafka, ETL. (Nifi, Talend etc.)
Data Warehousing and Modeling	Snowflake, Azure Synapse Analytics, Teradata, Oracle 19c, SQL Server 2022, Dimensional Modeling (Star Schema, Snowflake Schema), Data Vault, Erwin, ER/Studio
Databases	SQL Server 2022, Oracle 19c, Teradata V17, Teradata, HBase 2.3, Azure Cosmos DB 3.32, Snowflake DB.
Utilities/Tools	Eclipse, Tomcat, NetBeans, JUnit, SQL, SVN, Terraform, Log4j, SOAP UI, ANT, Maven, Alteryx, Visio, Jenkins, Jira, IntelliJ, Informaticav10.5
Cloud Services	GCP (Compute Engine, Cloud Storage, Dataproc, Cloud SQL, Cloud Functions, Cloud Monitoring, Filestore, Autoscaler, BigQuery, Deployment Manager, Dataflow, etc.) Azure (Azure Databricks, SaaS, Azure Synapse Analytics, Azure Service Bus, Azure HDInsight, Azure Synapse Analytics, Azure Data Factory, Azure Storage Azure Service), AWS
Data Visualization & ETL Tools	Tableau, Power BI, Google Data Studio, dbt, Informatica PowerCenter, SSIS, Talend, Azure Data Factory, Apache Nifi.
Version Control& Containerization tools	Git, GitHub, Bitbucket, Jenkins, Docker, Terraform
Software Life Cycle	Agile, Waterfall, RUP, RAD and JAD and Agile models.

Education	
University of North Texas, Master's in data science GPA: 3.78/4.0	Aug 2022 - May 2024
Work Experience	
Client – UPS, Atlanta, GA Data Engineer Responsibilities - Designed and built scalable ETL pipelines using dbt, Python, PySpark, and SQL to ingest, transform, and load structured and unstructured data into Snowflake. - Automated data workflows using Azure Data Factory to streamline data integration from multiple sources, including IoT devices, enterprise systems, and external APIs. - Developed optimized data models in Snowflake for analytics and reporting, ensuring high performance and scalability. - Implemented best practices for data partitioning, clustering, and query optimization to handle large-scale datasets effectively. - Configured and optimized Snowflake virtual warehouses for cost-efficiency and performance. - Leveraged Snowflake's Time Travel and Zero-Copy Cloning features for enhanced data management and backup strategies. - Utilized Snowflake's Data Sharing capabilities to enable seamless collaboration between internal teams and external partners. - Managed Snowflake's Access Control Framework to ensure data security and compliance with organizational policies and external regulations. - Conducted performance tuning of queries and workflows to minimize processing time and reduce compute costs. - Collaborated with business stakeholders to design interactive dashboards in Power BI, providing real-time insights into key performance indicators (KPIs) like delivery times, route efficiency, and package statuses. - Enabled self-service analytics capabilities for non-technical users by building user-friendly data models and views. - Leveraged Azure Synapse Analytics for data integration and analysis to enable cross-cloud compatibility and seamless scalability. - Used Azure Data Lake for staging and storing raw and processed data, ensuring efficient handling of large-scale datasets. - Implemented robust access controls, encryption mechanisms, and row-level security in Snowflake and Azure Key Vault to ensure compliance with internal and external data security regulations. - Deployed and monitored solutions using Azure Monitor and Azure Log Analytics to ensure the reliability and performance of data workflows. - Ensured data security and compliance with internal and external regulations by implementing robust access controls and encryption mechanisms. - Utilized Apache Kafka and Azure Event Hubs to enable real-time data streaming and processing, integrating with Snowflake and Azure Synapse Analytics to support time-sensitive decision-making and analytics. - Leveraged Apache Flink and Spark Streaming for real-time data processing, ensuring high throughput and low-latency data handling in large-scale distributed environments, aligned with business requirements. - Worked in an Agile environment, participating in sprint planning, daily stand-ups, and retrospectives to ensure timely delivery of features and enhancements. - Deployed and monitored solutions using Azure Monitor and Azure Log Analytics to ensure the reliability and performance of data workflows, integrating testing protocols with monitoring applications such as Dynatrace for proactive issue detection and resolution. - Collaborated with data scientists to integrate machine learning models into the data pipelines, utilizing Azure ML and Databricks for training and deployment to enhance predictive analytics capabilities across the business. - Used Azure DevOps for version control, CI/CD pipelines, and managing project backlogs. - Collaborated with cross-functional teams, including data scientists, product managers, and DevOps, to align data engineering efforts with organizational goals. Tech Stack: Python, SQL, dbt, Azure Data Factory, Snowflake, Power BI, Azure Synapse Analytics, Snowflake Data Sharing, Snowflake Time Travel, Zero-Copy Cloning, Data Partitioning, Clustering, Query Optimization, Schema Design, Jira, Confluence, Azure DevOps, Git, Row-Level Security, Role-Based Access Control (RBAC), Data Encryption Mechanisms.	Feb 23 - Present
Client – GSK Technologies, Hyderabad, IN	Apr 21 – Jul 22

Data Modeling Engineer

Responsibilities

- Built and optimized ETL pipelines using **Azure Data Factory** and Python, ensuring seamless integration of **healthcare data** from **EHR** systems, claims, and external data sources into a centralized **data lake**.
- Processed and ingested large-scale healthcare data into **Azure Synapse Analytics** and **Snowflake** for downstream analytics and reporting.
- Developed workflows to handle structured and semi-structured data (e.g., HL7, JSON) using **PySpark** and **Databricks**.
- Designed and implemented a dimensional data model (star schema) for the analytics layer, supporting healthcare reporting and **KPI** dashboards.
- Utilized **Magic Draw** to create **UML class diagrams** for effective **data modeling**.
- Created **normalized 3NF models** for operational data stores (ODS) to facilitate real-time data updates.
- Defined and developed master data management (MDM) strategies to ensure consistent and accurate patient and provider information.
- Leveraged **Azure Blob Storage** and **Azure Data Lake** for scalable data storage and processing.
- Explored NoSQL options using **Azure Cosmos DB** to manage semi-structured data for real-time queries.
- Utilized **Apache Spark** and **Hadoop** for distributed data processing and transformation.
- Automated data validation and transformation workflows using **dbt** for ensuring high data quality.
- Implemented logging and monitoring solutions using **Azure Monitor** to track data pipeline performance and resolve bottlenecks.
- Developed and published interactive dashboards in **Power BI**, providing actionable insights on patient care metrics, operational efficiency, and claims processing.
- Enabled self-service reporting by creating user-friendly data marts and pre-aggregated views in **Snowflake**.
- Processed HL7 v3 messages for seamless integration with EHR systems, ensuring compliance with healthcare interoperability standards.
- Integrated **Databricks ML** pipelines to preprocess healthcare data and engineer features for machine learning models.
- Deployed **Azure Machine Learning** models to operationalize predictions and improve decision-making processes in real time.
- Containerized data processing applications using **Docker** to ensure consistent development, testing, and deployment environments, enabling seamless integration with cloud platforms and improving scalability and efficiency across pipelines.
- Implemented **Delta Lake** for managing large-scale, high-quality data lakes, ensuring ACID transactions, time travel, and scalable data storage to support both batch and streaming analytics in the healthcare domain.
- Utilized **MLflow** and **Databricks** for model tracking, versioning, and managing machine learning lifecycle.
- Conducted data encryption and anonymization to comply with **HIPAA** and data privacy regulations.
- Implemented **Kafka** to enable real-time streaming of healthcare data, allowing timely analysis of events such as patient admissions, claims updates, and operational metrics.
- Used **Terraform** to automate the provisioning of cloud infrastructure, ensuring consistent and repeatable deployments of data pipelines and related resources across environments.
- Worked in an **Agile** Scrum environment, collaborating with data scientists, analysts, and business stakeholders to deliver incremental features.
- Participated in sprint planning, stand-ups, and retrospectives to ensure timely delivery of data engineering solutions.

Tech Stack:

Azure Data Factory, Python, Azure Synapse Analytics, Snowflake, PySpark, Databricks, Magic Draw, Apache Spark, Hadoop, Azure Blob Storage, Azure Data Lake, Azure Cosmos DB, dbt, Azure Monitor, Power BI, Kafka, Terraform, HL7 v3, MLflow, Azure Machine Learning.

Client— **Deccan Info Systems, Hyderabad, IN**

Feb 20 – Mar 21

Data Engineer

Responsibilities

- Built and maintained ETL pipelines using **Python**, **SQL**, and **GCP Dataflow** to process high-volume financial transactions and integrate them into a centralized **BigQuery** data warehouse.
- Automated data ingestion and transformation workflows using **GCP Data Fusion**, ensuring real-time processing of customer and transaction data.
- Integrated **Pub/Sub** for real-time data streaming from banking applications and third-party APIs.
- Designed and implemented a **star schema data model** in **BigQuery**, optimizing it for reporting and analytics.

- Developed dimensional models for customer segmentation, transaction categorization, and credit risk evaluation.
- Performed **query optimization** in **BigQuery** and **SQL**, reducing reporting latency by 25%.
- Utilized **Apache Beam** for distributed data processing in real-time and batch workflows.
- Leveraged **Dataproc** and **Hadoop** for managing and processing large-scale data.
- Implemented **ML pipelines** to predict customer churn and detect fraudulent transactions using **TensorFlow** and **Scikit-learn**.
- Deployed models in **Vertex AI** for operationalizing machine learning predictions in fraud detection and credit risk management.
- Used **MLflow** for tracking, versioning, and managing the lifecycle of machine learning models.
- Created and maintained **source-to-target mappings** and ensured **data lineage** for all transformations.
- Performed data validation and implemented quality checks using **Python scripts** to ensure accuracy and compliance with **data governance standards**.
- Collaborated with the compliance team to meet regulatory requirements, including **SOX** and **GDPR**.
- Built interactive dashboards using **Tableau** and **Google Data Studio** to monitor KPIs such as customer churn rate, fraud alerts, and credit scoring trends.
- Designed pre-aggregated views and **data marts** for business intelligence and self-service reporting.
- Configured Fivetran connectors to automate the ingestion of data from third-party APIs and SaaS platforms into **Snowflake** and **BigQuery**.
- Leveraged **GCP Cloud Storage** for scalable and secure data storage.
- Used **GCP IAM** for managing access controls and ensuring data security.
- Implemented **CI/CD pipelines** using **Jenkins** and **GitHub Actions** to streamline deployment of data pipelines and workflows.
- Orchestrated complex workflows and data pipelines using **Apache Airflow**, enabling efficient scheduling, monitoring, and maintenance of ETL processes.
- Automated infrastructure provisioning and configuration using **Terraform** to ensure consistency across environments.
- Collaborated with cross-functional teams, including data scientists and business analysts, to design and implement data-driven solutions for fraud detection and credit risk management.
- Followed **Agile methodologies**, participating in sprint planning, daily stand-ups, and retrospectives.

Tech Stack:

Python, SQL, GCP Dataflow, BigQuery, GCP Data Fusion, Pub/Sub, Apache Beam, Dataproc, Hadoop, TensorFlow, Scikit-learn, Vertex AI, MLflow, Tableau, Google Data Studio, GCP Cloud Storage, GCP IAM, Jenkins, GitHub Actions, Airflow, Terraform, Agile methodologies

License & Certifications

- **Snowflake - Snow core pro Certificate**
- **Azure Data Engineering**
- **Google Cloud Analytics**
- **IBM Business Intelligence (BI) Analyst Professional Certificate**
- **Power BI**
- **Data Bricks**
- **Machine Learning**
- **Gen AI**