


SENIOR AZURE DATA ENGINEER

Nikhil Reddy Burry

Mobile: 856-390-3706

Mail: nikhilreddyburry@gmail.com



PROFESSIONAL SUMMARY

- Accomplished Data Engineer with over 12 years of experience in designing and implementing scalable data ingestion pipelines using Big Data, Microsoft Azure Cloud, Python, PySpark, On-Prem to Azure Cloud Solutions.
- Designed and implemented Databricks Unity Catalog 12+ years of experience on Azure Data Engineer, Big Data with Hadoop, Hive, Map Reduce, Spark, Spark Core, Spark SQL, and Data Frames/Data Sets/RDD API
- Deployed and developed various data pipelines extensively and intricately harnessing the power of Azure Data Factory.
- Demonstrated expertise in leveraging Azure Databricks for distributed data processing, transformation, validation, cleansing, and ensuring data quality and integrity
- Event Hub for real-time streaming data ingestion
- Proficient in utilizing Azure Synapse Pipelines for orchestrating and managing complex data integration and transformation workflows
- Secured and managed sensitive cryptographic keys and secrets by leveraging Azure Key Vault, ensuring robust data protection and compliance in cloud-based environments.
- Developed end-to-end data workflows using Azure Logic Apps, Azure Functions, and serverless solutions.
- Demonstrated expertise in utilizing Azure Databricks to establish a centralized data governance framework, improving security, metadata management, and cross-workspace access control.
- Extensive experience in working with Azure Blob Storage for efficient storage and retrieval of unstructured and semi-structured data.
- Worked in complex SNOW SQL and Python Queries in Snowflake.
- Created DB, Schema, Virtual Warehouses, Tables, Views, Stages, Sequence, and DB replication in Snowflake Cloud Data Warehouse.
- Used Replication Task to have the data back up in Snowflake.
- Designed and implemented complex ETL workflows using Apache Airflow, ensuring seamless data pipeline orchestration across diverse data sources
- Developed custom scheduling strategies in Airflow to automate daily, weekly, and monthly data ingestion and transformation tasks, optimizing resource usage and ensuring timely data availability.
- Optimized Airflow DAGs for high performance and scalability, reducing execution time by 30% through efficient task parallelization and resource management.
- Worked on Snowflake advanced concepts like setting up virtual warehouse sizing, query performance tuning, Data Sharing, UDF, Zero Copy Clone, Time Travel, and Data Pipelines (Streams and Tasks).
- Followed Snowflake best practices like Clustering the View and Tables, created Materialized views, enabled the Result cache, Resizing and Multi Clustering to improve the performance in Snowflake.
- Designed dimensional model, data lake architecture, data vault 2.0 in Snowflake and used Snowflake logical data warehouse for compute.
- Engineered high-volume batch processing pipelines using Cosmos DB as a scalable data store, optimizing bulk insert and update operations to achieve throughput of 1 million records per minute.
- Implemented partitioning strategies and custom indexing policies in Cosmos DB to support efficient batch query operations, reducing processing time of nightly ETL jobs by 40%.
- Developed a fault-tolerant batch-processing framework integrating Azure Data Factory with Cosmos DB, incorporating check pointing and retry mechanisms to ensure data consistency during large-scale data migrations.
- Optimized Cosmos DB's bulk executor library for batch operations, resulting in a 60% reduction in RU consumption and enabling cost-effective processing of multi-terabyte datasets during off-peak hours.
- Highly skilled in utilizing Hadoop, HDFS, Map-Reduce, Hive, and SparkSQL and for efficient ETL tasks, real-time data processing, and analytics.
- Exceptional command over Kafka streaming technology, adeptly utilizing its distributed messaging capabilities to construct resilient and high-performing data flows.

- Designed and developed scalable ETL pipelines in Databricks (PySpark, SQL) and Snowflake, optimizing data ingestion, transformation, and storage
- Built end-to-end data pipelines using Delta Lake in Databricks, ensuring ACID compliance, time travel, and efficient data versioning.
- Developed real-time and batch data processing workflows using Databricks Delta Live Tables (DLT) and Snowflake Streams & Tasks for low-latency analytics
- Implemented orchestration frameworks using Apache Airflow, Databricks Workflows, and dbt for automated data pipeline execution and dependency management.
- Implemented Databricks Unity Catalog to enable centralized data governance, metadata management, and fine-grained access control across multiple workspaces
- Designed data lineage and auditing strategies in Unity Catalog to track data transformations and dependencies across ETL pipelines
- Designed and developed Data Mesh-based decentralized data platforms using Databricks Delta Lake and Snowflake, enabling domain-driven data ownership and governance
- Established data contracts and schema enforcement policies to maintain data consistency and interoperability across different teams and business units
- Designed and optimized Snowflake data warehouses, implementing clustering, materialized views, and partition pruning for improved query performance
- Developed and automated data ingestion pipelines into Snowflake using Snowpipe, COPY INTO, and Streams & Tasks for near real-time processing
- Configured RBAC and ABAC policies in Databricks Unity Catalog and Snowflake, enabling fine-grained access control at the catalog, schema, table, column, and row levels
- Automated governance rule enforcement using Databricks Workflows and Snowflake Stored Procedures, ensuring continuous compliance and monitoring
- Demonstrated ability to design and implement data integration strategies between Snowflake and external systems, leveraging technologies such as Apache Airflow or custom-built orchestration frameworks to ensure seamless data movement and synchronization

CERTIFICATIONS

Microsoft Certified:

- **AZ-900 MICROSOFT AZURE FUNDAMENTALS**
- **DP-203 MICROSOFT AZURE DATA ENGINEER ASSOCIATE.**
- **DATABRICKS DATA ENGINEER ASSOCIATE**

TECHNICAL SKILLS

Azure Services	Azure Analysis Services, Azure cosmos DB, Azure Data Factory, Azure Data Lake, Azure Devops, Azure Event Hub, Azure Key Vault ,Azure Log Analytics, Azure Logic Apps, Azure Monitor, Azure purview, Azure SQL ,Azure Synapse Analytics, Data bricks, Dimensional Data, Modeling, Microsoft Fabric ,Power BI.
Big Data Technologies	HDFS, MapReduce, Hive, Sqoop, Oozie, Zookeeper, Kafka, Apache Spark, Spark Streaming.
Operating Systems	Amazon Linux AMI, Linux (ubuntu, Centos, Red Hat, Debian), server, Jboss
Java / J2EE Technologies	Java, Servlets / JSP, Struts, JMSD, XML, BEA WebLogic, APACHE web server .
Virtualization Tools	VMware Workstation, Oracle VirtualBox, Vagrant
Version Control Tools	GIT, GitHub, SVN , Bitbucket, GitLab.
CI / CD Tools	Jenkins, GitHub Actions , Azure Devops.
Containerization	Docker, Kubernetes.
Web / App services	Nginx, Apache Tomcat, Jetty ,Apache HTTP server ,WebLogic
SDLC	Agile, Waterfall.
Bug Tracking Tools	JIRA

WORK EXPERIENCE

Client	Veolia
Role	Azure Data Engineer
Location	New Jersey
Project Duration	Sep2022- Mar2025

- Designed and implemented end-to-end data pipelines using Azure Data Factory to facilitate efficient data ingestion, transformation, and loading (ETL) from diverse data sources into Snowflake data warehouse.
- Implemented Azure Data Factory (ADF) extensively for ingesting data from different source systems like relational and unstructured data to meet business functional requirements
- Developed Pipelines in ADF using Linked Services/Datasets/Pipeline/ to Extract, Transform and load data from different sources like Azure SQL, Blob storage, Azure SQL Data warehouse, write-back tool and backwards
- Deployed Azure Data Lake Storage as a reliable and scalable data lake solution, implementing efficient data Partitioning and retention strategies to store and manage both raw and processed data effectively.
- Employed Azure Blob Storage for optimized data file storage and retrieval, implementing advanced techniques like compression and encryption to bolster data security and streamline storage costs.
- Integrated Azure Logic Apps seamlessly into the data workflows, ensuring comprehensive orchestration and triggering of complex data operations based on specific events, enhancing overall data pipeline efficiency.
- Have good experience working with Azure BLOB and Data Lake storage and loading data into Azure SQL Synapse analytics (DW).
- Implemented advanced data governance protocols using Fabric's Purview integration, establishing automated data lineage tracking and sensitive data classification across the entire data estate.
- Optimized query performance in Fabric's Synapse Data Warehouse by leveraging materialized views, result-set caching, and workload management, resulting in a 60% reduction in average query execution time for complex analytical workloads.
- Designed a scalable, multi-tenant data mesh using Fabric's OneLake and Dataverse, implementing row-level security and dynamic data masking to ensure data sovereignty while enabling cross-organizational analytics.
- Designed and developed metadata control tables for Data factory Pipelines to pass the parameters to the pipelines dynamically using pipeline parameters.
- Troubleshoot and resolved data processing issues promptly, minimizing downtime and ensuring uninterrupted data availability for analysis.
- Managed codebase and deployments using Git, facilitating seamless collaboration with team members, and ensuring version control integrity.
- Utilized SQL to normalize data structures and schemas, optimizing data storage and improving query performance for analytical workloads, while ensuring consistency and adherence to data governance standards.
- Developed and deployed Azure Functions to handle critical data preprocessing, enrichment, and validation tasks within the data pipelines, elevating the overall data quality and reliability.
- Design and implement end-to-end data solutions (storage, integration, processing, visualization) in Azure
- Design and implement database solutions in Azure SQL Data Warehouse, Azure SQL
- implemented medium to large scale solutions on Azure using Azure Data Platform services (Azure Data Lake, Data Factory, Data Lake Analytics, Stream Analytics, Azure SQL DW, HDInsight/Databricks, NoSQL DB)
- Design & implement migration strategies for traditional systems on Azure (Lift and shift/Azure Migrate, other third-party tools
- Implemented comprehensive monitoring and alerting for Airflow tasks using built-in logging and external tools like Prometheus and Grafana, ensuring proactive issue detection and resolution.
- Orchestrated complex Spark jobs within Airflow, leveraging its capabilities to handle large-scale data processing tasks efficiently, including both batch and real-time processing.
- Developed custom Airflow operators and plugins to extend functionality, enabling seamless integration with proprietary tools and third-party APIs.
- Implemented data quality checks and validation steps within Airflow pipelines to ensure data accuracy and consistency

- Ensured data security and compliance within Airflow workflows by implementing role-based access control (RBAC), encryption, and adherence to industry standards and best practices.
- Engage with business users to gather requirements, design visualizations, and provide training to use self-service BI tools.
- Develop conceptual solutions & create proof-of-concepts to demonstrate viability of solutions
- Technically guide projects through to completion within target timeframes
- Collaborate with application architects and DevOps
- Identify and implement best practices, tools and standards Architect and implement ETL and data movement solutions using Azure Data Factory, SSIS create and run SSIS Package ADF V2 Azure-SSIS IR
- Design and implement database solutions in Azure SQL Data Warehouse, Azure SQL Recreating existing application logic and functionality in the Azure Data Lake, Data Factory, SQL Database and SQL Datawarehouse environment. experience in DWH/BI project implementation using Azure DF
- Migrate data from traditional database systems to Azure databases
- Designed and deployed architecture for Azure Service fabric.
- Developed Orchestrated data pipelines using Microsoft Fabric's Data Factory, integrating OneLake storage and Synapse Data Engineering to streamline ETL processes across diverse data sources.
- Implemented Fabric's Synapse Data Warehouse and Data Science workspaces to develop and deploy real-time analytics dashboards, resulting in a 40% improvement in data-driven decision-making efficiency.
- Designed and developed Fabric's Power BI integration to create interactive reports and KPI monitoring tools, enhancing data storytelling capabilities and increasing stakeholder engagement by 25%
- Designs Logical and Physical Data Model for Staging, DWH and Data Mart layer. Recommend, Design and construct policies and standards that affect infrastructure operations and services but also improve overall business performance. (Strategic)
- Provide leadership and work guidance to less experienced personnel (Leadership). Work with similar Microsoft on-prem data platforms, specifically SQL Server and related technologies such as SSIS, SSRS, and SSAS
- Experience managing Azure Data Lakes (ADLS) and Data Lake Analytics and an understanding of how to integrate with other Azure Services.
- Identify potential problems and recommend alternative technical solutions
- Collaborate with application architects on infrastructure as a service (IaaS) application to Platform as a Service (PaaS)
- Design and implementation of High Availability (HA) systems (24x7x365) an Expertise in setup High Availability and Recoverability of databases using with SQL Server technologies including Always On Azure VM
- Architect, design and validate Azure infrastructure-as-a-Service (IaaS) environment
- Cloud Security Architecture, Governance, Best practices, and tools
- Azure Cloud Build and Automation using ARM template Deploying Azure Resource Manager JSON Templates from PowerShell worked on Azure suite: Azure SQL Database, Azure Data Lake, Azure Data Factory, Azure SQL Data Warehouse, Azure Analysis Service
- Worked on SQL Server Integration Services (SSIS), SSAS, SSRS,T-SQL skills, stored procedures, triggers SQL Server Analysis Services (SSAS) and SQL Server Reporting Services (SSRS) and Familiarity with Excel PowerPivot, PowerView, and PowerBI
- Collaborated closely with cross-functional teams including data scientists, data analysts, and business stakeholders, ensuring alignment with data requirements and delivering scalable and reliable data solutions.

Environment: Azure Data Factory, Azure Databricks, Snowflake data warehouse, Azure Event Hubs, Azure Functions, Azure Data Lake Storage, Azure Blob Storage, Azure Logic Apps, Azure Machine Learning, Microsoft Fabric, Power BI, Azure Analysis Services, Purview, Airflow, Snowflake, Unity Catalog

Client	Velocity clinical research
Role	Azure Data Engineer
Location	Remote
Project Duration	july2020 -Aug2022

- Experience in designing and developing POCs in Spark using Scala to compare the performance of Spark with MapReduce, Hive
- Hands - on experience in Azure Cloud Services, Azure Synapse Analytics, SQL Azure, Data Factory, Azure Analysis services, Application Insights, Azure Monitoring, Key Vault, and Azure Data Lake.
- Created Batch & Streaming Pipelines in Azure Data Factory (ADF) using Linked Services/Datasets/Pipeline/ to Extract, Transform and load data.
- Created and managed various types of Snowflake tables, including transient, temporary, and persistent tables, to cater to specific data storage and processing needs.
- Implemented advanced partitioning techniques in Snowflake to significantly enhance query performance and expedite data retrieval.
- Defined robust roles and access privileges within Snowflake to enforce strict data security and governance protocols.
- Implemented regular expressions in Snowflake for seamless pattern matching and data extraction tasks.
- Developed and implemented Snowflake scripting solutions to automate critical data pipelines, ETL processes, and data transformations,
- Developed event-driven data mesh to combines the scale and performance of data in motion with product-focused rigor and self-service capabilities, putting data at the front and center of both operational and analytical use-cases.
- Developed event-driven data mesh which reduces barriers between operational and analytical use-cases, while simultaneously enabling both real-time applications and batch-based jobs
- Created Azure Data Factory (ADF) Batch pipelines to Ingest data from relational sources into Azure Data Lake Storage (ADLS gen2) & incremental fashion and then load into Delta tables after cleansing
- Created Azure logic apps to trigger when a new email received with an attachment and load the file to blob storage
- Implemented CI/CD pipelines using Azure DevOps in cloud with GIT, Maven, along with Jenkins plugins.
- Build a Spark Streaming application to perform real-time analytics on streaming data.
- Worked on snowflake connector for developing python applications.
- Created roles and access level privileges and taken care of Snowflake Admin Activity end to end.
- Implemented Snowflake Continuous Data Loading, Snowflake Time Travel & Fail-safe and secure Data Sharing.
- Created conceptual, logical, and physical models for OLTP, Data Warehouse Data Vault and Data Mart Star/Snowflake schema implementations.
- Build the logical and physical data model for Snowflake as per the changes required.
- Loaded the tables from the azure data lake to azure blob storage for pushing them to snowflake.
- Develop a Spark Streaming application that integrates with event-driven architectures such as Azure Functions or Azure Logic Apps.
- Automated critical data workflows and monitoring tasks using Python scripting, saving over 20 hours of manual effort per week.
- Use Spark Streaming to process events in real-time, and trigger downstream workflows based on the results.
- Involved in creating Hive tables and loading and analyzing data using hive queries
- Designed and developed custom HiveUDF's
- Using the JSON and XMLSerDe's for serialization and de-serialization to load JSON and XML data into HIVE tables
- Involved in migration of ETL processes from Oracle to Hive to test the easy data manipulation.
- Implemented to reprocess the failure messages in Kafka using offset id
- Used HiveQL to analyze the partitioned and bucketed data.
- Worked on Infomatica Data Quality to resolve customers address related issues.
- Designed a STAR schema for the detailed data marts and Plan data marts involving shared dimensions (Conformed).
- Ensured the feasibility of the logical and physical design models.
- Developed a Spark job in Java which indexes data into azure functions from external Hive tables which are in HDFS.
- Written Hive queries on the analyzed data for aggregation and reporting

- Developed Sqoop Jobs to load data from RDBMS to external systems like HDFS and HIVE
- Developed Spark applications using PySpark and Spark-SQL for data extraction, transformation, and aggregation from multiple file formats.
- Worked on converting the dynamic XML data for injection into HDFS
- Transformed and Copied data from the JSON files stored in a Data Lake Storage into an Azure Synapse Analytics table by using Azure Databricks
- Azure Databricks, Azure Storage Account etc. for source stream extraction, cleansing, consumption and publishing across multiple user bases.
- Created resources, using Azure Terraform modules, and automated infrastructure management.
- Implemented Spark Scripts using Scala, Spark SQL to access hive tables into spark for faster processing of data
- Loading data from UNIX file system to HDFS
- Configured spark streaming to receive real time data from the Apache Flume and store the stream data using Scala to Azure Tables.
- Loaded the data into Spark RDD and do in memory data Computation to generate the Output response
- Used several RDD transformation to filter the data injected into SparkSQL
- Used HiveContext and SQLContext to integrate Hive meta store and SparkSQL for optimum performance.
- Used the version control system GIT to access the repositories and used in coordinating with CI tools

Environment Azure Databricks, Data Factory, Snowflake Data Warehouse, Logic Apps, Functional App, Snowflake, SparkSQL, Oracle, HDFS, MapReduce, YARN, Spark, Hive, SQL, Python, Scala, PySpark, Jenkins, Kafka, Power Bi, Spark Streaming.

Client	Wells Fargo
Role	Big Data Developer
Location	California
Project Duration	April 2018- jun 2020

- Imported Data using Sqoop to load Data from MySQL to HDFS on regular basis.
- Worked with Data Lakes and big data ecosystems (Hadoop, Spark, Hortonworks, Cloudera).
- Load and transform large sets of structured, semi structured, and unstructured data.
- Written Hive queries for data analysis to meet the Business requirements.
- Built HBASE tables by leveraging on HBASE Integration with HIVE on the Analytics Zone.
- Hands on experience in using Kafka, Spark streaming, to process the streaming data in specific use cases.
- Developed data pipeline using Flume, Sqoop to ingest customer behavioral data histories into HDFS for analysis.
- Worked on analyzing Hadoop cluster using different big data analytic tools including Hive, and MapReduce.
- Developed a data pipeline using Kafka, Spark, and Hive to ingest, transform and analyzing data.
- Wrote Hive queries for data analysis to meet the specified business requirements by creating Hive tables and working on them using Hive QL to simulate MapReduce functionalities.
- Written Hive queries for data analysis to meet the Business requirements.
- Built HBASE tables by leveraging on HBASE Integration with HIVE on the Analytics Zone.
- Hands on experience in using Kafka, Spark streaming, to process the streaming data in specific use cases.
- Developed data pipeline using Flume, Sqoop to ingest customer behavioral data histories into HDFS for analysis.
- Worked on analyzing Hadoop cluster using different big data analytic tools including Hive, and MapReduce.
- Developed a data pipeline using Kafka, Spark, and Hive to ingest, transform and analyzing data
- Implemented UNIX scripts to define the use case workflow and to process the data files and automate the jobs.
- Implemented automation for deployments by using YAML scripts for massive builds and releases
- Migrated the existing data to Hadoop from RDBMS (Oracle) using Sqoop for processing the data.
- Implemented CICD pipelines to build and deploy the projects in Hadoop environment.

- Using JIRA to manage the issues/project workflow.
- Worked on Spark using Python (PySpark) and Spark SQL for faster testing and processing of data.
- Used Spark Streaming to divide streaming data into batches as an input to Spark engine for batch processing.
- Used Zookeeper to coordinate, synchronize and serialize the servers within the clusters.
- Worked on Oozie workflow engine for job scheduling.
- Used Git as version control tools to maintain the code repository.
- Worked on SparkSQL using PySpark for analyzing and processing the data.
- Worked closely with the team in fixing the JVM related issues.

Environment: : SQoop, MYSQL, HDFS, Apache Spark Scala, Hive Hadoop, Cloudera, HBASE, Kafka, MapReduce, Zookeeper, Oozie, RDBMS, Python, PySpark, shell script, Ambari, JIRA.

Client	Southwest Airlines
Role	Data warehouse Developer
Location	Dallas ,Texas
Project Duration	July 2015-Mar 2018

- Conducted business requirements gathering, designed data flows, performed data quality analysis, and collaborated with the data warehouse architect to develop logical models, driving efficient data management.
- Created Erwin-based logical and physical data models, incorporating Agile methodology, in collaboration with technical teams and application users for efficient data modeling.
- Designed and developed SSIS packages for streamlined data movement from staging to EDW environments, enhancing data integration efficiency within the organization.
- Implemented normalization rules in database development and maintained referential integrity using primary and foreign keys.
- Produced interactive Drill Through and Drill Down reports in Power BI, incorporating drop-down menu options, data sorting capabilities, and subtotal definitions, enabling insightful data exploration and analysis.
- Utilizing Performance Point Server to create interactive dashboards and scorecards, facilitating data visualization and performance monitoring in data warehouse environments.
- Designed high level ETL architecture for overall data transfer from the OLTP to EDW with the help of SSIS.
- Demonstrating proficiency, possess a comprehensive grasp of the features, structure, attributes, hierarchies, and schema designs such as Star and Snowflake Schemas in Data Marts.
- Utilized Data Warehousing to develop a Data Mart for generating downstream reports and created a User Access Tool to enable ad-hoc reporting and data analysis through querying the proposed Cube, promoting self-service analytics.
- Designed, Developed and Deployed servers in MS SQL Server environment using the SSIS Packages and created jobs for efficient running of the packages.
- Implemented diverse data transformations, including Slowly Changing Dimension, Aggregate, Sort, Multicasting, Conditional Split, and Derived Column, using SSIS, enhancing data integration and manipulation capabilities.
- Developed and maintained a data store for reporting in SSRS, utilizing Global Variables, Expressions, and Functions to generate comprehensive reports for informed decision-making.
- Proficient in utilizing SSAS for designing and developing analytical models, enabling multidimensional data analysis, and reporting in data warehouse environments.
- Created reports using SSRS from OLTP and OLAP data sources and deployed on report server.
- Leveraging PL/SQL code for Data Extraction, Transformation and Loading and also for data Control flow and Exception handling.
- Proficient in utilizing SQL Profiler for performance tuning and query analysis, ensuring optimized query execution and enhanced database performance.
- Experienced in Building Cubes and Dimensions with different Architectures and Data Sources for Business Intelligence and writing MDX Scripting.

- Engaged in the creation of SSIS jobs to automate report generation and cube refresh packages, streamlining data processing and enhancing reporting capabilities.
- Proficient in NoSQL databases like MongoDB and Cassandra for efficient handling and analysis of large-scale unstructured and semi-structured data in data warehouse environments.
- Proficient in SharePoint, leveraging expertise in managing and collaborating on documents, workflows, and team sites to enhance productivity and streamline business processes.
- Developed stored procedures and triggers to facilitate consistent data entry into the database.
- Shared data outside using Snowflake to quickly set up to share data without transferring or developing pipelines.

Environment: Windows server, MS SQL Server 2014, SSIS, SSAS, SSRS, SQL Profiler, Performance Point Server, MS Office, SharePoint

Client	Tenet Health Care
Role	ETL & DataWare House Developer
Location	Dallas ,Texas
Project Duration	Jan 2013-Jun 2015

Cloudera HDP was used to develop and maintain data warehouse solutions, as well as to access its extensive range of Hadoop ecosystem tools.

- ETL/ELT processes were designed and put into place to provide smooth data loading, extraction, and transformation using Informatica Power Center.
- Utilized Hive for querying and analyzing large datasets stored in Hadoop Distributed File System (HDFS), using YARN resource management to optimize performance.
- Developed Unix scripting solutions to automate data processing tasks and enhance system efficiency.
- Conducted data integration between Oracle, Sybase, and Microsoft SQL Server databases for comprehensive data consolidation.
- Proficient in creating and implementing dimensional models for data warehousing solutions, ensuring maximum performance and flexibility.
- Implemented Apache Oozie workflow engine for managing and scheduling Hadoop jobs, ensuring efficient workflow execution.
- Leveraged SQOOP for importing and exporting data between HDFS and relational database systems, ensuring data consistency and integrity.
- Utilized Selenium for web automation testing to ensure the accuracy of data extraction and integration processes.
- Configured Maven and Eclipse for streamlined project management and development of ETL workflows.
- Employed TOAD for Oracle and Microsoft SQL Server for database management and development tasks.
- Implemented REST APIs for seamless integration between different data sources and systems.
- Utilized Unix scripting for automation and customization of data processing workflows.
- Conducted performance tuning and optimization of ETL/ELT processes to enhance overall system efficiency.
- Developed and maintained data models using Oracle SQL for efficient data storage and retrieval.
- Worked with stakeholders to gather requirements, define KPIs, and deliver reports and dashboards using tools like Microsoft SQL Server Reporting Services (SSRS) and Tableau.
- Participated in Agile development methodologies including sprint planning, daily stand-ups, and retrospectives.
- Set up and manage Linux environments for hosting Hadoop clusters and other data processing systems.

Environment: Cloudera HDP, HDFS, Linux, Hive, YARN, Apache Oozie, SQOOP, Oracle, Informatica Power Center, Sybase, Unix Scripting, Selenium, Maven, Eclipse, TOAD, Microsoft SQL Server, Oracle SQL, REST APIs, Tableau, SSRS.