

Vallapu P
Senior Data Engineer
Mobile: 607-218-2097

Email: praveenvallapu66@gmail.com



PROFESSIONAL SUMMARY

- Certified Azure Data Engineer (DP-203) and Databricks Data Engineer Associate with 12+ years of experience designing and implementing scalable cloud-based data solutions using Azure, Databricks, Snowflake, and Spark..
- Accumulating 10+ years of focused knowledge in Azure Data services like Azure Data Factory, Azure Data Lake storage, Blob Storage, Delta tables, Key Vault, Azure Databricks, Azure Synapse Analytics, Polybase, Event Hub, Function Apps, Logic Apps, Data Flow, Power Query, and Cosmos DB.
- 6 years of focused knowledge of big data technologies, Spark, Kafka, HBase, Scala, Oozie, and Zookeeper as well as working with Spark SQL, and the Core Spark API to investigate Spark capabilities and create data pipelines.
- Professional expert with 4 years of data warehouse with, a dedicated focus on designing, implementing, and optimizing solutions. Developed advanced SQL queries, ETL development, and cloud-based solutions, actively driving data integration and migration projects.
- Developed and deployed various ETL pipelines extensively and thoroughly with the help of Azure Data Factory.
- By strong knowledge of Azure Databricks to scalable and effective ETL data pipelines for data processing and transformation projects within Azure Databricks by using PySpark and SQL to create solid data workflows.
- Hands-on experience with Azure Key Vault to secure and manage private encryption keys and secrets ensuring strong data protection in cloud-based systems.
- Developed data workflows using Azure Logic Apps, Azure Functions, and serverless solutions.
- Working with Azure Event Hub for ingesting streaming data in real time.
- Expert in organizing developed intricate workflows for data integration and transformation using Azure Synapse Pipelines.
- A significant amount of expertise in developing Structured, Unstructured, and semi-structured data with Azure Blob Storage.
- Expertise in Data Engineering & Cloud Technologies – Proficient in designing and implementing scalable ETL pipelines, data processing workflows, and cloud-based data solutions using Azure Data Factory, Databricks, and Snowflake to support enterprise data initiatives.
- Strong Experience with Big Data & Distributed Processing – Skilled in leveraging Apache Spark (PySpark & Scala) for large-scale data transformation, real-time analytics, and performance optimization in cloud environments, ensuring efficient data processing and high availability.
- Hands-on experience in designing SSIS packages for handling SQL Server databases and objects between SQL Server instances.
- Proficiency in programming languages Python, Scala, PySpark, SQL, and PLSQL which help design scalable and efficient data solutions.
- Adept at designing cutting-edge, cloud-based data warehousing solutions using Snowflake on Azure, optimizing schemas, tables, and views for streamlined data storage and retrieval.
- Implemented real-time streaming applications using Apache Kafka on distributed Hadoop Clusters.

- Hands-on experience on Azure Logic Apps with a focus on triggers, including event-based triggers for real-time feedback and scheduled triggers for time-based automation workflows through task execution
- Hands-on experience dealing with a variety of file formats, such as CSV, JSON, Parquet, Binary, and ORC.
- Developed in leveraging Big Data technologies, including Hive, for efficient data querying and analysis.
- In-depth knowledge of HDFS for reliable and scalable storage and retrieval of large datasets.
- Adept at working with the Hadoop ecosystem, including Map Reduce, Yarn, and Sqoop, for seamless integration and processing of diverse data sources.
- Implemented performance tuning, by optimizing data processing performance, and increasing system efficiency for both OLAP and OLTP environments.
- Demonstrated expertise in implementing advanced data organization techniques, with effective Partitioning and Repartitioning strategies.
- In orchestrating data workflows using industry-standard scheduling tools such as Control-M and Apache Airflow. Implemented effective scheduling mechanisms to ensure timely execution of data pipelines, enhancing overall workflow efficiency.
- Expert in version control systems including Git and GitHub, ensuring co-development and effective code management.
- Software development and deployment procedures were streamlined by utilizing Azure DevOps, and automated CI/CD pipelines for efficient software delivery.
- Strong adherence to Agile methodologies, using collaborative and flexible development methodologies.
- Performance experience in cross-functional teamwork, iterative development, and delivery of high-quality solutions aligned with business objectives.

CERTIFICATIONS

Microsoft Certified: Azure Data Engineer Associate

Databricks Certified Data Engineer Associate

TECHNICAL SKILLS

Big Data Technologies	MapReduce, Hive, Python, PySpark, Scala, Kafka, Spark streaming, Oozie, Sqoop, Zookeeper
Hadoop Distribution	HDFS, SQL, YARN, PIG Latin, MapReduce, Hive, Sqoop, Spark, Yarn, Storm, Zookeeper, Oozie, Kafka, Storm, Flume
Azure Services	Azure data Factory, Airflow, Azure Data Bricks, Logic Apps, Functional App, Snowflake, Azure DevOps
AWS Services	EC2, S3, EMR, Redshift, DynamoDB, Aurora, VPS, Glue, Kinesis, Boto3
Languages	Java, SQL, PL/SQL, Python, HQL, PySpark.
Operating Systems	Windows (XP/7/8/10), UNIX, LINUX, UBUNTU, CENTOS.
Build Automation tools	Ant, Maven
Version Control	GIT, GitHub.
IDE & Build Tools, Design	Eclipse, Visual Studio.
Databases	MS SQL Server 2016/2014/2012, Azure SQL DB, Azure Synapse. MS Excel, MS Access, Oracle 11g/12c, Cosmos DB

PROFESSIONAL EXPERIENCE:

Client: Lumen Technologies

Sep 2022 to till Date

Location: Remote

Role: Azure Data Engineer

Job Description:

As an Azure Data Engineer at the Lumen, I played a key role in pioneering the migration of legacy data systems to modern Azure data platforms Utilizing my expertise in Azure data processing and ETL processing, improved data integrity, security, Led the design, development, and implementation of scalable data pipelines to ensure seamless data transfer while maintaining compliance.

Responsibilities:

- Worked as **Azure Data Engineer**, responsible for building and deploying scalable data ingestion pipelines using **Azure Data Factory**.
- Developed high-performance **ETL pipelines** using **Scala** within **Azure Databricks**, optimizing data ingestion and transformation processes.
- Utilized **Spark with Scala** to perform large-scale **data processing, transformation, and aggregation**, improving query performance for analytics.
- Implemented **parallel processing techniques** in **Scala** to enhance the efficiency of real-time data ingestion and ETL workflows.
- Wrote **custom Scala functions** for data validation and quality checks, ensuring data consistency and integrity before loading into **Snowflake**.
- Integrated **Scala-based Spark applications** with **Azure Event Hubs and Azure Data Lake**, enabling efficient real-time data streaming.
- Access and integrate data from multiple sources, including **SQL databases, CSV files, and REST APIs** to support the organization's data-driven initiatives
- Designed and optimized data processing workflows using **Azure Databricks**, leveraging **PySpark** and **Spark SQL** for highly distributed data processing and transformation workflows.
- Building a robust **ETL pipeline** with Linked services in **Azure Data Factory** ensures easy extraction and transformation insertion of data.
- Ensure data quality and integrity through data manipulation, cleansing, and transformation using **Azure Data Factory** and **Databricks**.
- Design and optimize **data warehouse systems**, including **Azure SQL Data Warehouse** structures, tables, and instances for efficient data storage, retrieval and analysis
- Work closely with **data analysts** and **business stakeholders** to understand their needs and implement appropriate data models and processes within **Snowflake**.
- Implemented and optimized **Spark** tasks for **data transformation**, integration, and machine learning tasks on large data sets.
- Designed and implemented data masking solutions to anonymize sensitive data, ensuring compliance with privacy regulations.
- Utilized data masking tools such as Informatica Data Masking, IBM InfoSphere Optim, and Oracle Data Masking to protect sensitive information.
- Hands-on experience with **Snowflake** to combine large volumes of data processing and analytics capabilities for simple and efficient **data analysis**.
- Configure **event-based trigger** mechanisms through Azure Logic Apps to optimize data pipelines and business processes.

- **Metadata management** solution using **Azure services** to manage, and analyze data flows and changes.
- Identify and resolve performance **bottlenecks** in data processing and storage by optimizing query performance and reducing **data storage**.
- Development and use of advanced **Tableau** and **Power BI visualizations** to transform raw data into insightful charts and visualizations for informed decision-making.
- Developed interactive and insightful data visualizations using Tableau, enabling stakeholders to make data-driven decisions.
- Actively participating in **ETL projects**, ensuring data integrity, and maintaining pipeline stability.
- Expertise in data processing, integration, and management projects using **shell scripting** for responsive and flexible data collaboration.
- Hands-on experience with **Git** for version control and collaborative development, ensuring effective code management, collaboration, and change tracking.

Environment Skills: Azure Data Factory, ETL, Azure Databricks, Azure Event Hubs, Azure Functions, Azure Data Lake Storage, Azure Key Vault, Azure Blob Storage, Azure SQL, Azure Logic Apps, Azure DevOps, Azure Monitor, Event grids, Power BI, Snowflake Features (zero-per-cloning, time travel, shared data management).

Client: Goldmansachs

July 2020 - Aug 2022

Location: Houston, Texas

Role: Azure Data Engineer

Job Description:

As an Azure Data Engineer at Goldmansachs , I played a vital role in leveraging cloud-based technologies to design, develop, and implement scalable and efficient data solutions to support various state government initiatives and projects. Working closely with stakeholders across different departments, I contributed to the development of robust data pipelines, ensuring the integrity, reliability, and accessibility of data for critical decision-making processes.

Responsibilities:

- Built and implemented scalable data ingestion pipelines through Azure Data Factory, accessing data from various sources such as SQL databases, CSV files, and REST APIs.
- Developed scripts for data transformations using Azure Databricks-PySpark and Spark SQL for distributed parallel processing.
- Designed and implemented distributed data processing frameworks using Apache Spark with Scala, reducing data transformation times by 40%.
- Developed Scala-based transformations in Databricks to process structured and semi-structured data from various sources.
- Used Scala-based data masking techniques, including encryption and tokenization, to protect sensitive financial data in compliance with regulatory standards.
- Optimized Spark SQL queries written in Scala to improve performance and scalability in Azure Synapse Analytics.
- Built Kafka-based real-time data pipelines with Scala and Spark Streaming, enabling near real-time analytics and business intelligence insights.
- Applied various data masking techniques including substitution, shuffling, encryption, and tokenization to secure sensitive data.

- The pipeline was connected using Linked services in Azure Data Factory to establish a connection between the source and target.
- Designed a real-time data processing pipeline using Azure Stream Analytics, Azure Event Hubs, and Apache Airflow, enabling real-time insights and decision-making for the business.
- Managed and orchestrated large-scale data workflows using Apache Airflow, ensuring timely and accurate data processing.
- Azure Data Factory and Databricks were used to ensure data integrity, purity, and transformation to ensure data quality and integrity.
- Implemented a cloud-based data warehousing solution using Azure Synapse Analytics to take advantage of its scalability and efficiency.
- Azure SQL Databases to store relational structured data, including schemas, tables, and views, and provide efficient data storage, retrieval, and analysis for analytics and reporting purposes
- Implemented appropriate data models and processes within Azure Synapse Analytics and collaborated with data analysts and business stakeholders to understand their needs.
- Optimized Spark applications for data transformation, aggregation, and machine learning tasks on large datasets.
- Azure Synapse Analytics was used to combine huge data processing and analytics capabilities for complex data analysis.
- Configured event-based trigger mechanisms using Azure Data Factory for a more efficient data pipeline and business processes.
- Adept at Azure Synapse Analytics' unique features such as zero-copy cloning, time travel, and shared data management.
- Azure Synapse Analytics and Azure Data Factory were combined to build a robust ETL pipeline, simplifying the migration of data from source to Azure-based data warehouses.
- Collaborated on ETL projects, ensuring data integrity and pipeline stability.
- A data pipeline using Kafka was developed and deployed using Zookeeper, Spark, and Hive for data optimization, management, and analysis.
- Expert in optimizing data processing, integration, and management projects, using shell scripting for responsive and flexible data collaboration.
- Git has been used for version control and collaborative development, facilitating efficient code management, collaboration, and change tracking.

Environment Skills: Azure Data Factory, Databricks, Snowflake, Azure DevOps, PySpark, Spark SQL, Azure Logic Apps, Azure Synapse Analytics, Tableau, Power BI visualizations, shell scripting, Power BI, Kafka, Zookeeper, Git.

Client: Citizens Bank

April 2018 to June 2020

Location: Dallas, Texas

Role: Data Engineer

Job Description:

As a Data Engineer working with a banking client Republic Financial Bank, I will be responsible for designing, developing, and maintaining scalable data pipelines and analytic solutions to support business processes using your expertise in big data technology and programming language, plays a key role in transforming raw data into actionable insights, and in developing data-driven decision-making processes within an organization.

Responsibilities:

- Utilized Apache Spark with Python to develop and execute Big Data Analytics and Machine learning applications, executed machine Learning use cases under Spark ML and Mllib.
- Identified areas of improvement in existing business by unearthing insights by analyzing vast amount of data using machine learning techniques.
- Interpret problems and provides solutions to business problems using data analysis, data mining, optimization tools, and machine learning techniques and statistics.
- Designed and developed NLP models for sentiment analysis.
- Led discussions with users to gather business processes requirements and data requirements to develop a variety of Conceptual, Logical and Physical Data Models.
- Expert in Business Intelligence and Data Visualization tools: Tableau, MicroStrategy.
- Worked on machine learning on large size data using Spark and MapReduce.
- Let the implementation of new statistical algorithms and operators on Hadoop and SQL platforms and utilized optimizations techniques, linear regressions, K-means clustering, Native Bayes and other approaches.
- Developed Spark/Scala, Python for regular expression (regex) project in the Hadoop/Hive environment with Linux/Windows for big data resources.
- Data sources are extracted, transformed, and loaded to generate CSV data files with Python programming and SQL queries.
- Stored and retrieved data from data-warehouses using Amazon Redshift.
- Refined time-series data and validated mathematical models using analytical tools like R and SPSS to reduce forecasting errors.
- Worked on data pre-processing and cleaning the data to perform feature engineering and performed data imputation techniques for the missing values in the dataset using Python.
- Created Data Quality Scripts using SQL and Hive to validate successful data load and quality of the data.
- Created various types of data visualizations using Python and Tableau.
- Consult on broad areas including data science, spatial econometrics, machine learning, information technology and systems and economic policy with R.
- Performed Data mapping between source systems to Target systems, logical data modeling, created class diagrams and ER diagrams and used SQL queries to filter data.
- Enabled speedy reviews and first mover advantages by using Oozie to automate data loading into the Hadoop Distributed File System and PIG to pre-process the data.
- Used various techniques using R data structures to get the data in right format to be analyzed which is later used by other internal applications to calculate the thresholds.
- Maintaining conceptual, logical and physical data models along with corresponding metadata.
- Worked on data migration from an RDBMS to a NoSQL database and gives the whole picture for data deployed in various data systems.

Environment Skills: Python, Hadoop, Map Reduce, Spark, Spark Mllib, Tableau, SQL, Excel, VBA, SAS, MATLAB, AWS, SPSS, Cassandra, Oracle, MongoDB, SQL Server 2012, DB2, T-SQL, PL/SQL, XML, Tableau.

Client: Everlane

July 2015 to March 2018

Role: ETL Developer

Job Description:

I was a key player in the design, development, and upkeep of data warehouse systems at Everlane , helping the company fulfil its reporting and data analytics needs. My duties included every stage of the data warehousing lifecycle, including requirements analysis, design, development, deployment, and continuing maintenance.

Responsibilities:

- ETL data flows were Developed and deployed with **SSIS**, facilitating seamless data migration and transformation from sources such as **SQL Server, Access, and Excel**
- Optimizing **SQL Server** performance through the creation of stored procedures, **triggers**, and functions, as well as the implementation of indexing and tracking techniques.
- Dimensional **data modelling** knowledge was used to design and develop **DataMart** including the presentation of facts and concepts
- Monitored **ETL** processes, quickly detecting and resolving errors and incidents using techniques such as priority limits, breakpoints, checkpoints, and logging
- Hands-on experience in creating **SSAS cubes**, defining aggregation, and KPIs, and helping implement data mining models for **SSAS** products.
- Developed various reports such as parameterized, charts, graphs, linked, dashboards, scorecards, and drill-down/drill-through reports in **SSAS** cubes using **SSRS**
- Implemented **ETL** data flows using **SSIS** that facilitates data migration and transformation from sources such as **SQL Server, Access, and Excel**.
- Collaborated with cross-functional teams to gather requirements, design data solutions, and deliver actionable insights through SQL queries and Alteryx workflows.
- Gained experience in dimensional **data modeling** and specified facts and dimensions for DataMart design, as well as the development of fact tables and dimension tables using slowly changing dimension (SCD) techniques
- Monitored errors and incidents in **ETL processes**, using techniques such as priority constraints, breakpoints, checkpoints, and logging.
- Gained experience developing **SSAS** cubes, implementing aggregation, defining KPIs, segmenting cubes, and developing, supporting the implementation of, and data mining models for **SSAS** products.
- Developed reports such as parameterized, charts, graphs, links, dashboards, scorecards, and drill-down/drill-through reports on **SSAS** cubes using **SSRS**.

Environment Skills: MS SQL Server, Visual Studio Legacy Versions, SSIS, SharePoint, MS Access, Git, SQL Server 2008/2012 Enterprise, SSRS, SSIS, T-SQL, Windows Server 2003, Performance Points Server 2007, Oracle 10g, Visual Studio 2010.

Client: CVS Health,

Location: Dallas, Texas

Jan 2013 to Jun 2015

Role: Data Warehouse Developer

Responsibilities

- Creation, manipulation and supporting the SQL Server databases.
- Involved in the Data modeling, Physical and Logical Design of Database.
- Contributed to the front end's integration with the SQL Server backend.
- Created Stored Procedures, Triggers, Indexes, User-defined Functions, Constraints etc on various database objects to obtain the required results.

- Import & Export of data from one server to other servers using tools like Data Transformation Services (DTS)
- Wrote T-SQL statements for retrieval of data and was involved in performance tuning of TSQL queries.
- Transferred data from various data sources/business systems including MS Excel, MS Access, Flat Files etc to SQL Server using SSIS/DTS using various features like data conversion etc. In addition, Created derived columns from the present columns for the given requirements.
- Supported team in resolving T-SQL and SQL Reporting Services-related issues, and Proficiency in designing and formatting a variety of reports, including Cross-Tab, Conditional, Drill-down, Top N, Summary, and Sub reports.
- Performed routine maintenance procedures, such as backups, index rebuilds, and statistics updates, to maintain the data warehouse's health and performance.
- Provided via the phone, application support. Developed and tested Windows command files and SQL Server queries for Production database monitoring in 24/7 support.

Environment: IBM WebSphere DataStage EE/7.0/6.0 (Manager, Designer, Director, Administrator), Ascential Profile Stage 6.0, Ascential Quality Stage 6.0, Erwin, TOAD, Autosys, Oracle 9i, PL/SQL, SQL, UNIX Shell Scripts.