

SUMMARY

AI/ML Engineer with 3+ years of experience delivering machine learning and data engineering solutions across ITSM, cloud, and large-scale platforms. Specialized in NLP (NER, sentiment analysis, QA systems, LLMs) and predictive analytics. Strong track record in building scalable ML pipelines, anomaly detection systems, and real-time analytics using Kubernetes, MLflow, Spark, and Kafka. Skilled in AWS, GCP, and Azure for high-throughput data engineering and deployment.

SKILLS

Programming: Python (NumPy, Pandas, Dask), SQL, R, Java, C++

Machine Learning & AI: Supervised & Unsupervised Learning, Ensemble Methods (XGBoost, LightGBM), Neural Networks (CNNs, RNNs, LSTMs, Transformers), Generative AI, Recommender Systems, Predictive Analytics, Time-Series Forecasting, Autoencoders

AI Specializations: Natural Language Processing (Text Classification, NER, Sentiment Analysis, Question Answering), Computer Vision (Image Classification, Object Detection, Segmentation)

MLOps & Deployment: Docker, Kubernetes, MLflow, Kubeflow, TensorFlow Serving, Weights & Biases, CI/CD (Jenkins, GitHub Actions, GitLab CI), Model Deployment (Flask, FastAPI, REST APIs, gRPC), Monitoring (Prometheus)

Cloud & Databases: AWS (SageMaker, Lambda, EC2, S3, Redshift), GCP (BigQuery, Vertex AI), Azure (Azure ML, Cognitive Services), Snowflake, PostgreSQL, MySQL, MongoDB, Redis

Visualization & Analytics: Matplotlib, Seaborn, Plotly, Tableau, Power BI

Data Engineering: ETL Pipelines (Airflow, Luigi), Distributed Data Processing (Spark, Hadoop), Real-Time Streaming (Kafka, Flink, Kinesis), Data Warehousing (Snowflake, BigQuery, Redshift), Data Lakes (Delta Lake, Lakehouse), Data Modeling & Governance

EDUCATION

Master of Science in Information Systems	Dec 2024
Northeastern University, Boston, MA	
Bachelor of Science in Computer Science	Oct 2020
Mumbai University, Mumbai, IN	

EXPERIENCE

ServiceNow, MA AI/ML Engineer	Feb 2025 - Present
- Architected and productionized advanced NLP pipelines (NER, sentiment analysis, QA) leveraging transformer-based LLMs, enabling multi-intent classification and entity extraction at scale, resulting in a 28% uplift in ticket classification accuracy. - Devised predictive analytics ecosystems by integrating XGBoost/LightGBM ensembles with time-series forecasting models, identifying ITSM anomalies and reducing downtime by 22%. - Operationalized BERT, T5, and domain-adapted LLMs into ServiceNow’s intelligent search layer with relevance feedback loops, enhancing contextual retrieval precision by 30%. - Architected fully containerized MLOps pipelines on Kubernetes, orchestrated via MLflow/Kubeflow for automated retraining, version lineage tracking, and zero-downtime model rollouts, compressing release cycles by 40%. - Implemented real-time anomaly detection pipelines fusing Kafka log streams, Spark streaming analytics, and Redis in-memory state stores, achieving low-latency MTTR improvements for infrastructure incident management. - Spearheaded microservice-driven AI workflow automation using FastAPI and REST/gRPC interfaces, integrating seamlessly with ITSM, HR, and Customer Service verticals across ServiceNow’s SaaS ecosystem.	
Experiential AI, MA Data Analyst/Engineer	Jan 2024 - Aug 2024
- Streamlined 10+ enterprise workflows by embedding AI-powered orchestration layers with n8n and Make.com, eliminating repetitive manual dependencies and reclaiming 15 engineer-hours/week. - Engineered a hyperlocal environmental intelligence platform by ingesting sensor telemetry via GCP Pub/Sub, Flink Clusters, BigQuery, sustaining 1M+ records/minute throughput with millisecond-level SLAs. - Constructed high-performance research data marts on BigQuery, augmented with Spark SQL pipelines for distributed preprocessing, enabling near-real-time analytics for 10M+ end users. - Delivered real-time & historical insights to government stakeholders using Looker dashboards powered by federated query engines, directly reducing reporting latency and operational cost by \$10K/month.	
Asfaliea, India Data Engineer I	Dec 2020 – Aug 2022
- Devised multi-source data ingestion (server logs, app telemetry, clickstream events) into AWS S3 via Airflow DAGs, generating unified analytical views and cutting manual ETL overhead by 25%. - Built event-driven ETL workflows on AWS Glue, Lambda, Spark, and Snowflake Streams, scaling ingestion throughput by 15% while maintaining schema evolution guarantees. - Directed cloud migration blueprinting IAM policy hierarchies, S3 tiered storage strategies, and re-platformed orchestration workflows, enhancing reliability metrics by 30%. - Instituted validation frameworks with PyTest and CI/CD hooks to validate transformations, schemas, and anomaly detection eradicating 100+ recurring data defects in production. - Enhanced Snowflake performance (clustering, pruning, caching) and built CloudWatch monitoring with SLA Airflow triggers and ML anomaly detection, accelerating sales forecasting and preventing revenue-impacting downtime.	

PROJECTS

LogiFlow: Logistics Data Pipeline | GCP, Pulumi, FastAPI, Apache Spark, Docker, Airflow

- Formulated a logistics analytics pipeline on GCP using Python, Dataflow, and BigQuery, increasing throughput by 40%.
- Automated supply-chain data-processing workflows by integrating Mage.ai for data validation within Cloud Scheduler, Dataflow pipelines, improving analytics accuracy by 35%.

PodcastGPT: RAG-Based AI for Intelligent Podcast Search | FAISS, OpenAI API, LangChain, FastAPI

- Innovated PodcastGPT by integrating FAISS, OpenAI GPT, and LangChain, and incorporated GitHub Actions for CI/CD to accelerate query retrieval speed by 15%.
- Orchestrated a scalable API using FastAPI, Docker, AWS Lambda, and Step Functions to optimize podcast search efficiency.