

LAKSHMI VENKAT SAI SARATH KUNAPAREDDY

Azure Data bricks Developer, PySpark Developer, AWS Redshift, Oracle PLSQL Developer, Power BI
"Technical Lead in Data Engineering and Data Analytics with Expertise in Retail, Health Care
Insurance."

Expertise in Data Warehouse, Data Modeling and Data Marts

United States, Mob: +1 (224) 800-2495, Email: kunapareddyivsaisarath@gmail.com

LinkedIn ID: <https://www.linkedin.com/in/lakshmi-venkata-sai-sarath-kunapareddy>

PROFILE SUMMARY

- Over 10 years of experience as a **Data Engineer** with expertise in **Spark, PySpark**, Oracle, SQL, PLSQL and **Azure Databricks** for large-scale data processing and analytics.
- Proficient in using Azure services such as **Azure Data Factory, Azure Databricks, Azure SQL Database**, and **Azure Data Lake Storage** for efficient data ingestion and management.
- Extensive experience with **Oracle SQL, PLSQL, Redshift, and PostgreSQL**, including **performance tuning and complex query optimization**.
- Proficient in creating and managing Advanced **ADF Pipelines** with **Linked Services** and **Datasets** to extract, load, and transform data from diverse sources such as **Azure SQL, ADLS, Blob Storage, and Oracle Data Warehouse**.
- Skilled in designing and orchestrating **ETL processes** with **Azure Data Factory**, ensuring reliable and efficient data movement and transformation.
- Adept at developing Spark applications using **Spark SQL** and **PySpark in Azure Databricks**, handling ETL processes for various data formats (**CSV, Parquet**) to support data analysis and transformation.
- Proficient in **building scalable data pipelines** to process, transform, and analyze large datasets using **PySpark** and **Azure Data Factory**.
- Skilled in implementing **delta lakes** for managing structured and unstructured data with **Azure Data Lake Storage Gen2**, ensuring high performance and data integrity.
- Expertise in integrating and managing **Azure Key Vault** for securing sensitive data such as connection strings and credentials in Azure workflows.
- Demonstrated experience in **end-to-end data flow management** using **Azure Data Factory triggers**.
- Skilled in **optimizing Spark jobs** for distributed computing environments, ensuring cost-effective and high-performance data processing.
- Proficient in creating **CI/CD pipelines** using **Azure DevOps** and **GitHub**, automating the deployment of data engineering solutions.
- Extensive experience in **Python development**
- Strong background in **Agile methodologies**, actively participating in sprint planning, daily stand-ups, and retrospective meetings to ensure project delivery within timelines.
- strong expertise in **Oracle SQL, PL/SQL** for data modeling, query optimization, and creating scalable procedures, packages, and functions.
- Experienced in data visualization (**Power BI, Oracle BIP**), advanced performance tuning, and managing data across **Snowflake, PostgreSQL, and AWS (Redshift, S3)**.
- Hands-on experience in **machine learning, data mining**, and Agile methodologies, delivering scalable and impactful solutions.
- Experience working with AWS EMR, RDS, Kinesis and other instances in **Amazon Web Services**.
- Database Migration: Experience in planning and executing database migration projects, including data transfer, transformation, and validation to ensure a smooth transition.

EDUCATION

Roosevelt University, Chicago, USA (Graduated)

Dec-2024

- Master's in computer information science

JNTU, Kakinada, AP, India. (Under Graduated)

May-2016

- BTech. in Electronics and communications engineering

TECHNICAL SKILLS:

Big Data Technologies	:	Spark, PySpark,
Cloud Databases	:	SQL Server, SQL, PLSQL, Oracle, AWS Redshift, GCP Big query,
Snowflake. Reporting Tools	:	Power BI, Oracle BIP reports, SAP BO, Power BI Paginated Reports, Tableau
Cloud Technologies	:	Azure Cloud, AWS Cloud, GCP Cloud, SPARK
SQL Databases	:	SQL, PLSQL, PostgreSQL
NOSQL Databases	:	MongoDB
ETL Tools	:	Azure Data bricks (ETL), Informatica cloud (IICS).
Scripting languages	:	Linux Shell scripting, Unix, Red hat Linux, python.

EXPERIENCE:

Projects:

Client: Aetna Health Insurance USA (CVS Health Corporation USA).

Aug 2022– Dec-2024

Project Title: Data warehousing and Datamart (Data Engineering Team, Power BI)

Languages and Tools: Redshift, Spark, SQL, PLSQL, Python, pyspark, Linux Shell scripting, Oracle, HCS, Cloud, Technologies. **ETL Tools:** Azure Data Bricks.

Role: Technical Lead- (Senior Developer & Admin)

Project Description: Developed enterprise-wide data solutions by integrating structured and unstructured data, including Epic, into Oracle DataMart. Designed scalable ETL pipelines with Spark/PySpark and built Azure Databricks-based data lakes for advanced analytics. Delivered actionable insights via Tableau, Power BI, and SAP Crystal Reports, ensuring scalability and reliability.

- Designed and developed **custom ETL pipelines** using **PySpark in Azure Databricks**, handling large-scale datasets for structured and unstructured data.
- Designed and implemented a **data pipeline on Azure** to ingest, transform, and analyze retail data from various sources, ensuring high data quality and reliability.
- Utilized **Pyspark for data transformation** and automation tasks within the data pipeline, enhancing data processing capabilities.
- Automated **ETL workflows using Azure Data Factory** to streamline data preparation and loading, enabling real-time analysis and reducing manual intervention.
- Designed and implemented **Redshift** schemas optimized for analytical queries.
- Created star and snowflake schemas to support BI tools and reporting.
- Migrated on-premises data warehouses to **Amazon Redshift**.
- Tuned query performance using techniques like **distribution keys, sort keys**, and compression encodings.
- Managed workload using Redshift Workload Management (WLM) to prioritize query execution.
- Built a data warehousing solution with **Redshift** to store cleaned and processed retail sales data, supporting fast querying and report generation.
- Leveraged **Azure Databricks** for advanced analytics and scalable data transformation using **Apache Spark** and Python, enabling efficient processing of high-volume sales data.
- Integrated **Azure Blob Storage, S3 bucket** to store raw and processed data, ensuring cost-effective, high-availability storage for data archival and recovery.
- Using python created **procedures and ETL scripts** at the time of migration
- A vast creation of python scripts by using pandas and NumPy modules for connecting from cloud technologies and transformations.
- Data modeling for Classifying the fact tables, dimension tables.
- Developed reusable Power BI Dashboards and created consolidated views of reports shared across organizational teams.
- Created and optimized Power BI data models, ensuring faster report performance and supporting dynamic visualization.
- Designed and developed custom ETL pipelines using **PySpark in Azure Databricks**, handling large-scale datasets for structured and unstructured data.
- Partnered with stakeholders to **gather requirements, analyze data needs, and design end-to-end scalable data** solutions aligned with project goals.
- Designed and built **ETL pipelines** using **Azure Data Factory** to extract data from **SFTP and Snowflake, Oracle** transforming and loading it into **Azure Data Lake Storage (ADLS) Gen2**.
- Developed and optimized **Spark-based transformations** in **Azure Databricks**, processing large datasets for analytics and reporting.
- Automated data ingestion workflows by configuring **event-driven triggers** in **Azure Data Factory**, ensuring real-time updates and data availability.
- Converted raw data from multiple file formats (e.g., **CSV, JSON, Parquet**) into **Delta Lake** tables to support advanced analytics and improve query performance.
- Created reusable **notebooks in Databricks** to perform data cleansing, transformations, and enrichment for integration into downstream systems.
- Implemented **incremental data processing** and **change data capture (CDC)** mechanisms for efficient updates to the data warehouse.
- Collaborated with **cross-functional teams to monitor, debug**, and optimize **ETL pipelines**, reducing execution times and enhancing reliability.
- Established logging and monitoring frameworks using **Azure Monitor** and **Log Analytics** to track data pipeline health and performance.
- Migrated **Oracle SQL/PLSQL** code to **Snowflake**, including **procedures, functions, and queries**, while validating data accuracy post-migration.
- Enhanced database performance through query optimization, partitioning, indexing, and managing historical data in archive tables.
- Leveraged **Delta Lake on Databricks** to ensure data reliability and consistency, supporting scalable ACID-compliant workflows.
- Automated incremental data processing for real-time and batch datasets using **PySpark**, handling schema evolution and

dynamic data ingestion.

- Designed and implemented fact and dimension tables in data **warehouses** to support business intelligence and analytics.
- Managed SQL/PLSQL-based data migration projects across platforms (e.g., **Oracle to PostgreSQL**), ensuring seamless ETL workflows.
- Led sprint planning, coordinated team activities, and provided technical guidance to ensure on-time delivery of data engineering projects.

Epsilon (PUBLICIS GROUP)

Jan 2022 – Aug 2022

Client: Takeda Pharmaceutical Company USA

Title: Data Integration and Data Injection (Datawarehouse & Data Engineering Team)

Languages and Tools: Oracle, SQL, PLSQL, Azure data bricks, Pyspark, Python, Linux Shell scripting, MongoDB, AWS redshift, IICS.

Role: Associate Technical Lead - PLSQL Developer, ETL Developer.

Project Description: The Takeda Pharmaceutical Company Limited is a Japanese multinational pharmaceutical company. It is the third largest pharmaceutical company in Asia, behind Sinopharm and Shanghai Pharmaceuticals, and one of the top 20 largest pharmaceutical companies in the world by revenue. Where I have been working in Data Injection Team loading the data from different sources mainly from S3 bucket by using python loading the data through Azure Data Bricks cloud ETL tool, IICS to redshift.

- Implemented and maintained serverless architectures using AWS Lambda, Glue, and DynamoDB for scalable, cost-efficient solutions.
- Ensured seamless integration of services within a serverless environment.
- Azure Databricks: Designed and implemented scalable data pipelines using PySpark for transformations, validations, and aggregations; integrated with Azure Synapse, Blob Storage, and Delta Lake.
- Data Quality: Built Data Quality Frameworks in Databricks to validate, clean, and enrich datasets, ensuring compliance with business logic.
- Azure Synapse Analytics: Migrated and optimized data models, built data marts, and developed efficient T-SQL queries for reporting and analytics.
- Azure Storage: Configured Azure Data Lake and Blob Storage to handle large-scale data ingestion, archival, and retrieval needs.
- Monitoring and Automation: Automated workflows using Azure Monitor, ADF, and Databricks job scheduling, enabling real-time and batch data processing.
- Security: Ensured data security with RBAC, Azure Key Vault, encryption, and compliance with organizational standards.
- ETL and Data Integration: Built scalable ETL pipelines using Informatica Cloud, Redshift, and PySpark, handling diverse data formats like JSON, Avro, and Parquet.
- AWS Data Processing: Utilized AWS EMR, DynamoDB, and S3 for high-throughput data transformations, optimized queries, and transaction processing.
- Collaboration and Reporting: Created Power BI dashboards, supported cross-functional teams, and aligned Azure-based solutions with business goals.
- Experience in Power BI (Desktop and Service) with different data sources such as SQL Server, Oracle, Azure SQL Server
- Experienced in installing and configuring both personal and enterprise Data Gateway for power bi reports.
- Experienced in importing data by using different connectors like OData, Web, Microsoft SQL Server, Share point
- Experience in creating a consolidated view of all the reports in a single forum and shared to different groups, teams, and to the entire organization through Teams Channel.
- Created many procedures and complex SQL Queries for the different Requirements.
- Doing a shadow work based on the huge requirements.
- Created many tables like partitioned tables, sub partitioned tables, indexed tables like Global indexes, local indexes based on table spaces.
- Constructing sales aggregates packages and reclassification packages to do daily.
- Reviewing the code and sample technical design templates.
- Writing a python program to load data from S3 to redshift

ALSHAYA OUTSOURCING COMPANY INDIA PRIVATE LIMITED, Bangalore, India

Oct 2020 - Jan 2022

Client: ALSHAYA OUTSOURCING COMPANY INDIA PRIVATE LIMITED Kuwait

Title: Data warehouse and Reporting (ER Team)

Languages and Tools: Spark, PYSPARK, Oracle, SQL, PLSQL, Python, Linux Shell scripting, MongoDB, Cloud: AWS Cloud, Azure Data bricks

Role: Associate Technical Lead - PLSQL Developer, ETL Developer.

Project Description: M.H. Alshaya Co. is a multinational retail franchise operator headquartered in Kuwait. It operates 90 consumer retail brands across the Middle East and North Africa, Russia, Turkey and Europe. In addition to its retail operations, the Alshaya group has holdings in real estate, construction, hotels, automotive and general trading, largely centered in the Middle East. It is a franchise operator for over 70 retail brands, including Mother care, H&M, Debenhams, American Eagle Outfitters, Payless Shoes, Pottery Barn, Starbucks, Dean & DeLuca. Working in a Datawarehouse Team handling the large sets of data for business analytics.

- Utilized AWS CloudWatch and other monitoring tools for proactive resource management and performance tuning.
- Conducted cost-optimization audits for AWS services to minimize resource wastage and enhance ROI.
- Spark and Databricks in AWS: Designed and deployed ETL pipelines using Apache Spark on AWS EMR and Databricks for large-scale data processing, transformation, and aggregation.
- AWS Data Engineering: Built scalable data ingestion workflows using AWS Glue, S3, Lambda, and Redshift; implemented partitioned storage and optimized queries with Athena.
- Azure Databricks: Developed end-to-end data pipelines, data quality frameworks, and automated workflows with Databricks Workspaces, Synapse, and Delta Lake integration.
- Power BI: Created dashboards and reports integrating multiple data sources, implemented Row-Level Security (RLS), and delivered actionable insights for stakeholders.
- Oracle SQL and Database Management: Optimized database performance with partitioned, indexed tables, complex SQL queries, and efficient data migrations between on-premises and cloud platforms.
- Monitoring and Optimization: Conducted Spark job tuning, data flow monitoring with AWS CloudWatch, and Databricks notebooks to enhance job execution and system performance.
- Team Leadership: Led disaster recovery drills, mentored junior team members, and ensured alignment of technical implementations with business requirements.
- Client Engagements: Collaborated with leading retail clients to deliver tailored solutions, improve store performance, and scale revenue growth.
- Achievements: Overachieved targets for 24 months, doubled revenue within two years, and consistently delivered data-driven insights for decision-making.
- Optimized business processes by engineering packages, procedures, and handling data migration from on-premises to cloud solutions, ensuring data integrity.
- Achieved 50% average growth, doubling revenue from Rs. 15 million to Rs. 30 million over two years, while mentoring junior staff and driving team success.
- Collaborated with top retail clients to deliver tailored solutions, improving performance through strong business acumen and communication.
- Proficient in SQL, Power BI, and Excel for data analysis, reporting, and creating actionable insights for sales and stock performance.
- Led end-to-end project execution, ensuring alignment with business objectives and seamless migration between cloud and on-premises databases.

Cinntra Infotech Solutions Pvt. Ltd, Bangalore, India.

May 2019 - Oct 2020

Client: Lifestyle Fashions, Max Fashions -USA

Project Title: Data warehouse and Reporting (ER Data Team)

Languages and Tools: Spark, Pyspark, Oracle, SQL, PLSQL, Python, Linux Shell scripting, Informatica (IICS). **Reporting Tools:** Oracle BIP, Power Bi Reports

Role: software engineer-PL/SQL Developer, ETL Developer

Project Description: Lifestyle is a leading retail destination for branded beauty, fashion and home products. Lifestyle has offered customers a unique shopping experience featuring an eclectic range of cosmetic products, perfumes, accessories and home ware, all under one roof. The purpose of this project is to for reporting, where reports should be submitted based on day-to-day life for businesspeople and department heads either CEO'S or directors. It's purely an OLTP to OTAP which includes large amounts of historical data, aggregated from OLTP Databases (RMS, WMS, OMS, RESA) systems.

- Proficient in SQL development, including creating procedures, functions, cursors, collections, and packages to meet business requirements.
- Regularly refresh data in Power BI on a scheduled basis and distribute reports to users.
- Lead a team in code validation to ensure performance optimization and address any drawbacks before deployment.
- Develop complex queries daily to support various business needs.
- Skilled in constructing dynamic SQL packages and automating query execution.
- Experienced in creating and managing various table types, including partitioned, sub-partitioned, and indexed tables.
- Build sales aggregates and reclassification packages for daily incremental sales analysis.
- Generate reports using SQL queries to fulfill diverse business requirements.
- Handle complex support tickets promptly and effectively.
- Export reports to PDF and share links with clients for easy access.
- Designed and implemented a comprehensive report package for calculating displacement vs. stock on hand

across multiple lifestyle stores.

- Implemented logic for sales analysis, including yesterday, week to date, month to date, year to date, and season to date reporting.
- Proficient in processing large datasets during reclassifications using DML statements, cursor, collection methods.
- Develop templates for unit test cases, technical designs, tables, and packages.
- Implement exchange partitions in packages to optimize data transmission speed.
- Developed packages to track stock on hand and back on hand, facilitating the calculation of damaged stock based on hierarchy.
- Established relationships between various tables to optimize data management.
- Automated package operations using advanced logic to ensure adaptability in any scenario.
- Customized visualizations to meet client specifications effectively.
- Implemented diverse filters such as visual level, page level, and report level filters to enhance reporting flexibility, utilizing Sync Slicer.
- Produced SQL and PLSQL-based reports surpassing the capabilities of Oracle OBIEE and BIP.
- Designed calendars to aggregate sales data for reporting during festival days and end-of-season sales.
- Engineered procedures leveraging store opening days to compute growth formulas based on previous year sales.
- Utilized partitioned tables to enhance data retrieval and insertion efficiency.
- Proficient in utilizing EXCHANGE PARTITIONS for efficient bulk data insertion.
- Skilled in crafting complex SQL queries and interpreting explain plans to address performance issues.
- Created reports using Oracle BI Publisher to fulfill diverse business needs.
- Proficient in scheduling alterations and migrating reports from Oracle 11g to 12c.
- Acknowledged twice for outstanding Team performance, contributing significantly to the company's growth and success.

Client: LOLC (Insurance Management Application) –USA

LQ IT Services PVT LTD, Hyderabad, India

Jul 2016 – Apr 2019

Project Title: Data warehouse and Reporting

Role: software engineer (Developer)

Project Description: It is a web Application. the goal of the project is to maintain all activities of insurance system. any customers without visiting to the office manually he can view the policy details as: policy amounts, policy premium dates, Bonus details etc. through web. the main purpose of the project is any customer can pay the premium amount in any branch and they can pay the amount online.

- Developed backend Database programs, including views, functions, procedures, and packages, utilizing SQL and PL/SQL languages.
- Proficient in understanding project scope, requirements, and business rules for seamless data migration.
- Successfully migrated data from flat files to tables using SQL*Loader, ensuring efficient data transfer.
- Managed data loads and conversions utilizing SQL*Loader, with a focus on resolving errors caused by poor data quality.
- Optimized performance by tuning complex queries and strategically creating indexes on essential columns.
- Implemented database objects using advanced concepts such as collections, dynamic SQL, and bulk operations.
- Extensive experience in PL/SQL database programming, including creating, testing, and analyzing triggers, procedures, and functions.
- Skilled in data manipulation language (DML) for updating database records.
- Designed and developed end-to-end programs for ETL workflows, ensuring smooth data processing from extraction to loading.
- Created mapping and templates for informatica workflow.
- Monitoring the informatica jobs and their duties for the successful Delivery on time.
- Appointed as a L1 support from the development team to support team to lead the issues on jobs.
- Appreciations from the clients for point of contact.
- Designing the mappings in informatica for various transformations and workflows.

Date:

Lakshmi Kunapareddy