

Qiming Xie

Data Engineer

TX | (508) 271-7892 | qiming.x@newjobshunt.com

SUMMARY

Data Engineer with over 6 years of experience specializing in designing and implementing cloud-based data solutions using AWS. Proficient in developing and optimizing both ETL and ELT pipelines with tools like Apache Airflow, DBT, and AWS Glue, achieving up to 50% reduction in processing times. Skilled in handling real-time and batch data processes utilizing Hadoop, Spark, and Databricks, alongside comprehensive database management across various SQL and NoSQL platforms. Committed to maintaining data integrity and compliance with standards such as GDPR and HIPAA through rigorous data governance and security practices. Expert in automating CI/CD processes, containerization, and executing cloud cost optimizations to boost system performance and scalability.

SKILLS

Methodologies: SDLC, Agile, Waterfall

Programming Languages: Python, R, SQL, Java

IDEs: PyCharm, Jupyter Notebook, IntelliJ IDEA, Visual Studio

Big Data Ecosystem: Hadoop, MapReduce, Hive, Apache Spark, HDFS, Yarn, Data Lake, Data Warehouse

Cloud Technologies and Infrastructure: AWS (S3, DynamoDB, Lambda, EC2, Redshift, Glue, Athena, CloudWatch, Aurora, EMR)

ETL & Data Processing Tools: AWS Glue, Apache Airflow

Streaming & Analytics Tools: AWS Kinesis, AWS Athena

Monitoring & Automation: AWS CloudWatch, AWS Lambda

Packages & Libraries: NumPy, Pandas, Matplotlib, SciPy, Seaborn, TensorFlow

Development and Operations Tools: Git, GitLab, Jenkins (CI/CD), Docker, Kubernetes, Grafana, Jira, GitHub

Databases: MySQL, PostgreSQL, MongoDB, Cassandra, Elasticsearch, RDBMS, NoSQL

EDUCATION

Bachelor in Project Manager | Tianjin University of Technology, China

EXPERIENCE

ServiceNow, TX | Aug 2024 – Current | Cloud Data Engineer

- Developed scalable cloud-based data architectures using AWS, optimizing data storage, retrieval, and processing. Integrated data pipelines with business applications, improving query performance and reducing latency by 30%.
- Built and optimized ETL/ELT pipelines using Apache Airflow, AWS Glue, and DBT. Streamlined data ingestion, transformation, and validation, reducing processing time by 40% while improving data consistency and accuracy.
- Optimized cloud data warehouses and lakes by resolving bottlenecks, improving query execution, indexing, and storage strategies.
- Enforced data security measures by implementing data governance frameworks, encryption protocols, IAM policies, and access controls to protect sensitive data, prevent breaches, and ensure compliance with industry standards.
- Implemented cloud data solutions on AWS, leveraging Snowflake, Redshift, and Databricks for batch and real-time processing.
- Standardized technical documentation for cloud architectures, data pipelines, data governance policies, and security frameworks. Developed runbooks, SOPs, and best practices to streamline onboarding, troubleshooting, and cross-team collaboration.

The Hartford, TX | Nov 2022 – July 2024 | Data Engineer

- Optimized ETL pipelines and data workflows, reducing financial data processing time by 30%.
- Improved query performance by 40% through data partitioning and indexing in AWS Redshift and Snowflake analytics.
- Built real-time data streaming solutions with Apache Kafka and Apache Spark, enhancing fraud detection accuracy to 95%.
- Ensured regulatory compliance (GDPR, PCI-DSS, Basel III) with encryption and audit logging, cutting security violations by 20%.
- Integrated risk modeling pipelines and machine learning algorithms, reducing fraud detection false positives by 25%.

Unum, TX | May 2020 – Oct 2022 | Data Engineer

- Implemented a data governance framework to ensure compliance with GDPR and HIPAA, improving security and auditability.
- Developed Python models to detect fraudulent insurance claims, saving \$2M annually.
- Migrated legacy data systems to a cloud-based architecture using AWS S3, increasing scalability and reducing costs by 30%.
- Built ETL workflows with Apache Airflow to integrate data from policy management systems and customer databases.
- Optimized data retrieval with indexing and query tuning, reducing response times and improving performance.
- Standardized data collection across multiple sources to enhance accuracy and support data quality initiatives.

DXC Technology, TX | June 2019 - Apr 2020 | Data Engineer

- Migrated on-premise data infrastructure to AWS Cloud (S3, Redshift, DynamoDB, Aurora, EC2), improving scalability and reducing operational costs by 40%.
- Engineered and optimized a Data Lake using HDFS, integrating structured and unstructured data to enhance analytics.
- Automated ETL processes using advanced tools, reducing processing time by 30% and streamlining data handling and integration.
- Developed and maintained real-time data streaming architectures, improving data ingestion and analytics capabilities.
- Administered and optimized MongoDB and Cassandra databases, ensuring high availability, fault tolerance, and scalability.
- Designed and implemented frameworks to validate data accuracy across multiple sources, improving quality and reliability.

CERTIFICATION

AWS cloud practitioners | AWS solution architecture associated | AWS developer associate | HashiCorp Terraform |

Oracle Gen AI professional | Google data analyst | Udacity data engineer