

RAKSHITH L
RAKSHITH.LKR9@GMAIL.COM
205-687-8989
Sr. Data Engineer

PROFESSIONAL SUMMARY

- Over 10+ of hands-on expertise as a Senior Data Engineer specializing in Database Development, ETL Development, Data Modeling, Report Development, and Big Data Technologies.
- Managed data catalogs and implemented data governance, curating extensive datasets to ensure 99% data integrity and streamlined accessibility for cross-departmental use.
- Proficient in programming languages like Python (Pandas, NumPy, Py Spark, scikit-learn, Py Torch), SQL (including PL/SQL for Oracle), Scala, and PowerShell.
- Have Extensive Experience in IT data analytics projects, Hands on experience in migrating on premise ETLs to Google Cloud Platform (GCP) using cloud native tools such as BIG query, Cloud Data Proc, Google Cloud Storage, Composer.
- Practical understanding of the Data modeling (Dimensional & Relational) concepts like Star - Schema Modeling, Snowflake Schema Modeling, Fact and Dimension tables.
- Proficient in Apache Spark, Apache Airflow, Hadoop, Hive, Sqoop, Kafka, Impala, and Apache Beam for large-scale data processing and analytics.
- Experience with cloud data warehouses like Redshift,
- Snowflake, and big Query for scalable data storage and retrieval.
- Expert in creating ETL packages, migrating data, cleaning data, and backing up data files, and synchronizing daily transactions using Informatica, Alteryx, Autosys etc.
- Skilled in data visualization tools such as Tableau, Power BI, Google Data Studio, and Quick Sight for creating insightful reports and dashboards.
- Hands-on experience with data integration tools like Informatica, Talend, and SSIS for seamless data flow across systems.
- In-depth knowledge of various database systems including SQL Server, Cosmos DB, Oracle, PostgreSQL, Cassandra, MySQL, and DynamoDB for efficient data storage and retrieval.
- Proficient in handling data formats like JSON, XML, and Avro for data interchange and storage.
- Hands on Bash scripting experience and building data pipelines on Unix/Linux systems.
- Expert in developing SSIS Packages to extract, transform and load (ETL) data into data warehouse/data marts from heterogeneous sources.
- Hands on experience with different programming languages such as Python, SAS.
- Engineered scalable web services and RESTful APIs with Python, integrating complex business logic that facilitated real-time data interactions and improved system responsiveness by 40%.
- Experience in designing Data Marts by following Star Schema and Snowflake Schema Methodology.
- Managed extensive BI solutions, utilizing SQL and NoSQL databases, Informatica, and tools like Tableau and Power BI to drive data strategy, improving data visualization and decision-making processes.
- Very keen on knowing newer techno stack that Google Cloud platform (GCP) adds.
- Worked with SAS Enterprise guide and have written numerous passes thru Hive SQL queries.
- Proficient in Excel Advanced functions, pivot tables, and V Lookups for data analysis and reporting.

- Familiarity with version control systems like Git, GitHub, and Bitbucket for collaborative development and code management.
- Proficient in project management tools like Bugzilla, Confluence, SharePoint, JIRA, Agile, Scrum, and Kanban for efficient project execution and collaboration.
- Analyzed business needs, designed efficient processes, and managed development teams to deliver successful projects.
- Implemented data pipelines, ensured system integration, and adapted to new technologies for continuous improvement.
- Tackled complex data challenges with strong analytical skills and a drive to contribute to a dynamic and innovative environment.

TECHNICAL SKILLS:

Cloud Platforms:	Google Cloud Platform (GCP)
Big Data Processing and Analytics:	Apache Spark, Apache Airflow, Hadoop, Hive, Sqoop, Kafka, Impala, Apache Beam
Programming and Scripting:	Python (Pandas, NumPy, PySpark, scikit-learn, PyTorch), SQL (including PL/SQL for Oracle), Scala, PowerShell
Data Integration and ETL Tools:	Glue, Informatica, Talend, SSIS
Containerization and Orchestration:	Docker, Kubernetes
Version Control and Collaboration:	Git, GitHub, Bitbucket
CI/CD:	Azure DevOps, Jenkins, AWS CodePipeline
Data Warehousing and Database Management:	Redshift, Snowflake, BigQuery, SQL Server, CosmosDB, Oracle, PostgreSQL, Cassandra, MySQL, DynamoDB
Data Visualization and BI Tools:	Tableau, Power BI, Google Data Studio, QuickSight
Security and Access Control:	SSL/TLS, AES encryption standards
Miscellaneous Tools and Technologies:	JSON, Excel Advanced functions, pivot tables, V Lookups, Bugzilla, Confluence, SharePoint, JIRA, Agile, Scrum, Kanban, Window, Linux, UNIX, macOS

PROFESSIONAL EXPERIENCE:

Client: Wells Fargo, Summit, NJ

May 2022– Present

Role: Azure Data Engineer /Snowflake / BI Developer

Responsibilities:

- Developed Logical and physical data model using ERWIN/SAP Power Designer and mapped the data into database objects.
- Designed, developed, and maintained ETL processes using Informatica PowerCenter to ensure efficient data integration and transformation.
- Gathered report requirements and determined the best solution to provide the results in either a Reporting Services Report or Power BI or Tableau.

- Expert in Data Migration, Data conversion, Data modeling, Normalization and SQL Server Enterprise manager.
- Primarily involved in Data Migration using SQL, SQL Azure, Azure storage, and Azure Data Factory, SSIS, PowerShell.
- Writing Snowflake script to load the data from Azure blob storage to Snowflake database in the respective tables.
- Processed data in Alteryx for tableau reporting.
- Creating logic and functionality in the Azure Data Lake, Data Factory, SQL Server Database, and SQL Data Warehouse environment. Experience in DWH/BI project implementation using ADF.
- Developed Azure Data Factory pipelines to integrate data from Apache Hive, EDL, APIs, Dynamic365, SFTP and various flat-file types into Azure SQL.
- Develop stored procedures/views in Snowflake and use in Talend for loading Dimensions and Facts.
- Hands on experience with Azure DevOps, CICD pipeline in GitHub Actions across a broad variety of technologies with deployment targets including Azure Database and EDW.
- Design ETL Framework, Developed and deployed SSIS packages for ETL from OLTP and various sources to staging and staging to Data warehouse using Foreach Loop Container, Execute Package task, Execute SQL Task, Sent Mail task, Lookup, Fuzzy Lookup, Derived Columns, Condition Split, CDC Control Task, CDC Splitter, CDC Source, Slowly Changing Dimension and more.
- Development, monitoring and maintenance of data integration and processing tools and pipelines Azure Data Factory Designed and build python and metadata driven solution for managing ELT loading into Snowflake (all rules stored as metadata so can be managed in a central metadata model).
- Designed and developed ETL workflows and datasets in.
- Design ETL framework to source data from various sources and load it into fact/dimension tables using SSIS/Informatica/Talend/ADF (Azure Data Factory).
- Developed different kinds of reports such as Sub Reports, Charts, Matrix reports, linked reports using SSRS / Tableau/Power BI.
- Utilized Power BI (Power View) to create various analytical dashboards that depict critical KPIs.
- Expertise in writing complex DAX functions in Power BI and Power Pivot.
- Developed Reports for business end users using Report Builder with updating Statistics using Tableau/Power BI & SSRS.
- Use Snowflake advance features like Multi Clustering, Warehouses, Cluster keys to improve the performance to meet the daily client deliverable.
- Created drill down and drill through report to go to details and child report within power bi based on hierarchy.
- Utilized Power BI gateway to keep dashboards and reports up to-date with on premise data sources.
- Migrated clients off EMR (Hadoop, Sqoop, Hive, Redshift) to AWS Data Lake/Snowflake Azure.
- Utilized Power Query in Power BI to Pivot & Un-pivot the data model for data cleansing and data massaging.
- Designed, developed, and deployed Kafka-based solutions for real-time data processing and event-driven architecture implementation.
- Strong knowledge of Kafka architecture, data streaming concepts, and distributed systems
- Developed and maintained Apache Spark applications for data processing and analysis.

- Built functionality into the ELT framework to scale-up Snowflake on-demand.
 - Created many calculated columns and measures using DAX in Power BI based on report requirements.
 - Design ELT framework to ingest data into data lake using nifi
- Tech Stack:** MS-SQL Server 2019/2014/2012/2008 R2, Azure Synapse, Snowflake, Azure SQL, ADF (Azure Data Factory), Azure Data Bricks, Power BI, Tableau, MS SQL Server Integration Services (SSIS), MS SQL Server Reporting Services (SSRS), MS SQL Server Analysis Services (SSAS), T-SQL, SQL Profiler, CI/CD , Azure DevOps, Oracle, Oracle apex.

Client: Johnson & Johnson, Santa Clara, CA

June 2019 - April 2022

Role: Data Engineer

Responsibilities:

- Utilized Google Cloud Platform (GCP) services including big Query, Dataflow, and Data proc for data processing, analysis, and orchestration, optimizing performance and scalability.
- Implemented Pub/Sub and Cloud Storage for real-time data ingestion and storage, ensuring reliable and scalable data pipelines.
- Managed Cloud SQL and Cloud Datastore for structured and unstructured data storage, maintaining data integrity and accessibility.
- Designed and implemented data streaming pipelines using Cloud Pub/Sub, Apache Beam, and Apache Kafka, enabling real-time data processing and analytics.
- Leveraged Apache Spark and Hadoop for distributed data processing and analytics, handling large-scale datasets efficiently.
- Orchestrated data workflows using Apache Airflow, ensuring automation and scheduling of data pipelines for timely processing.
- Integrated data sources using Sqoop and Informatica, facilitating seamless data extraction, transformation, and loading processes.
- Utilized Python with Pandas and NumPy for data manipulation, analysis, and modeling, enhancing data processing capabilities.
- Developed data visualizations and dashboards using Data Studio and Google Analytics, providing actionable insights to stakeholders.
- Managed code repositories and collaborated with teams using GitHub, ensuring version control and code quality.
- Orchestrated containerized applications using Docker and Kubernetes, ensuring scalability and reliability of deployed data solutions.
- Automated infrastructure deployment and management using Terraform, optimizing resource utilization and cost efficiency.
- Implemented continuous integration and deployment pipelines using Jenkins, ensuring seamless delivery of data solutions.
- Utilized ELK Stack (Elasticsearch, Logstash, Kibana) for log analysis and monitoring, ensuring data visibility and troubleshooting capabilities.
- Handled data serialization using Avro, ensuring efficient data storage and processing.
- Managed relational databases including PostgreSQL and NoSQL databases like Cassandra, ensuring data availability and performance.
- Developed and deployed machine learning models using TensorFlow, enhancing data analytics and predictive capabilities.
- Utilized VS Code for code development and debugging, ensuring code efficiency and reliability.
- Contributed to Agile and Kanban methodologies, participating in sprint planning, daily stand-ups, and backlog grooming to deliver data solutions efficiently.

- Collaborated and documented technical specifications using Confluence, ensuring knowledge sharing and documentation of data solutions.

Tech Stack: GCP, Big Query, Apache Beam, Apache Spark, Apache, Airflow, Hadoop, Kafka, Sqoop, Informatica, Python, Scala, Data Studio, Google Analytics, GitHub, Terraform, Jenkins, ELK Stack, Avro, PostgreSQL, Cassandra, TensorFlow, VS Code, Agile, Kanban, Confluence.

Client: Mercury Insurance, San Antonio, TX

Feb 2017 - May 2019

Role: Data Engineer

Responsibilities:

- Utilized Hadoop, Spark, and Hive to process and analyze large volumes of data, optimizing performance and scalability for big data applications.
- Implemented Sqoop for efficient data transfer between Hadoop and relational databases, ensuring seamless data integration and synchronization.
- Developed complex SQL and PL/SQL queries to extract, transform, and load data from diverse sources into data warehouses, improving data accessibility and analysis capabilities.
- Managed AWS resources including EC2 instances, S3 buckets, RDS databases, and Lambda functions, leveraging cloud services for scalable and cost-effective data solutions.
- Utilized Python with NumPy and Pandas for data manipulation, statistical analysis, and machine learning model development, enhancing data processing workflows.
- Maintained version control and collaborated with teams using Git, ensuring code quality, and facilitating efficient project management.
- Designed and implemented ETL processes using SSIS (SQL Server Integration Services), ensuring data quality and consistency across data pipelines.
- Managed bug tracking and issue resolution using Bugzilla, ensuring data integrity and timely resolution of data-related issues.
- Contributed to Agile and Kanban methodologies, participating in sprint planning, daily stand-ups, and backlog grooming to deliver data solutions efficiently.
- Developed interactive dashboards and visualizations using Tableau, providing stakeholders with actionable insights and data-driven decision-making capabilities.
- Collaborated with SharePoint for document management and collaboration, ensuring data governance and compliance with organizational standards.
- Implemented data security measures and access controls, ensuring data privacy and compliance with regulatory requirements.
- Conducted performance tuning and optimization of database queries and processes, improving data processing efficiency and reducing latency.
- Participated in data architecture design and data modeling activities, ensuring scalability, flexibility, and performance of data solutions.
- Provided technical expertise and support to cross-functional teams, contributing to the successful delivery of data projects and initiatives.

Tech Stack: Hadoop, Spark, Hive, Sqoop, SQL, PL/SQL, AWS, EC2, S3, RDS, Lambda, Python (NumPy, Pandas), Git, SSIS, Bugzilla, Agile, Kanban, Tableau, SharePoint.

Value labs, Hyderabad, India

June 2013 - Nov 2016

Role: Data Analyst/ Engineer

Responsibilities:

- Utilized Python, pandas, NumPy, and scikit-learn for data cleaning, preprocessing, analysis, and machine learning model development, resulting in improved data quality and predictive accuracy.
- Leveraged SQL to query, manipulate, and extract insights from large datasets stored in Oracle databases, ensuring data integrity and optimizing data retrieval performance.
- Demonstrated expertise in Excel Advanced functions, pivot tables, and V Lookups to create interactive dashboards and reports for stakeholders, facilitating data-driven decision-making processes.
- Implemented Apache Spark for big data processing and analytics, handling large-scale datasets efficiently and performing distributed computing tasks for faster data processing.
- Designed and implemented data integration workflows using Talend, ensuring seamless data flow between heterogeneous systems and maintaining data consistency across platforms.
- Managed version control and collaborated with teams using Git, tracking changes and resolving issues efficiently to maintain code quality and project progress.
- Utilized Bugzilla for bug tracking and issue management, ensuring timely resolution of data-related issues and maintaining data accuracy.
- Applied data engineering techniques in Hadoop ecosystem, including HDFS, Hive, and HBase, to store, process, and analyze large volumes of structured and unstructured data.
- Collaborated with cross-functional teams to develop and deploy data pipelines and ETL processes, ensuring data availability and reliability for business analytics and reporting.
- Contributed to data governance initiatives by establishing data quality standards, monitoring data quality metrics, and implementing data cleansing and enrichment strategies.

Tech Stack: Python, pandas, NumPy, scikit-learn, SQL, Excel Advanced functions, pivot tables, V Lookups, Spark, Talend, Oracle, Hadoop, Hive, HBase, Git, Bugzilla.