

Srikanth S
Senior Data Scientist
+1 (214) 494-9454 | srikanth.soma06@gmail.com
[linkedin.com/in/srikanths](https://www.linkedin.com/in/srikanths)

PROFESSIONAL SUMMARY:

- 10+ years of experience in machine learning engineering and data science, with deep expertise in deploying ML models in production environments.
- Designed and maintained end-to-end ML pipelines, including data ingestion, model training, deployment, and monitoring.
- Strong programming skills in Python, with extensive use of Scikit-learn, TensorFlow, Keras, and PySpark for building scalable ML models.
- Experienced in cloud platforms: AWS (S3, EC2, Lambda, Snowflake); adaptable to GCP services such as BigQuery, Cloud Composer, and GCS.
- Proficient in DevSecOps practices including CI/CD, version control (GitHub), Docker, and automation with tools like Jenkins.
- Worked with Kubernetes, Airflow, and Terraform for container orchestration and pipeline automation in hybrid cloud environments.
- Hands-on experience integrating APIs and real-time data pipelines (AWS SQS, Adobe Analytics, Snowflake) for ML deployment and monitoring.
- Proficient in data visualization tools such as Tableau, Matplotlib, and ggplot2 to create visually powerful and actionable interactive reports and dashboards.
- Proven ability to enforce data quality, security, and governance including PII tagging and compliance (CCPA, internal policies).
- Collaborated with cross-functional teams of ML Engineers, Data Scientists, Data Engineers, and IT to build robust, scalable analytics solutions.
- Excellent problem-solving mindset, with a strong focus on model reliability, business alignment, and continuous improvement.

TECHNICAL SKILLS:

Languages	Python, SQL, R
Machine Learning & Deep Learning	Scikit-learn, TensorFlow, Keras, PyTorch, XGBoost, CNN, LSTM, NLP (NLTK, spaCy), Feature Engineering, Hyperparameter Tuning, Model Deployment, PCA, SMOTE
MLOps	GitHub, Jenkins, Docker, Kubernetes, Terraform, Tekton, CI/CD Pipelines, Model Monitoring, ML Workflow Automation
Cloud Platforms	AWS (EC2, S3, Lambda, Snowflake, Boto3, SQS, Secrets Manager), GCP (BigQuery, Google Cloud Storage, Cloud Composer, Cloud Run)
Data Engineering & Processing	PySpark, Apache Spark (MLlib), Hadoop (HDFS, Hive, Pig), Airflow, DataProc, OpenShift, API Integration, Data Pipelines
Databases	Snowflake, PostgreSQL, SQL Server, Oracle
Visualization & Reporting	Tableau, Matplotlib, Seaborn, ggplot2
Project Tools & Agile	JIRA, Confluence, Agile/SCRUM, Alation (Data Catalog), Adobe Analytics
Others	Informatica PowerCenter, Erwin (Data Modeling), OpenCV (Computer Vision), StatsModels

PROFESSIONAL EXPERIENCE:

Client: Morgan Stanley Alpharetta, GA

March 2023 - Present

Role: Senior Data Scientist

Roles & Responsibilities:

- Automated data cataloging using Python and Alation, enabling streamlined metadata management and tagging of PII data for regulatory compliance.
- Implemented robust PII tagging in enterprise data catalog; improved data governance and supported secure analytics environments in line with compliance protocols.
- Developed production-grade ML models (Logistic Regression, KNN, Gradient Boosting) using Python; leveraged Pandas, Scikit-learn, and Matplotlib for modeling, evaluation, and visualization.
- Built modular AI recommendation engines integrated with real-time pipelines to drive personalized insights for business intelligence teams.
- Designed and optimized predictive ML models (classification, regression, clustering); integrated results into automated data workflows.
- Conducted iterative analysis and enhancement of ML algorithms; collaborated with analysts to refine model features and boost performance.
- Partnered with stakeholders to translate analytics into business action plans, leveraging predictive modeling to identify customer segments and business KPIs.
- Applied ensemble techniques (bagging, boosting) to enhance model accuracy; automated deployment via AWS, simulating DevSecOps practices.
- Delivered end-to-end ML workflows from Snowflake data ingestion to model deployment; integrated feature engineering and evaluation using Scikit-learn.
- Built Python EDA modules to automate data profiling and visualizations, facilitating feature insights across large datasets.
- Built and tuned Random Forest models to analyze behavioral web data; deployed trained models to predict customer engagement outcomes.
- Developed deep learning models (CNN) using TensorFlow and Keras for classification problems; fine-tuned model performance with hyperparameter tuning.
- Engineered advanced ML features using PCA and optimized model accuracy through hyperparameter tuning and feature selection pipelines.
- Built and deployed LSTM models for time-series forecasting; implemented model versioning and evaluation metrics for ongoing performance tracking.
- Embedded AI modules into software systems to enhance performance and analytics functionality; collaborated with product teams for seamless integration.
- Applied deep learning frameworks (TensorFlow, PyTorch) to production ML pipelines for improved model robustness and accuracy.
- Built and optimized CNNs using pre-trained Keras models (VGG, ResNet) for classification tasks; adjusted epochs and optimizer parameters for peak accuracy.
- Applied deep learning with Keras (VGG16, VGG19, ResNet) and OpenCV for real-time image classification and processing use cases.
- Developed Python connectors for Adobe Analytics API to extract web metrics; integrated analytics data into Snowflake for real-time reporting.
- Engineered Snowflake stored procedures for sensitive data identification and anonymization using salted hashing, ensuring CCPA compliance.
- Automated data exchange with AWS services (S3, Secrets Manager, SQS) using Boto3, enabling scalable and secure ML data pipelines.
- Built data ingestion pipelines to process subscription data from AWS SQS into Snowflake; supported near real-time analytics and metadata storage.
- Automated reporting workflows to deliver actionable analytics to WarnerMedia brands; leveraged Python scripts for efficient data extraction.
- Developed serverless functions (Lambda) for privacy automation; processed and posted consumer files to Malibu privacy systems using secure APIs.
- Led all stages of ML project lifecycle including data collection, preprocessing, modeling, evaluation, and deployment in AWS-based environments.
- Used Python stack (NumPy, Pandas, SciPy, Matplotlib) for advanced data transformations, statistical modeling, and exploratory analysis.

- Performed dual-language analysis (Python, R) to ensure model validity and cross-functional interpretability.
- Built robust ETL pipelines to extract and clean data from Snowflake; ensured high-quality training datasets for ML development.
- Addressed class imbalance in fraud detection models using SMOTE, under-sampling, and over-sampling techniques for balanced training.
- Applied PCA and one-hot encoding for dimensionality reduction and feature scaling; integrated results into ML pipelines for model training.

Client: Cardinal Health Dublin, OH

April 2021 - February 2023

Role: Senior Data Scientist

Roles & Responsibilities:

- Built scalable ML solutions on AWS cloud, developing infrastructure for big data processing and model deployment.
- Developed Spark-based Python modules using PySpark and MLlib for predictive analytics on distributed Hadoop clusters.
- Engineered distributed Random Forest models via PySpark MLlib, enabling high-speed processing of large-scale datasets.
- Optimized model performance using cross-validation and parameter tuning; integrated validation within automated pipelines.
- Designed, implemented, and deployed AI models to production environments, solving complex business problems with measurable outcomes.
- Boosted model performance using ensemble techniques (Bagging, Boosting); deployed models via CI/CD pipelines on AWS.
- Built monitoring dashboards with Tableau to track real-time model performance and health across deployed environments.
- Designed end-to-end ML pipelines with automation for training, validation, deployment, and monitoring across AWS and Spark ecosystems.
- Implemented DevOps practices using GitHub and Jenkins for version control, CI/CD automation, and secure deployment workflows.
- Collaborated across teams using agile tools like JIRA, Confluence, and Git-based versioning (GitHub, Bitbucket).
- Applied ML algorithms for business process automation and data-driven decision support systems.
- Led full ML lifecycle from data sourcing to deployment; ensured scalable workflows for big data preprocessing and analytics.
- Developed fraud detection models using behavioral analytics and supervised learning, ensuring high-accuracy classification in real-time.
- Built ETL workflows to extract SQL Server data, ingest into Hadoop HDFS, and transform via Hive and Pig for ML processing.
- Applied advanced feature engineering to boost model efficiency; utilized PCA, correlation analysis, and variable transformation.
- Performed comprehensive data cleaning and transformation using Pandas and NumPy to ensure model-ready datasets.

Client: American Express, Phoenix, AZ

July 2019 - March 2021

Role: Data Scientist

Roles & Responsibilities:

- Built ML models using Python (Scikit-learn, NumPy, Pandas) and R for customer behavior analysis and risk scoring.
- Developed sentiment analysis pipelines using Python NLP libraries (NLTK), showcasing transferable skills for unstructured data processing.
- Built predictive models (Logistic Regression, Random Forest, KNN) with ensemble techniques to identify customer churn patterns.
- Applied RFM analysis and clustering algorithms to segment customers and deliver targeted marketing insights.
- Partnered with data engineers to build robust ETL workflows; extracted and transformed high-volume datasets using optimized SQL.
- Wrote complex SQL queries to extract business-critical data; ensured high-quality inputs for ML models and dashboards.

- Retrieved and preprocessed large datasets from Hadoop environments; built ML-ready data structures for downstream modeling.
- Conducted multivariate statistical analysis to identify key variables influencing model outcomes.
- Evaluated models using advanced classification metrics (AUC-ROC, F1, Confusion Matrix); validated model stability pre-deployment.
- Designed and deployed recommendation engines using collaborative filtering; integrated Google Analytics data into the modeling workflow.
- Applied collaborative filtering for personalized content delivery, aligning with Stefanini's focus on actionable ML insights.
- Created executive dashboards using Tableau and Matplotlib to visualize ML outcomes and customer trends.

Client: Robinhood Menlo Park, CA

January 2017 - June 2019

Role: Machine Learning Engineer

Roles & Responsibilities:

- Designed statistical validation strategies (ANOVA, Chi-square, cross-validation) to assess model effectiveness and guide tuning.
- Deployed and optimized ML models on AWS EC2; fine-tuned hyperparameters to maximize predictive accuracy in production settings.
- Developed robust preprocessing pipelines to ensure data consistency across training and live inference environments.
- Delivered serverless model deployment using AWS Lambda; incorporated feedback loops for continuous retraining and improvement.
- Built Tableau dashboards to monitor model outputs, business trends, and benchmark KPIs post-deployment.
- Operated in Agile/Scrum workflows; collaborated cross-functionally with engineers, analysts, and DevOps for sprint-based ML delivery.
- Led data intake from stakeholders; developed standardized ETL pipelines to support uniform model training inputs.
- Retrieved and transformed data from Oracle databases for model development and analytics workflows.
- Applied advanced preprocessing in Python (Pandas, Scikit-learn), including imputation and data balancing for fair model training.
- Addressed class imbalance via over/under-sampling techniques; improved model generalizability and stability.
- Conducted feature engineering using PCA, normalization, and categorical encoding; optimized dimensionality for efficient modeling.
- Processed and modeled 12M+ EMR records using comorbidity scoring frameworks to support healthcare analytics solutions.

Client: Developp Solutions, India

June 2013 - September 2016

Role: Data Analyst

Roles & Responsibilities:

- Extracted and consolidated structured datasets from Oracle and internal platforms to support analytics and reporting initiatives.
- Developed advanced SQL procedures and triggers to support data transformation, validation, and loading workflows.
- Facilitated deployment and migration of ETL mappings across development, testing, and production environments for business continuity.
- Created dimensional and relational data models using Erwin to represent enterprise data structures and support analytics architecture.
- Led dimensional modeling for the enterprise data warehouse, ensuring schema alignment with reporting and analytics needs.
- Performed data quality and cleansing using Oracle ETL tools; standardized and enriched address data for downstream applications.
- Designed Informatica ETL pipelines using transformations (Joiner, Router, Aggregator, Expression) to load data into staging and fact tables.
- Processed and summarized large datasets using SQL and Informatica for efficient decision support and dashboarding.
- Produced comprehensive documentation including data lineage (STTM) and test cases to ensure traceability and quality in data operations.