SIDDHARTHA.C

DATA ENGINEER

Mobile: +1 469-284-9966 | Email: Siddharthac761@gmail.com | Linkedin

SUMMARY:

Results-driven Data Engineer Data Analyst with 5+ years of experience designing and optimizing scalable data solutions for enterprise environments. Proven expertise in ETL development, cloud data engineering, big data analytics, and business intelligence. Skilled in AWS, SQL, Python, and Apache technologies to streamline data pipelines, automate workflows, and extract insights that drive business growth. Adept at data visualization, predictive analytics, and statistical modeling, enabling data-driven decision-making. Passionate about improving data accessibility, optimizing BI tools, and transforming raw data into actionable insights.

SKILLS:

Methodology: SDLC, Agile, Waterfall
Programming Language: R, Python, SQL, T-SQL, Scala
IDE's: PyCharm, Jupyter Notebook, RStudio, Visual Studio Code, IntelliJ IDEA, Sublime Text, Atom, Spyder, Eclipse, DataGrip
Big Data Ecosystem: Hadoop, MapReduce, Hive, Apache Spark, Pig, Kafka, Flink, Presto
ETL Tools: SSIS, AWS Glue, Data Integration,
Cloud Technologies: AWS (EC2, S3), Azure, Google Cloud Platform (GCP), Snowflake,
ML Algorithm: Linear Regression, Logistic Regression, Decision Trees, Supervised Learning, Unsupervised Learning, Classification, SVM, Random Forests, Naive Bayes, KNN, K Means, CNN
Packages: NumPy, Pandas, Matplotlib, SciPy, Scikit-learn, Seaborn, TensorFlow
Reporting Tools: Tableau, Power BI, SSRS
Data Management Platforms: Data Warehouse, Data Lake, ODS (Operational Data Store), Clinical data management,
CI/CD Tools: Jenkins, Control-M, Spinnaker, CircleCI, Apache Airflow, Azure
Database: MongoDB, MySQL, Oracle
Other Tools: Git, MS Office, MS Excel
Operating Systems: Windows, Linux

EDUCATION:

Master of Science in Information Science | University of North Texas, Denton, TX
Bachelor of Technology in Computer Science | KL University, AP, India.

EXPERIENCE:

Cigna, TX | April 2024 - Current | Data Engineer

- Designed and optimized ETL workflows using AWS Glue & Redshift, reducing data processing time by 30% and enhancing analytics capabilities.
- Built and managed high-volume data pipelines across cloud platforms, improving data ingestion efficiency by 40%.
- Automated data transformation processes, reducing manual intervention by 50% and improving system reliability.
- Engineered a robust data processing framework using Hadoop and MapReduce, improving big data operations efficiency by 45%.
- Implemented complex ETL processes using SSIS and AWS Glue, enabling seamless data integration and reducing load times by 35%.
- Developed custom machine learning models using TensorFlow and Scikit-learn to forecast consumer behavior, increasing accuracy by 30%.
- Crafted and maintained scalable and secure data pipelines in AWS, incorporating EC2 and S3 to enhance data storage and retrieval processes.
- Utilized Python, R, and SQL within Jupyter Notebook and PyCharm to perform data analysis and predictive analytics, improving decision-making insights by 25%.
- Spearheaded migrating legacy systems to Snowflake on Azure, optimizing data storage and access, which led to a 20% reduction in operational costs.
- Led cross-functional teams adopting CI/CD practices using Jenkins and CircleCI, improving deployment cycles and system reliability.
- Developed interactive dashboards in Power BI, improving KPI tracking & business reporting for executives.
- Conducted predictive analytics and statistical modeling, identifying trends that optimized marketing ROI by 20%.
- Optimized SQL queries & indexing, reducing query execution time by 25% and improving data retrieval speeds.
- Led data governance initiatives, ensuring compliance with GDPR, SOC 2, and internal security policies.

Morgan Stanley, TX | June 2023 - April 2024 | Data Analyst

- Collaborated with business users and IT teams to gather requirements, comprehensively understanding and articulating business processes and logic.
- Extracted data for reporting from diverse sources, including SQL Server, SQL Server Analysis Services, Azure SQL Database, and Azure SQL Data Warehouse.
- Engineered scalable ETL solutions with Apache Airflow, improving pipeline efficiency by 35%.
- Built real-time streaming applications using Apache Kafka, reducing data latency and improving response times.

- Integrated big data analytics using Pyspark, enabling machine learning models for predictive business insights.
- Designed and implemented automation test scripts using Python, enhancing the efficiency and reliability of testing Processes.
- Developed real-time analytical solutions using Apache Kafka and Flink, reducing data streaming latency by over 50%.
- Enhanced data visualization capabilities using Power BI and SSRS, creating dynamic reports that improved stakeholder understanding of key business metrics.
- Orchestrated the deployment of data lakes on the Google Cloud Platform, enhancing data aggregation and analysis capabilities.
- Automated data workflows using Apache Airflow, increasing process efficiency and reducing manual errors by 40%.
- Executed advanced statistical analyses using RStudio and Scala, identifying key trends and improving predictive accuracy by 18%.
- Optimized and maintained high-volume databases in MongoDB and Oracle, ensuring robust data integrity and availability.
- Collaborated with IT and business stakeholders to refine business intelligence strategies, leading to a 15% improvement in data-driven decision-making.
- Supported database architects by contributing to SQL database maintenance logs and reporting any identified issues.
- Constructed and optimized database components, including stored procedures, functions, triggers, indexes, and views, to facilitate data processing and cleansing.

Orion Technolab, India | May 2018 - Dec 2021 | Data Analyst

- Led a software development project utilizing the Waterfall methodology, ensuring successful delivery.
- Utilized SAS and SAS Enterprise Miner for advanced statistical analysis and predictive modeling, resulting in a 15% improvement in forecasting precision.
- Combined NumPy and Pandas functionalities to streamline data preprocessing tasks, maintaining data quality and consistency for analytical projects.
- Controlled the development and execution of a predictive analytics platform using Hadoop ecosystem tools, resulting in a 25% increase in marketing campaign effectiveness.
- Integrated Apache Spark with existing data systems to enhance processing speeds by up to 60%.
- Utilized machine learning algorithms such as SVM and Random Forests to develop classification models, improving customer segmentation accuracy by 22%.
- Designed and implemented a clinical data management system in a hybrid cloud environment, ensuring compliance with healthcare regulations and improving data accessibility by 30%.
- Conducted thorough data quality assessments using Python and Pandas, leading to a 25% improvement in data accuracy across projects.
- Automated report generation and data visualization using Tableau and Power BI, providing actionable insights that increased operational efficiency by 20%.
- Implemented data governance frameworks in line with GDPR requirements, enhancing data security and privacy measures and reducing compliance risks by 35%.
- Conducted extensive data analysis using Python, R, and SQL to extract, clean, and transform large datasets.
- Employed advanced data visualization tools, including Tableau, Power BI, and QlikView, to create insightful dashboards and reports, aiding in strategic decision-making.
- Applied scripting languages like HTML and CSS for web-based data presentation and analysis.
- Leveraged cloud platforms such as AWS, Azure, and GCP to manage, process, and analyze large-scale datasets, enhancing scalability and efficiency in data operations.
- Implemented rigorous data cleaning and wrangling processes, resulting in a 20% reduction in data errors.

ACADEMIC PROJECTS:

Data Analysis of Sales Data from a Superstore

- Analyzed store performance by identifying top customers, bestselling products, and high-revenue-generating states by assessing revenue and profits. Used Tableau, Google Cloud Platform (GCP), and SQL for in-depth data analysis.

Real-Time Data Pipeline for E-Commerce

- Designed a real-time data ingestion pipeline using Kafka & Spark Streaming, reducing order processing delays by 50%.
- Deployed AWS-based ETL workflows, enabling real-time dashboarding for executive decision-making.
- Created an automated data validation process, ensuring 99.8% data accuracy across sales and inventory records.

Cloud-Based Data Warehousing Optimization

- Migrated legacy SQL-based data warehouse to Snowflake, improving query execution speed by 200%.
- Designed data lake solutions to handle structured and unstructured data, enhancing analytical capabilities.
- Automated data ingestion and transformation processes, reducing pipeline failures by 30%.

CERTIFICATIONS:

AWS Certified Data Analytics – Specialty
Google Cloud Professional Data Engineer
Microsoft Azure Data Fundamentals