

Saikrishna Bondalapati | Data Engineer

Location: NC, USA || Ph: 475-766-9337 || Mail: manisankar1528@gmail.com || [LinkedIn](#)

SUMMARY

Data Engineer with 5+ years of experience designing and optimizing large-scale data infrastructure for retail, healthcare, and pharmaceutical domains. Proven ability to deliver 15–50% cost reductions (AWS/Snowflake) and 20–30% performance gains through scalable ETL pipelines, real-time streaming architectures, and efficient dimensional modeling. Combines technical expertise in cloud platforms with a focus on data governance, automation (CI/CD), and cross-functional collaboration to drive data driven decision making

SKILLS

Programming Languages: Python, Power Automate, HubSpot, Veeva CRM (Salesforce), Anaconda

Packages: NumPy, Pandas, Matplotlib, SciPy, Scikit-learn, TensorFlow, Seaborn

Visualization Tools: Tableau, Power BI, Advanced Excel, Power Query

Database: MySQL, PostgreSQL, SQL Server

Cloud Technologies: AWS - DynamoDB, S3, SNS, SQS, Lambda, Kinesis, Athena, Aurora, DMS, FireHose, Glue, Redshift, API Gateway, Snowflake, DBT, Airflow, Pyspark

Data Modelling: Dimensional Modeling, Star Schema, Snowflake Schema, Data Marts

Technical Skills: Machine Learning Algorithms, Statistics, Probability distributions, MS Excel, Hypothesis Testing, Regression Analysis, Data Visualization, Data Warehousing, Data Transformation, A/B Testing, GitHub

Methodologies: Agile, Scrum, Jira, Waterfall

PROFESSIONAL EXPERIENCE

Data Engineer

Asheville, NC, USA

Royal Food Import Corporation

Jan 2024 – Feb 2025

- Integrated AWS Kinesis with DynamoDB, Glue, S3, and Snowflake to enhance real-time data streaming and track historical price changes using SCD Type 2, optimizing data accessibility & reducing data storage costs by 15%.
- Developed Python-based ETL pipelines to migrate data to AWS S3, leveraging Snowpipe for ingestion into Snowflake, enabling real-time analytics in Power BI and reducing customer churn by 15%.
- Designed cross-region S3 data migration, optimizing storage costs (-15%) & maintaining version history for data redundancy.
- Integrated Apache Airflow, DBT & AWS Glue to automate ETL workflows, reducing processing time from 4 to 2 hours and improving data quality monitoring.
- Built ETL pipelines for SQL Server to AWS S3 via AWS DMS, processing raw data with AWS Glue & Snowflake, leveraging DBT for transformations.
- Developed CI/CD workflows with GitHub, DBT & Airflow, enabling version-controlled ETL pipelines & automated deployments for data reliability.
- Designed scalable dimensional models (star/snowflake schemas) for data warehouses, optimizing query performance and improved BI dashboard load times by 40%.
- Implemented robust data governance and security measures by configuring role-based access control (RBAC) in Snowflake and IAM policies in AWS to secure data across S3, Glue, and Snowflake.
- Built fault-tolerant ETL pipelines using AWS Glue with auto-retry, checkpointing, and Snowflake's multi-cluster warehouses for high availability.
- Configured AWS CloudWatch alerts for Glue job failures and Snowflake Query Performance Monitoring for query optimization.
- Worked on cross-functional Agile teams to deliver Retail data solutions, facilitating Scrum ceremonies and collaborating with business stakeholders to align project goals with food and safety compliance standards.

Data Engineer

Ridgefield, CT, USA

Boehringer Ingelheim Pharmaceuticals, Inc

Jan 2023 – Nov 2023

- Integrated the HCP data changes located in AWS DMS with NPI Data changes in AWS DMS through S3 bucket to identify discrepancies between previous and current data using **AWS Glue Data Catalog**.

- Strengthened data governance and security by implementing data validation checks, access control policies, and monitoring frameworks, improving compliance.
- Developed data cleansing and transformation scripts in Python and Spark, improving data quality and reducing data exploration time by 25%.
- Designed star-schema models and fact/dimension tables for healthcare analytics, improving OLAP query performance by 20% and enabling efficient Power BI/Tableau reporting
- Collaborated with cross-functional teams (product, data ops, data science) to build data-driven strategies, improving operational efficiency.
- Analyzed and managed comprehensive payor data for Boehringer Ingelheim, ensuring accuracy in reimbursement rates, coverage details, and claims processing, leading to a 15% improvement in operational efficiency.
- Designed patient data models to track treatment efficacy, adherence, and adverse events, contributing to improved patient care and safety.
- Developed a suite of automated reports in PowerBI for clinical teams, enhancing **Real-Time Tracking** of HEDIS measures and improving patient care strategies.
- Designed and implemented a real-time data architecture for seamless Tableau dashboard integration, enabling tracking of healthcare quality metrics such as patient satisfaction scores, readmission rates, and adherence to clinical guidelines.

Data Engineer

Dott Services Limited

Hyderabad,IND

July 2019 - July 2022

- Enhanced data quality and business metric calculations by leveraging **Spark** with AWS Glue Jobs, ensuring efficient pre-processing and transformation of raw data before loading it into Glue Data Catalog Tables.
- Collaborated with **cross-functional teams** to gather data requirements and delivered high-quality data models and optimized data warehouse tables, enabling efficient analysis by data analysts and data scientists.
- Designed an ETL pipeline to deliver high-quality data for regression models, enabling the analytics team to identify key factors impacting sales, resulting in **18% sales increase** in Q3 2021.
- Built a real-time data pipeline to process 3 years of sales data, this led to recommendations for 3+ senior sales managers to which customers should be prioritized, leading to an incremental sales lift of 15% in the next 180 days.
- Developed and implemented comprehensive in-built documentation for ETL workflows and data models using DBT, increasing team development productivity by 25% and facilitating a 30% faster onboarding process for new team members.
- Engineered a serverless AWS data pipeline integrating AWS Lambda, API Gateway, S3, and Snowflake, which enhanced the efficiency of low-throughput sequence data processing by 50%.
- Built scalable data pipelines for real-time and batch processing, integrating structured and unstructured data into a data lake to improve data accessibility and analytics workflows.
- Optimized data warehouse performance by improving data modeling, query optimization, and storage efficiency (15% cost savings).
- Developed and optimized SQL queries for efficient partitioning and indexing, improving query performance and analytics workloads across relational databases (Snowflake, SQL Server, PostgreSQL).
- Improved data cleansing by 25% using custom Python scripts, enhancing data accuracy & predictive analytics.
- Engineered cross-region data migration strategies, leveraging data lake architecture for disaster recovery, versioning, and 15% cost optimization.
- Implemented ETL pipelines to migrate data across on-prem and cloud platforms, transforming raw data into structured data warehouse layers for business reporting.
- Implemented a cloud-based data storage and processing solution using Amazon Web Services (AWS), including Amazon S3 for data storage, Amazon Athena for querying, and AWS Glue for ETL processes, resulting in a 50% reduction in infrastructure costs and improved scalability.

EDUCATION

University of Connecticut – Stamford, CT

Dec 2023

Major: **Master of Science (MS)** in Business Analytics and Project Management

GPA: 3.7/4.0

Indian Institute of Technology, Dhanbad

July 2019

Bachelor of Technology in Mine Economics and Engineering

GPA: 3.22/4.0