

SANDEEP SRIVATSAVA SINGARAJU

srivatsava.sandeep@gmail.com | (937) 784-9815 | www.linkedin.com/in/sandeep-srivatsava

SUMMARY:

Data Engineer with **5+ years** of experience designing, developing, and optimizing data pipelines and architectures across cloud (AWS, Azure, GCP) and on-premises environments. Proficient in Python, SQL, and Java for data processing, transformation, and analysis. Experienced in building scalable ETL workflows using Airflow, Luigi, and AWS Glue, and working with Big Data technologies like Hadoop, Spark, and Kafka. Strong expertise in data modeling, data warehousing (Snowflake, Redshift, BigQuery), and real-time streaming (Kafka Streams, Apache Flink). Adept in DevOps tools such as Docker, Kubernetes, Terraform, and CI/CD pipelines for cloud-based automation.

TECHNICAL SKILLS:

Cloud Platforms: Azure (Data Factory, Synapse, Data Lake, SSIS, ETL/ELT), AWS (S3, EC2, EMR, Redshift, RDS, Lambda, Kinesis, SNS, SQS, AMI, IAM, CloudFormation, Glue), GCP (BigQuery, Dataflow, Dataproc, Pub/Sub)

Big Data / Hadoop Ecosystem Hadoop (HDFS, Yarn, MapReduce, Hue, Pig, Hive, HCatalog, HBase, Sqoop, Zookeeper, Flume, Cloudera Manager, Kerberos), Spark (Core, SQL, PySpark, Scala), Kafka (Kafka Streams), Impala

Data Warehousing: Snowflake, Redshift, BigQuery.

ETL / Orchestration: Airflow, Luigi, AWS Glue, Informatica, NiFi, Talend.

Databases: Relational - Oracle, Microsoft SQL Server, MySQL, DB2, Teradata | NoSQL - MongoDB, Cassandra.

Programming Languages: Python, SQL, Java, Scala.

DevOps: CI/CD - Jenkins, GitLab | Containerization & Infrastructure - Docker, Kubernetes, Terraform.

Visualization: Tableau, Power BI, Looker.

Methodologies: Agile (Scrum), Waterfall, UML, Design Patterns, SDLC.

Collaboration Tools: JIRA, Confluence, Slack, Microsoft Teams.

Professional Development: Apache Flink, Drill, Tachyon.

WORK EXPERIENCE:

Client: McKesson Corporation, Irving, TX

April 2024 – Present

Role: Azure Data Engineer

Key Contributions:

- Designed and optimized **ETL pipelines** using **Azure Data Factory** for efficient data ingestion and transformation.
- Managed large datasets with **Azure SQL Database**, **SQL Managed Instance**, and **Cosmos DB**.
- Deployed scalable solutions via **Azure Kubernetes Service** and automated deployments through **Azure DevOps CI/CD**.
- Utilized **Pyspark**, **Scala**, and **Python** for data transformation and implemented advanced visualizations with **Power BI**.
- Integrated data workflows using **REST APIs**, **GraphQL**, and **Apache NiFi**.
- Coordinated distributed systems with **Zookeeper**; executed fast queries with **Google Big Query** and **Elasticsearch**.
- Processed big data with **HDInsight** using **Hadoop**, **Hive**, **Kafka**, and **Spark**, orchestrated workflows with **Apache Airflow**.
- Leveraged **Terraform** and **Bicep** for infrastructure-as-code deployments.
- Built cloud-based data warehousing solutions using **Snowflake**.

Environment: Azure Data Factory, Azure SQL Database, AKS, Azure DevOps, Pyspark, Scala, Power BI, REST APIs, Hadoop, Hive, Kafka, Cosmos DB, Apache NiFi, Apache Airflow, Zookeeper, Google Big Query, Elasticsearch, GitLab CI/CD, Snowflake.

Client: Comerica Incorporated, Dallas, TX

May 2023 – March 2024

Role: AWS Data Engineer

Key Contributions:

- Designed and implemented **ETL pipelines** using **AWS Glue**, **Apache NiFi**, and **Talend** to extract, transform, and load financial data into **Amazon S3** and **Redshift**.
- Managed relational and NoSQL databases with Amazon RDS, Aurora, and DynamoDB, ensuring high availability and secure access.
- Developed real-time data processing workflows using AWS Lambda triggered by events from Kinesis and S3.
- Processed large-scale financial datasets using Amazon EMR with Apache Spark, Hive, and Hadoop ecosystem tools (HDFS, MapReduce, Sqoop, Impala).
- Built machine learning models with Amazon SageMaker to predict trends and enhance risk management.
- Enabled full-text search and analytics with Amazon Elasticsearch and created interactive dashboards with Power BI and Tableau.

- Integrated security measures, including encryption and access controls, to protect sensitive financial data.

Environment: Amazon S3, Redshift, AWS Glue, EMR, RDS, Aurora, Lambda, Kinesis, SageMaker, DynamoDB, Elasticsearch, Apache NiFi, Talend, Hadoop, Spark, Hive, HDFS, MapReduce, Sqoop, Impala, Zookeeper, Flume, Power BI, Tableau.

Client: QBE Insurance Group, Hyderabad, India

March 2021 – December 2022

Role: GCP Data Engineer

Key Contributions:

- Designed and optimized ETL pipelines using Google Cloud Dataflow to integrate diverse data sources into BigQuery, resulting in a 20% reduction in processing time.
- Implemented event-driven architectures with Google Cloud Pub/Sub for real-time streaming, ensuring timely processing of insurance claims and transactions.
- Managed large-scale datasets with Google Cloud Dataproc, Apache Hadoop, and Spark to support advanced analytics.
- Developed and optimized data models in BigQuery for high-performance querying of extensive insurance datasets.
- Leveraged Cloud SQL and Spanner to maintain high availability and consistency for customer and claims data.
- Created real-time dashboards using Google Looker and BigQuery, enabling proactive monitoring of key performance indicators.
- Collaborated with data scientists to deploy predictive models with TensorFlow and the Google AI Platform for improved claims forecasting and customer retention.
- Enhanced pipeline flexibility with Apache Beam on Google Cloud Dataflow, streamlining insurance claims analysis.
- Integrated Snowflake as a cloud data platform to consolidate data storage and analytics, supporting seamless scalability.
- Automated data workflows with Python, Pandas, NumPy, and PySpark, ensuring efficient data manipulation and reporting.
- Utilized Hadoop technologies (HDFS, MapReduce) to process vast volumes of insurance data, supporting efficient claims processing and fraud detection.

Environment: Google Cloud Dataflow, BigQuery, Pub/Sub, Dataproc, Apache Hadoop, Spark, Cloud SQL, Spanner, Looker, TensorFlow, Google AI Platform, Apache Beam, Snowflake, Python, Pandas, NumPy, HDFS, MapReduce.

Client: ABB Ltd, Bangalore, India

June 2019 – February 2021

Role: Data Engineer

Key Contributions:

- Developed Databricks ETL pipelines using notebooks, Spark DataFrames, Spark SQL, and Python scripting.
- Automated cloud resource provisioning on Azure using Terraform for Infrastructure-as-Code.
- Designed and configured back-end databases and applications, leveraging SAP SD module transactions to generate sales reports.
- Managed large datasets using Pandas and SQL, ensuring data integrity and performance.
- Built and maintained Docker container clusters managed by Kubernetes; implemented CI/CD pipelines with Jenkins.
- Transformed and cleansed data using SQL, Python, and PySpark; executed ETL jobs with HiveQL, Presto SQL, and Spark SQL.
- Constructed data pipelines with Python, PySpark, Hive SQL, and Presto; integrated data from both relational and non-relational sources (including Cassandra and HDFS).
- Created interactive dashboards with Power BI and Tableau for real-time analytics.
- Developed scalable data processing workflows using Java and Apache Beam on Google Cloud Dataflow.
- Engineered Spark applications with Scala and Spark SQL for multi-format data extraction, transformation, and aggregation.

Environment: Databricks, Spark DataFrames, Spark SQL, Python, Terraform, Azure, SAP, Pandas, SQL, Docker, Kubernetes, Linux, Bash, Git, Jenkins, PySpark, Hive, Cassandra, HDFS, Power BI, Tableau, Scala, Java, Apache Beam, Google Cloud Dataflow.

CERTIFICATIONS:

- Google Cloud Certified – Associate Cloud Engineer
- AWS Certified Data Engineer – Associate
- Microsoft Certified: Azure Data Fundamentals

EDUCATION:

University of Dayton, Masters in Computer Science, USA

SASTRA University, Bachelor of Technology in Electronics and Instrumentation Engineering, India