

SNEHA REDDY G. AI/ML ENGINEER 4+ YEARS EXPERIENCE Dallas, TX snehareddy9720@gmail.com +1 726-219-6859 LinkedIn	
Summary	AI/ML engineer with 4+ years’ experience turning messy data into dependable systems using Python, PyTorch, scikit-learn, SQL, Spark, and Airflow, strong at problem framing, feature design, calibration, testing, and clear partner communication. Comfortable with Docker/Kubernetes and MLOps (MLflow, CI/CD), focused on applied NLP and tabular risk modeling. Delivered fraud screening models reaching AUC 0.89 and reducing manual reviews ~20%, plus a retrieval-augmented Q&A that improved top-3 recall ~12% and lower median handle time ~25% while keeping p95 latency under ~150 Ms.
Skills	Programming and Frameworks: Python, SQL, Bash, Java, React js Machine Learning: scikit learn, XGBoost, LightGBM, CatBoost, feature engineering, model selection, cross validation, calibration, hyperparameter tuning Deep Learning: PyTorch, TensorFlow, Keras, fine tuning, distributed training, mixed precision, model compression, ONNX NLP and GenAI: Transformers, embeddings, RAG, prompt design, evaluation, sentence transformers, LangChain, LlamaIndex, vector search, rerankers MLOps and Platforms: MLflow, model registry, Feast, DVC, CI CD for ML, Kubeflow basics, Argo Workflows, monitoring with Evidently and Why Labs Data Engineering: Spark, Airflow, Prefect, dbt, Delta Lake, Apache Iceberg, Kafka, Kinesis, Parquet, data modelling, Serving and APIs: FastAPI, gRPC, async I O, batching, caching, Triton or vLLM, TensorRT, A B testing hooks Cloud and Infra: AWS including Sage Maker, S3, Lambda, IAM, ECR, EKS, Kubernetes, Docker, Helm, GCP and Azure familiarity, Terraform basics Warehousing and Analytics: Databricks, Snowflake, Big Query, Redshift, SQL performance tuning, dashboards Observability: Prometheus, Grafana, Open Telemetry, alerting, log aggregation Testing and Quality: pytest, Great Expectations, data contracts, reproducibility, experiment tracking DevOps Tooling: Git, GitHub Actions, GitLab CI, artifact management, secrets management Security and Compliance: PII handling, access controls, VPC, encryption at rest and transit, SOC2 and GDPR awareness
Experience	AI/ML Engineer <i>Deloitte</i> Aug 2024 – Present USA - Built RAG service using FAISS and cross-encoder rerankers, improving top-3 recall by 15% and reducing support agents’ average handle time by 28% across seven teams. - Containerized PyTorch models with Triton, ONNX, and TensorRT, cut P95 latency from 220ms to 135ms, while autoscaling Kubernetes reduced compute spend by 23% monthly. - Implemented MLflow and Feast for tracking and feature management, versioning 500GB features and automating 60% of hyperparameter sweeps, cancel time-to-deploy from 14 days to 5. - Designed Airflow and Spark ETL pipelines processing 1.2TB daily, with Great Expectations data contracts, throttled schema breakages by 60% and ameliorated SLA adherence to 70%. - Established Prometheus and Evidently dashboards and alerts, lowering model incidents from seven to two, reducing mean time to detect by 41% and recovery by 36%. - Partnered with security to implement privacy checks for NLP models, achieved 60% PII redaction accuracy and cleared internal audit on first pass with no findings. Jr.AI/ML Engineer <i>IBM</i> Jan 2020 – Jul 2023 India - Designed Spark and Airflow recommendation pipeline covering preparation, candidate generation, ranking, and cold-start via embeddings, streamlined CTR by 11% and pruned P95 latency to 150ms. - Redesigned XGBoost fraud detection with calibrated probabilities and investigator queues, lifted AUC from 0.84 to 0.91, tightened false positives 32%, accelerated dispute resolution by 29%. - Tuned search relevance with a lightweight ranker and holdout metrics, increased NDCG@10 by 9%, maintained sub-120ms P95 through feature selection and compact architectures in production. - Established Great Expectations data quality program detecting schema drift, leakage, and anomalies, discontinue incident rate from seven to two month and lowered MTTR by 41%. - Built React dashboards and Java Spring Boot APIs for governance, implemented RBAC and audit trails serving 180+ users, optimized PostgreSQL queries, improving report latency 46%. - Delivered Java microservices with Kafka and Redis caching for recommendation telemetry, processed 1.8M events/day, strengthened idempotency, tuned SQL indexes to reduce query CPU by 35%.
Education	M.S., Computer Science University of North Texas May 2025 B. Tech, Electronics & Communication Engineering Gokaraju Rangaraju Institute of Engineering and Technology Sep 2020
Projects	Claims Anomaly Scoring with Gradient Boosting - Spark and XGBoost pipeline on 2.7M claims, with 180 engineered features, cross validated AUC 0.89, analyst review volume shaved 22% on heldout set. - Isotonic calibration upscaled Brier score 17%, SHAP flagged three risky codes, revising sampling rules, lifting confirmed fraud discovery rate from 3.1% to 3.8%. Campus Policy Q&A with Retrieval-Augmented Generation - Indexed 28,000 handbook pages using FAISS, cross-encoder reranker reached 84% top-3 recall, dropping unresolved queries in pilots from 105 to 49 weekly. - vLLM served 7B model with quantization, median latency 210ms at 90 RPS batching, hallucinations decreased 31% using retrieval citations, safety moderation, and tighter prompt constraints.
Certifications	- AWS Certified Machine Learning – Specialty - Google Cloud Professional Machine Learning Engineer - Certified Kubernetes Administrator (CKA) - NVIDIA DLI — Optimizing Inference with TensorRT