

Srujankumar Polepally
polepallysrujan@gmail.com
+1 (603)-264-6061 [Linkedin](#), [Medium](#), [Github](#)

SUMMARY

Results-driven Data Analyst with 7+ years of experience in collecting, analyzing, and visualizing large datasets to drive business insights and strategic decision-making. Proficient in SQL, Python, Power BI, Tableau, and cloud platforms (AWS, Azure, GCP), with a strong foundation in data modelling, reporting, and predictive analytics. Adept at designing and optimizing ETL pipelines, data transformation processes, and real-time analytics to improve data accuracy and accessibility. Experienced in leveraging Apache Spark, Snowflake, and Azure Synapse to process and analyze complex datasets efficiently. Passionate about translating data into actionable insights, enabling organizations to make data-driven decisions and optimize business performance.

Key Strengths:

- ✓ **Big Data & Cloud:** Expertise in Hadoop (Spark, Hive, Impala) and cloud-native data engineering on Azure (Synapse Analytics, Data Factory, Databricks), GCP (Big Query), AWS (Redshift).
- ✓ **ETL/ELT & Automation:** Built scalable ETL pipelines with Apache Airflow, Azure Data Factory, and DBT. Automated workflows reducing processing times by 60%.
- ✓ **Streaming & Real-Time Analytics:** Developed low-latency data ingestion with Apache Kafka, Azure Event Hubs, and Pub/Sub, enabling real-time decision-making.
- ✓ **AI/ML & Predictive Analytics:** Integrated AI/ML workflows into data pipelines using Azure ML, BigQuery ML, and predictive analytics, enabling advanced decision-making and automation.
- ✓ **Data Modelling & Storage Optimization:** Engineered efficient data models using Snowflake, Cosmos DB, and Big Query, improving query performance by 40%.
- ✓ **CI/CD & DevOps:** Implemented CI/CD pipelines with Azure DevOps, Terraform, and Docker, reducing deployment failures by 30%.
- ✓ **Data Governance & Compliance:** Ensured GDPR & HIPAA compliance with RBAC, encryption, and Collibra-based governance.
- ✓ **Leadership & Mentorship:** Led cross-functional teams, mentoring junior engineers and improving productivity by 20%.
- ✓ **Scalability & Containerization:** Designed cloud-native, event-driven architectures with Azure Functions, Google Cloud Functions, and containerized solutions using Docker & Kubernetes (AKS) for fault-tolerant and scalable data processing.

TECHNICAL SKILLS:

Category Technologies	Tools
Big Data & Frameworks	Apache Spark (PySpark, Scala), Hadoop (HDFS, MapReduce), Hive, Impala, HBase, Sqoop, Flume, Delta Lake, Flink
Cloud Platforms	Azure (Azure Synapse, Data Factory, Databricks, ADLS, AKS), AWS (Redshift, S3, Step Functions, Lambda, EMR), GCP (BigQuery, Dataflow, Dataproc, Cloud Storage)
Data Processing & ETL	Apache Airflow, Azure Data Factory, Google Cloud Dataflow, DBT, Informatica, Talend, Collibra
Real-Time & Streaming	Apache Kafka, Spark Streaming, Azure Event Hubs, Google Pub/Sub, Apache Beam, Kinesis
Databases & Data Modeling	Snowflake, Azure Synapse, BigQuery, PostgreSQL, Oracle, SQL Server, Cosmos DB, DynamoDB, MongoDB
Programming & Scripting	Python (Pandas, NumPy), SQL, Scala, Java, Shell Scripting
Machine Learning & AI	Azure ML, BigQuery ML, Spark ML, TensorFlow, Scikit-learn, Natural Language Processing and Computer Vision, Large Language Models (LLM), Machine Learning with Generative AI
DevOps & CI/CD	Azure DevOps, Jenkins, GitHub Actions, Terraform, Docker, Kubernetes (AKS, EKS, GKE)
Data Governance & Security	RBAC, Encryption, GDPR & HIPAA Compliance, Data Masking, Azure Purview
Business Intelligence & Visualization	Power BI, Tableau, Looker (LookML), QlikView, Workday Reports
Advanced Analytics & Forecasting	Time-Series Forecasting, Regression Analysis, Predictive Analytics, Metadata Framework, Statistical Analysis, Bayesian Inference, ARIMA, SARIMA, Prophet, Machine Learning Time-Series Algorithms, customer segmentation & targeted marketing analytics.

PROFESSIONAL EXPERIENCE

Sr Data Engineer
TechVedika, India

March 2021 to July 2023

Responsibilities:

- Designed and automated scalable ETL pipelines using Apache Spark, Azure Data Factory, and Snowflake, optimizing data ingestion from Python, MySQL, and SQL Server into Azure Data Lake and Snowflake, leveraging SQL for querying, data transformation, and reducing data processing times by 40%
- Developed real-time and batch processing pipelines using Apache Spark (PySpark, Scala), Kafka, and machine learning models, enabling low-latency streaming analytics and event-driven architectures for high-volume data ingestion, utilizing Python for data processing and automation.
- Optimized Spark RDD transformations, reducing job execution times and computational overhead, improving overall pipeline performance using Python and SQL for efficient data manipulation.
- Established CI/CD pipelines for ETL workflows, improving deployment efficiency and reducing operational downtime. Worked with stakeholders to gather requirements and define KPIs, optimizing reporting dashboards in Power BI and Tableau. Reducing manual effort by 50% and increasing the accuracy of business-critical reports.
- Applied machine learning and generative AI techniques to enhance data processing and analytics capabilities, providing actionable insights from large datasets. resulting in a 20% improvement in forecasting accuracy for sales and inventory management.

Data Engineer
Clean Harbors, India.

Jan 2017 to Nov 2019

Responsibilities:

- ETL Development & Automation: Designed and implemented ETL pipelines using Apache Spark, Azure Data Factory, and Python, automating data ingestion, transformation, and integration from multiple RDBMS and APIs into cloud-based data lakes and warehouses.
- Data Integration & Processing: Developed scalable data integration workflows to extract, transform, and load data from structured, semi-structured, and unstructured sources into Azure Data Lake, Snowflake, and Hive, ensuring seamless data processing.
- Cloud Data Architecture: Leveraged Azure Databricks, Synapse Analytics, and HDInsight to design, deploy, and manage cloud-based data solutions, optimizing storage, performance, and scalability for large-scale analytical workloads and big data processing.
- SQL & Query Optimization: Wrote and optimized complex SQL queries across Snowflake, Hive, and SQL Server, improving query performance through indexing, partitioning, and efficient data modeling.
- Data Quality & Governance: Established data quality frameworks, implementing automated validation checks, monitoring tools, and governance policies to ensure clean, accurate, and compliant data in alignment with GDPR and industry standards.
- Optimized SQL queries and data models in Snowflake and SQL Server, reducing report generation time by 30%.

Data Analyst
myGwork, India

Aug 2014 to Dec 2016

Responsibilities:

- Collected, cleaned, and analyzed large datasets from SQL databases, Excel, and APIs, providing actionable insights to improve business strategy and decision-making. Developed interactive dashboards and reports in Power BI, Tableau, and Excel, enabling stakeholders to visualize KPIs and trends for data-driven decisions.
 - Applied statistical techniques like regression analysis and time-series forecasting to identify trends and predict business outcomes. Optimized complex SQL queries for data extraction, manipulation, and aggregation, ensuring fast and accurate reporting.
 - Cleaned and pre-processed data using Python (Pandas) and SQL, ensuring high-quality datasets for analysis. Collaborated with stakeholders to define KPIs and develop customized reports aligned with organizational goals. Automated reporting processes using advanced Excel functions (VLOOKUP, PivotTables, Power Query, macros) to improve efficiency.
 - Partnered with cross-functional teams (marketing, finance, operations) to integrate data and deliver actionable insights. Built and maintained data models in Excel and SQL, integrating CRM, ERP, and financial data for a comprehensive view of business performance.
-

Education:

Master's(M.sc) Degree in Data Science

Aug 2023 – Dec 2024

[New England College](#), New Hampshire, USA

Master's Degree in Micro Data Analysis (Business Intelligence)

Jan 2020 - Jan 2021

[Högskolan Dalarna](#), Borlänge, Sweden

Thesis: Creating Financial Database for Education & Research: Using WEB SCRAPING Technique.

Publications: - [Weblink](#)

Projects:

Predictify Healthcare Application

- The application, that takes your current lifestyle, food habits and exercise patterns as an input, sit on your mobile and validate that data, compare it with the open health data available, and predict that if you continue today's lifestyle and food habits, you'll be in terms of health in 10-15-20 years.
- Gathering complex data from different sources and storing it into databases. Data Sources: WHO, Agency for Healthcare Research and Quality, Centers for Medicare and Medicaid Services (CMS), Medicaid/Medicare & Hospital Management.
- OMOP (Observational Medical Outcomes Partnership) is used designed to standardize healthcare data across different sources, enabling large-scale observational research and analytics.

- Use Generative AI & LLM models to automate CI/CD deployment pipelines with Terraform and Azure DevOps, cutting deployment time by 40%. Optimized data models in Snowflake and Azure Synapse, improving query performance by 35% and reducing cloud storage costs by 25%.
- Established data governance frameworks using Azure Purview and Collibra, ensuring GDPR and HIPAA compliance. Migrated on-premises data warehouses to cloud environments (GCP & Azure), improving scalability and reducing maintenance costs.

Instant Payment Fraud Detection & Prevention

- The rise of instant payments (UPI, Faster Payments, etc.), financial institutions must detect fraud within milliseconds. This system continuously monitors live transactions and flags suspicious activities based on user behavior analysis. Data from mobile banking apps, ATMs, and online transactions are streamed into Kafka topics and processed using Spark Structured Streaming in Azure Databricks.
- Automated CI/CD deployment pipelines with Terraform and Azure DevOps, cutting deployment time by 40%. Optimized data models in Snowflake and Azure Synapse, improving query performance by 35% and reducing cloud storage costs by 25%.
- Established data governance frameworks using Azure Purview and Collibra, ensuring GDPR and HIPAA compliance.
- Migrated on-premises data warehouses to cloud environments (GCP & Azure), improving scalability and reducing maintenance costs.

Real-Time Credit Risk Assessment

- Banks and financial institutions need to assess credit risk before approving loans. Traditional credit scoring models rely on historical data, but real-time models can incorporate live financial behavior, such as transaction patterns, spending habits, and income fluctuations.
- This project ingests real-time financial transactions using Apache Kafka and processes them with Azure Databricks and Spark Streaming. A machine learning model trained in past defaults and repayment behaviors predicts the risk of loan applicants. If the model detects high-risk behavior, the system triggers alerts for manual review or adjusts the interest rate dynamically.
- Optimized SQL queries and data models in Snowflake and SQL Server, reducing report generation time by 30%.
- Collaboration & Data-Driven Insights: Worked closely with Data Scientists, Analysts, and Business Units to build and optimize data pipelines supporting predictive models, business intelligence reports, and advanced analytics.