

Tejaswi Vempatapu Data Engineer USA +1(940) 629- 8592 tejaswivempatau@gmail.com LinkedIn 	
SUMMARY	Data Engineer with 5.5+ years of experience designing high-performance data pipelines, real-time streaming architectures, and cloud-native data platforms across AWS, Azure, and Hadoop. Skilled in PySpark, Kafka, Talend, and Redshift, with a proven record of accelerating analytics, optimizing CI/CD workflows, and driving enterprise-grade data governance using tools like Apache Ranger and Collibra.
TECHNICAL SKILLS	Programming: Python, SQL, Scala, PySpark Big Data Frameworks: Apache Spark, Hadoop Cloud Platform: AWS(EMR, S3, Glue, Redshift), Azure (Data Factory) ETL Tools: Talend, Apache NiFi Data Modeling: Star Schema, Snowflake Schema, Dimensional Modeling, Amazon Redshift Streaming & Real-Time Processing: Apache Kafka, Spark Streaming, Event-Driven Pipelines Databases: PostgreSQL, MongoDB Data Visualization: Power BI, Tableau Data Governance & Security: IAM, Apache Ranger CI/CD & Version Control: CircleCI, Git Project Tools: Jira, Confluence
PROFESSIONAL EXPERIENCE	U.S Bank, USA Oct 2024 – Current Data Engineer - Developed high-performance Spark ETL pipelines on AWS EMR, improving processing speed by 35% and reducing SLA breaches by 40%. - Built Kafka-based real-time streaming for fraud detection, processing 10M+ events/day with <3s latency. - Architected an AWS data lake (S3, Glue, Redshift), eliminating silos and improving query speed by 40%. - Designed star/snowflake schemas and Power BI semantic models, enabling 100+ standardized KPIs for finance reporting. - Migrated legacy ETL to AWS Glue, cutting maintenance by 25% and accelerating delivery cycles. - Enforced data governance using Collibra, IAM, and Ranger, ensuring SOX, GDPR, and PCI-DSS compliance. - Automated CI/CD with CircleCI, reducing deployment failures by 60% and halving release cycles. - Built serverless pipelines with AWS Lambda & S3, enhancing scalability and reducing infrastructure cost. - Deployed containerized Spark jobs via Docker/Kubernetes, improving portability and resource efficiency. - Led development of predictive fraud models, boosting accuracy by 30% and reducing risk exposure. Mphasis, India Nov 2021–Aug 2023 Client: Delta Airlines Company Data Engineer - Migrated 120+ TB of Hadoop workloads to Azure Data Lake Gen2, cutting infrastructure costs by 30%, improving data retrieval speeds by 45%, and enabling scalable access across 15+ analytics teams. - Processed 12M+ semi-structured JSON and AVRO records daily using optimized PySpark pipelines, reducing data transformation time by 40% and increasing downstream reliability by 33%. - Automated 25+ complex ingestion workflows using Azure Data Factory and Airflow, achieving 99.95% uptime and decreasing manual intervention time from hours to under 15 minutes per incident. - Created 10+ high-fidelity logical and physical data models in powering Power BI dashboards used by 200+ stakeholders for route planning, crew utilization, and delay prediction analysis. - Designed a Kafka-based log aggregation pipeline to process 5M+ logs/day, improving operational alerting speed by 50% and cutting incident resolution time from 3 hours to under 90 minutes. - Integrated Apache Ranger and custom policy templates across Azure data pipelines, aligning with GDPR and PCI-DSS standards and reducing non-compliance incidents by 65%. Tata Consultancy Services, India Aug 2018 – Nov 2021 Client: British Airways & USAA Bank Data Engineer - Architected Spark and data pipelines to manage and refine 15+ TB of daily financial and airline data, driving a 40% boost in data throughput and ensuring timely compliance reporting. - Devised real-time fraud monitoring frameworks using Kafka Streams and Spark Structured Streaming, leading to a 25% improvement in threat detection accuracy and cutting alert latency by 80%. - Streamlined SQL execution for Redshift and Athena, cutting complex query runtimes by 30%, which accelerated insights for high-priority decision-making across finance and customer operations. - Rolled out automated delivery pipelines and Shell scripting for 50+ big data jobs, increasing deployment frequency and slashing release time by 65%. - Assembled a cloud-compatible lake house architecture blending Hive, S3, and HDFS, enabling unified access for 12+ departments and cutting data retrieval lag by 45%. - Instituted data protection mechanisms—masking, encryption, and access tokenization—that secured sensitive records and ensured 100% audit pass rate across GDPR and PCI DSS benchmarks.
EDUCATION	Master of Science in Advanced Data Analytics Dec 2024 University of North Texas, TX, USA Bachelor of Technology in Electronics and Communication Engineering Apr 2018 GITAM University, India
CERTIFICATION	AWS Cloud Practitioner