

TYSON HARRISON

harristyson.ai@gmail.com • <https://github.com/techtitan91> • Virginia, United States

Experienced AI/ML Engineer with a strong track record of building and deploying intelligent systems across diverse industries, including healthcare, fintech, and SaaS. Skilled in developing end-to-end ML pipelines and applying advanced techniques in natural language processing, computer vision, and generative AI. Known for translating complex problems into practical AI-driven solutions that scale in production. Passionate about staying ahead of emerging trends in AI and turning innovation into measurable business outcomes

Experience

MAY 2020 – PRESENT

Lead AI/Machine Learning Engineer | [Kallidus](#)

- Led multiple AI/ML initiatives, driving end-to-end pipeline development, from data ingestion and model training to deployment and monitoring, while mentoring teams and delivering production-ready solutions with measurable business impact.
- Designed and deployed an end-to-end voice-based meeting assistant, generating minutes, summaries, and actionable insights; easily integrable with Zoom and similar platforms.
- Built and optimized LLM-powered pipelines (RAG, QLoRA, LangChain, Pinecone) within a serverless AWS Lambda architecture for scalable and cost-effective inference.
- Implemented Dockerized ML workflows and integrated with existing systems for production readiness and reproducibility.
- Established CI/CD pipelines for model versioning and automated deployment, reducing delivery time from weeks to days. Applied MLOps best practices using MLflow, Kubernetes, Dashboards and Docker.

SEP 2017 – APR 2020

Senior AI/Machine Learning Engineer | [ToMarket](#)

- Led design and deployment of ML infrastructure and algorithms, owning projects end-to-end from concept to production.
- Drove cross-functional collaboration with engineering, data, and product teams to ensure successful ML solution delivery.
- Delivered Medical Claim Automation using NLP and deep learning, reducing claim rejections and boosting efficiency.
- Built automated data labeling pipelines and deployed NER models to extract sensitive data from text/audio, supporting customer call workflows.
- Developed scalable OCR systems for parsing unstructured tabular documents across varied formats.

JUN 2015 – AUG 2017

Machine Learning Engineer | [Advekit](#)

- Contributed to customer segmentation and sentiment analysis models, enhancing targeted marketing and improving retention rates.
- Automated data scraping and preprocessing pipelines using BeautifulSoup, Selenium, NumPy, and Pandas, significantly reducing data preparation time.
- Actively pursued continuous learning and shared insights on emerging trends, driving innovation within the team.

Skills

Programming & Frameworks

Python, PyTorch, TensorFlow, Keras, Scikit-Learn, Django, FastAPI, Pandas, NumPy, SciPy, PySpark, Spark MLlib, LangChain, LangGraph, N8N, Voice AI, Dask.

Deployment & DevOps | Visualization & Monitoring:

Docker, Kubernetes, CI/CD GitLab, Jenkins), Model Deployment & Monitoring, RESTful API Development, Cloud Platforms AWS, Azure, GCP, Grafana, Tableau, Power BI, Matplotlib, Seaborn, Model Drift & Data Drift Monitoring, Dashboard Design for ML Operations.

Machine Learning & AI:

Supervised & Unsupervised Learning, Deep Learning (CNNs, RNNs, Transformers), Natural Language Processing (NLP, Sentiment Analysis, Large Language Models (LLMs), Agentic AI, Computer Vision & Image Classification, Predictive Modeling & Forecasting, Recommendation Systems (Collaborative Filtering, Deep Learning), Reinforcement Learning, Knowledge Graphs, Retrieval-Augmented Generation (RAG).

Data Engineering & Pipelines:

Data Wrangling & Feature Engineering, Exploratory Data Analysis (EDA), ETL Pipelines, Data Ingestion & Transformation, Airflow, Databricks, SQL & NoSQL (Elasticsearch, MongoDB, PostgreSQL), Data Versioning & Validation

Education

Bachelor of Science in Computer Science | UOS

Projects

CISCO: Cisco AgenticOps radically simplifies workflows, breaks down silos, and enables real-time, cross domain execution with AI agents—built for speed and scale.

Yardee: Our platform combines advanced hybrid RAG (Retrieval Augmented Generation) technology with agentic AI to synthesize information across both structured databases like databases and CSV files and unstructured documents like PDFs. It doesn't just search - it reasons, understands context, and can take intelligent actions like creating tickets or booking appointments.